

DL/Vision/CNNs

Microsoft ML OpenHack Series

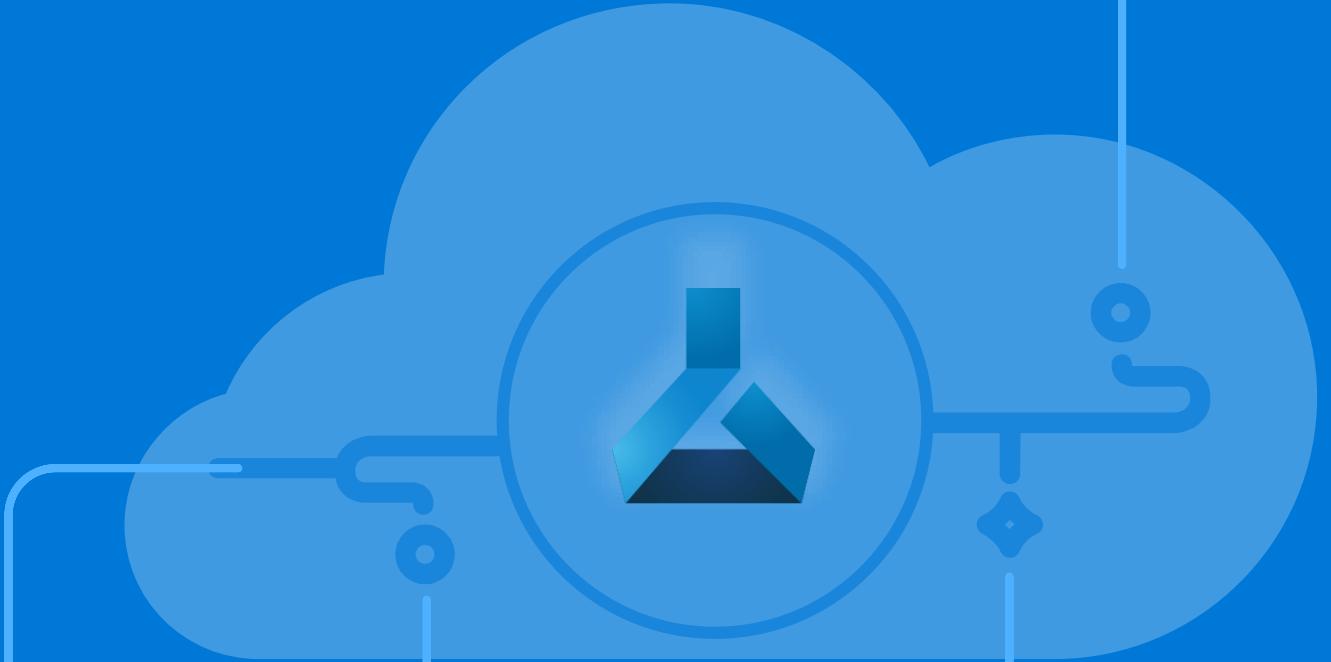
Tim Scarfe, Ph.D

Principal Software Engineer @ Microsoft
Commercial Software Engineering (CSE)

@ecsquendor 

 Subscribe

youtube.com/machinelearningatmicrosoft



- Home
- Trending
- Subscriptions
- Library
- History
- Watch Later
- Purchases 51
- Liked videos



Machine Learning at Microsoft

755 subscribers

CUSTOMISE CHANNEL

YOUTUBE STUDIO (BETA)

HOME

VIDEOS

PLAYLISTS

CHANNELS

DISCUSSION

ABOUT



Uploads ► PLAY ALL



A Joe Rogan style round table meditation on comput...
81 views • 3 weeks ago



Paper Reviews Call 004 -- (Mikolov)...
153 views • 3 weeks ago



Paper Reviews Call 006 - GloVe
179 views • 3 weeks ago



Paper Reviews Call 005 - Compact Prediction Trees
60 views • 3 weeks ago



Paper Reviews Call 003 -- Learning fine-grained image...
188 views • 1 month ago

FEATURED CHANNELS



Siraj Raval

SUBSCRIBED



Microsoft

SUBSCRIBE



Brandon Rohrer

SUBSCRIBED



Microsoft Research

SUBSCRIBE

- SUBSCRIPTIONS
- CloudTechTV ...
 - Jordan Knight
 - 3Blue1Brown
 - Robert Miles
 - ciruselvirus
 - Michael Lanzetta
 - The Artificial In...
 - Show 11 more

MORE FROM YOUTUBE

YouTube Movies

POPULAR CHANNELS ON YOUTUBE



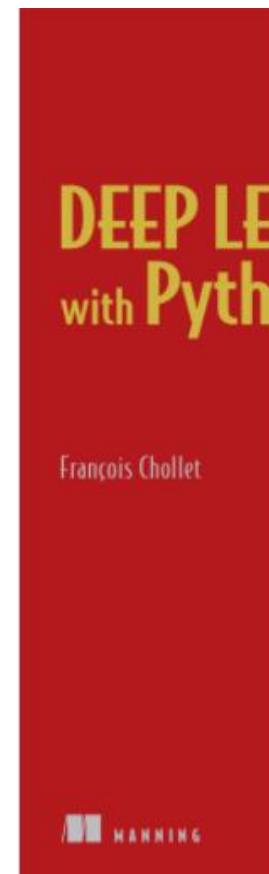
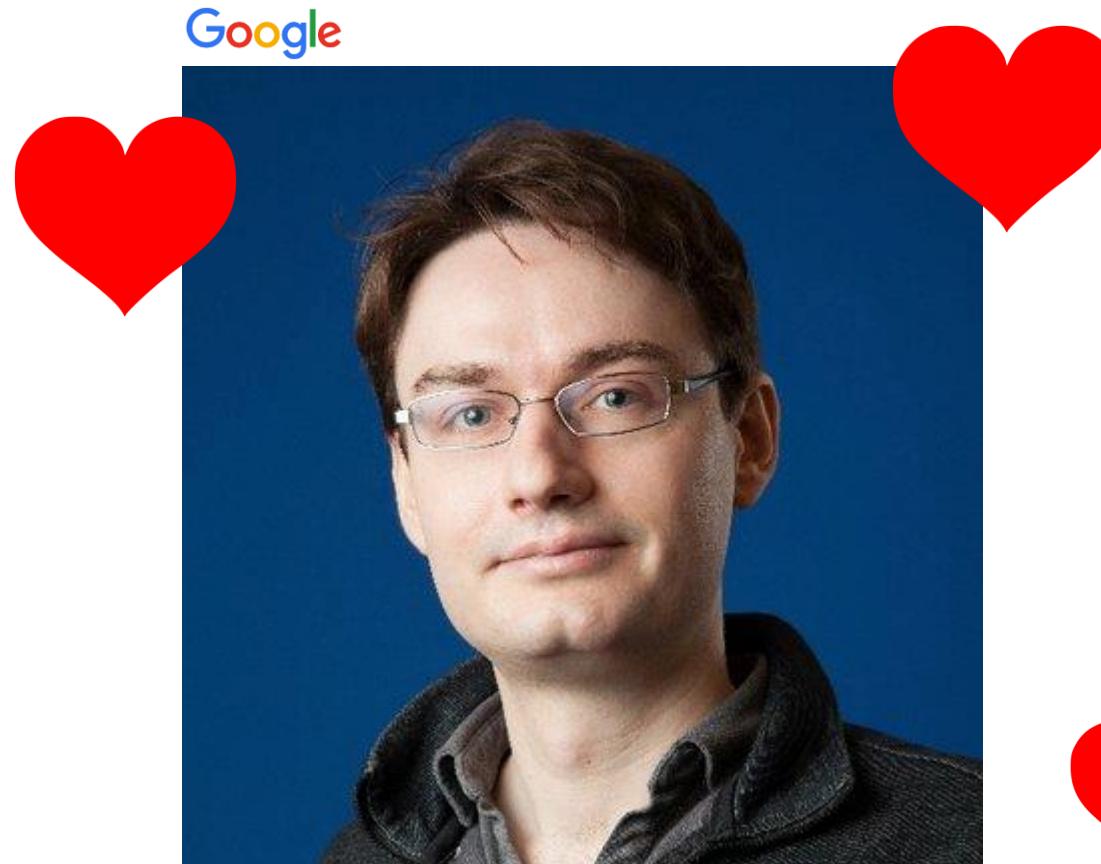
A PRACTICAL GUIDE TO DEEP LEARNING

Tess Ferrandez - Microsoft - @TessFerrandez

Tess Ferrandez

Principal Software Engineer @ Microsoft

**Francois Chollet. The man. The legend.
Please buy his book!**



<https://github.com/fchollet/deep-learning-with-python-notebooks/>

<https://www.manning.com/books/deep-learning-with-python>

I have marked his material with © Francois Chollet

External Test Image Population



axes



boots



carabiners



crampons



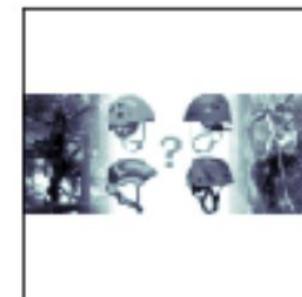
gloves



hardshell_jackets



harnesses



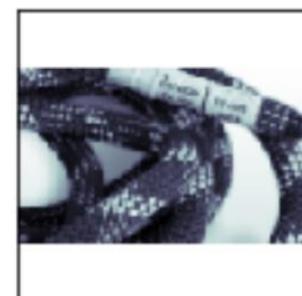
helmets



insulated_jackets



pulleys



rope



tents

https://storagetimscarfe.blob.core.windows.net/openhack/gear_images_test.zip

Experiment	F1 Validation	F1 Test/External	Difference
Histogram 20 bins Trees	0.55	0.11	0.44
Raw Pixels trees	0.66	0.15	0.51
SVM HOG	0.29		
Pixels->PCA + SVM	0.9	0.18	0.72
CNN raw pixels	0.91	0.23	0.68
Raw Pixels + Noise + Trees	0.34		
<u>Pixels+Noise->PCA + SVM</u>	0.68	0.03	0.65
CNN raw pixels + noise	0.94	0.03	0.91
Noise + Augmented 10k trees	0.15		
Noise + Augmented 20k PCA SVM	0.49	0.04	0.45
Noise + Augmented 20k HOG PCA SVM	0.48	0.1	0.38
Augmented 20K CNN	0.83	0.29	0.54
Augmented 20K HOG + PCA	0.61	0.25	0.36
Augmented 20K + PCA + SVM	0.8	0.21	0.59
Noise + Augmented 20K CNN	0.80	0.12	0.68
Transfer learning - Train just on external test set		0.39	
Transfer learning Augmented 20K CNN	0.89	0.50	0.39

Screenshot of a GitHub repository page for `mlopenhack`.

The repository is forked from `cauldnz/mlopenhack`. It has 1 unwatched star, 2 forks, and 1 issue.

Code navigation tabs: Code, Pull requests 0, Projects 0, Wiki, Insights, Settings.

Branch: master → `mlopenhack / notebooks / Challenge3.ipynb`

Comment by Tim Scarfe: There was a bug in the last one, I accidentally applied noise to the te... 577ba01 5 days ago

2 contributors

File details: 7.42 MB, Download, History, View, Delete.

Content of `Challenge3.ipynb`:

First download the gear images using;

- curl -O https://challenge.blob.core.windows.net/challengefiles/gear_images.zip
- curl -O https://storagetimscarfe.blob.core.windows.net/openhack/gear_images_test.zip
- unzip *.zip

Or in windows powershell;

- wget "https://challenge.blob.core.windows.net/challengefiles/gear_images.zip" -OutFile gear_images.zip
- wget "https://storagetimscarfe.blob.core.windows.net/openhack/gear_images_test.zip" -OutFile gear_im...
- unzip gear_images.zip

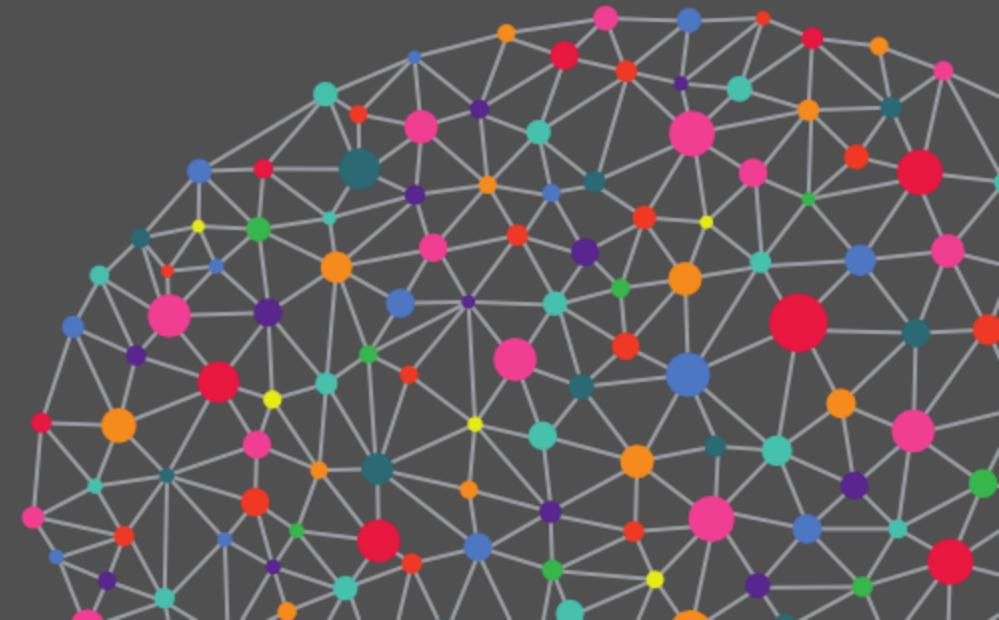


<https://www.youtube.com/watch?v=ZNC-ZNHcQ88>

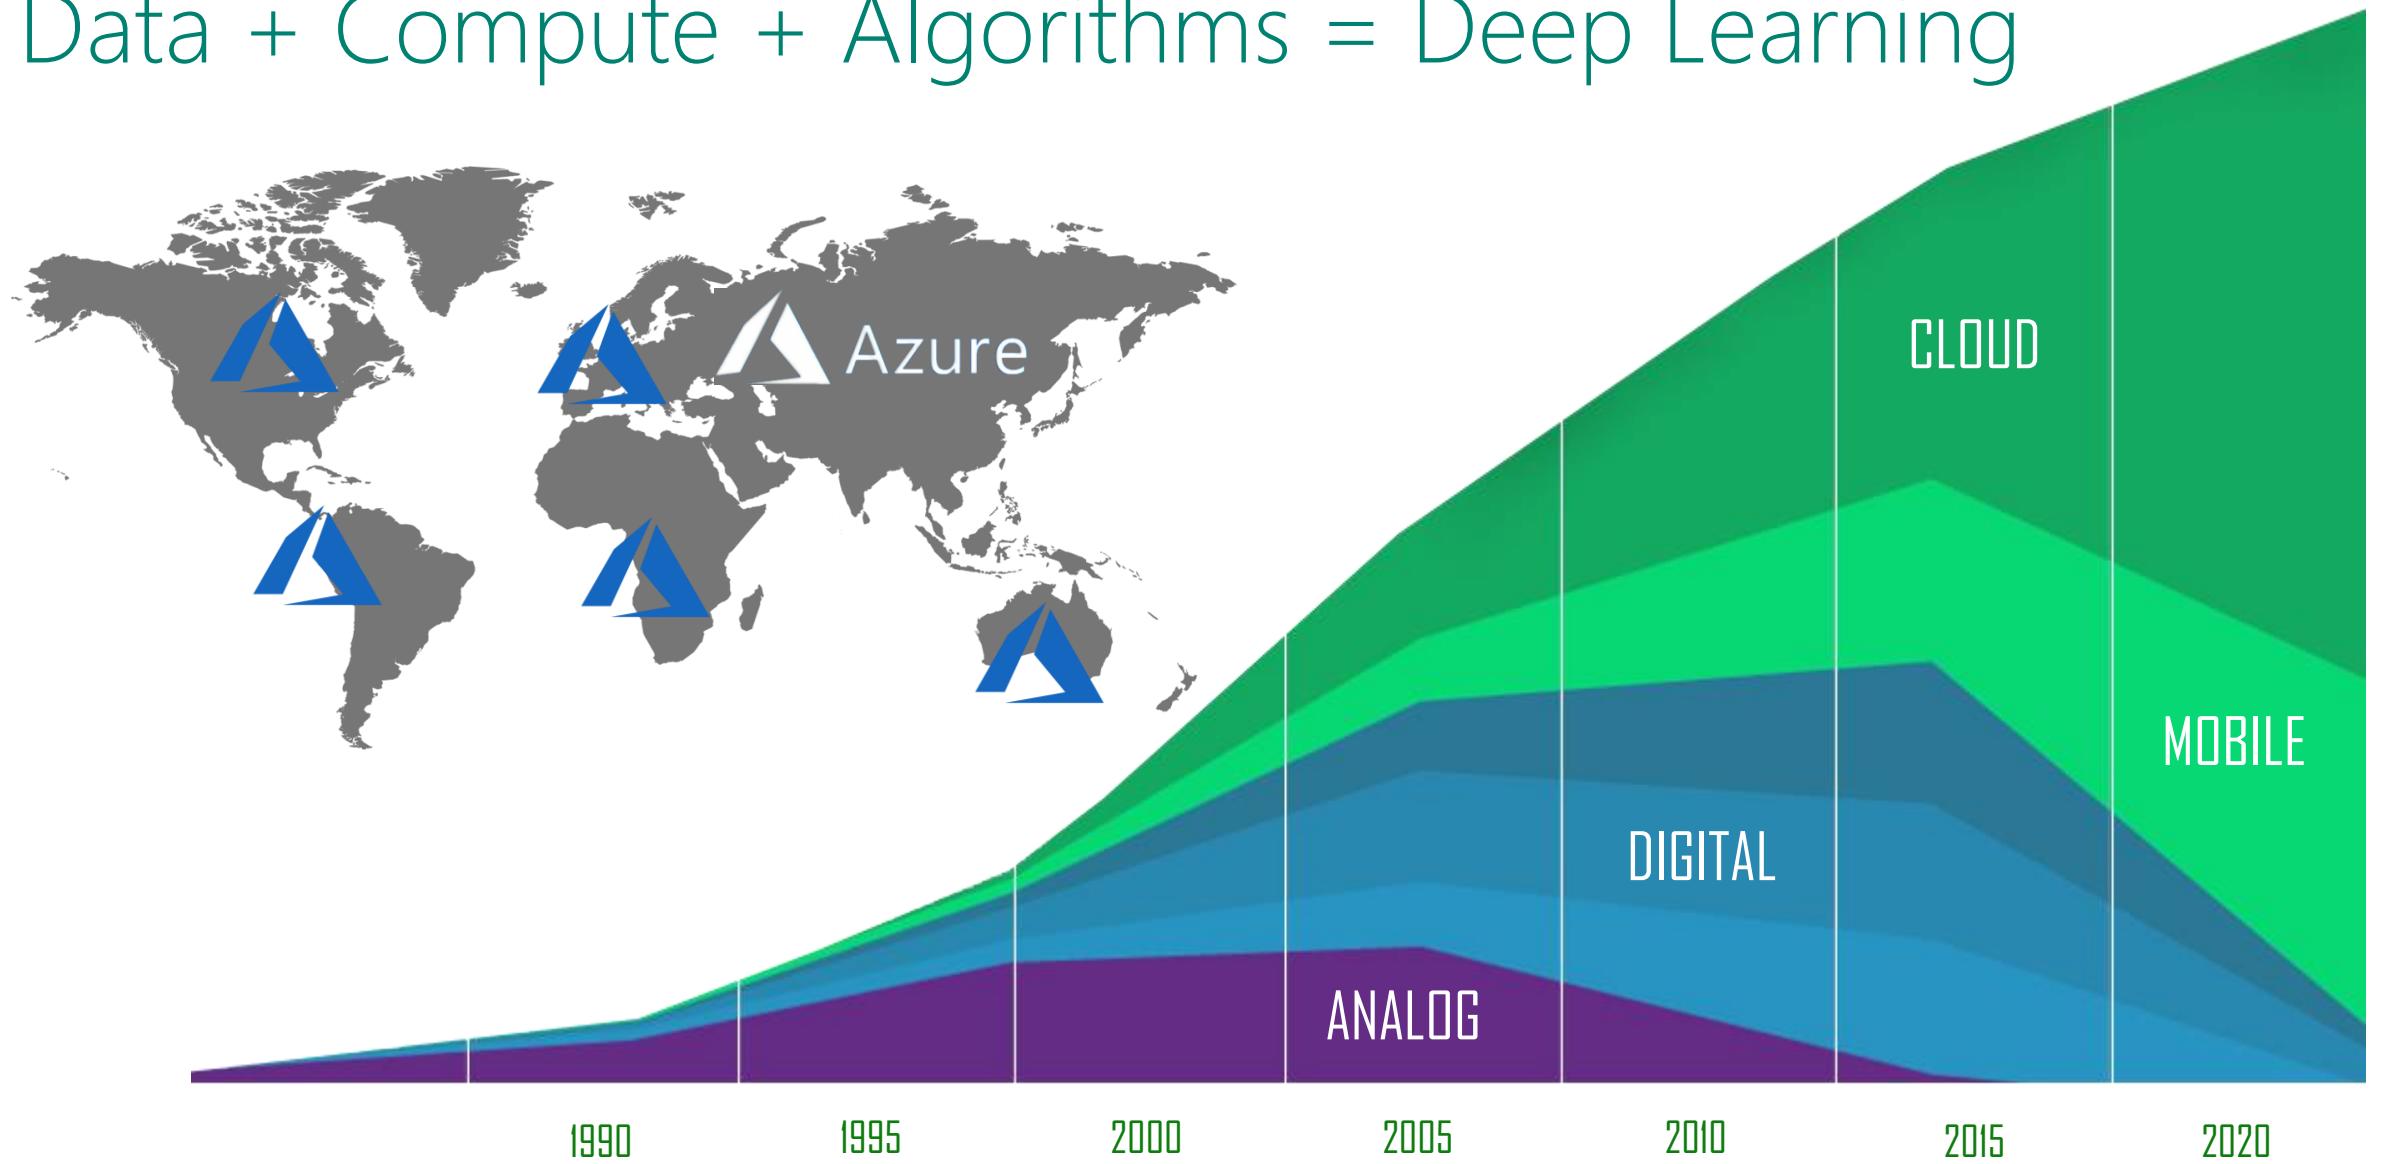
TALK OUTLINE



- High level concepts of deep learning
- Advanced topics
- Old days of vision
- How does convolution work?
- MNIST example in Keras code
- Data Augmentation on Cats+Dogs dataset
- Transfer Learning
- Visualising CNNs



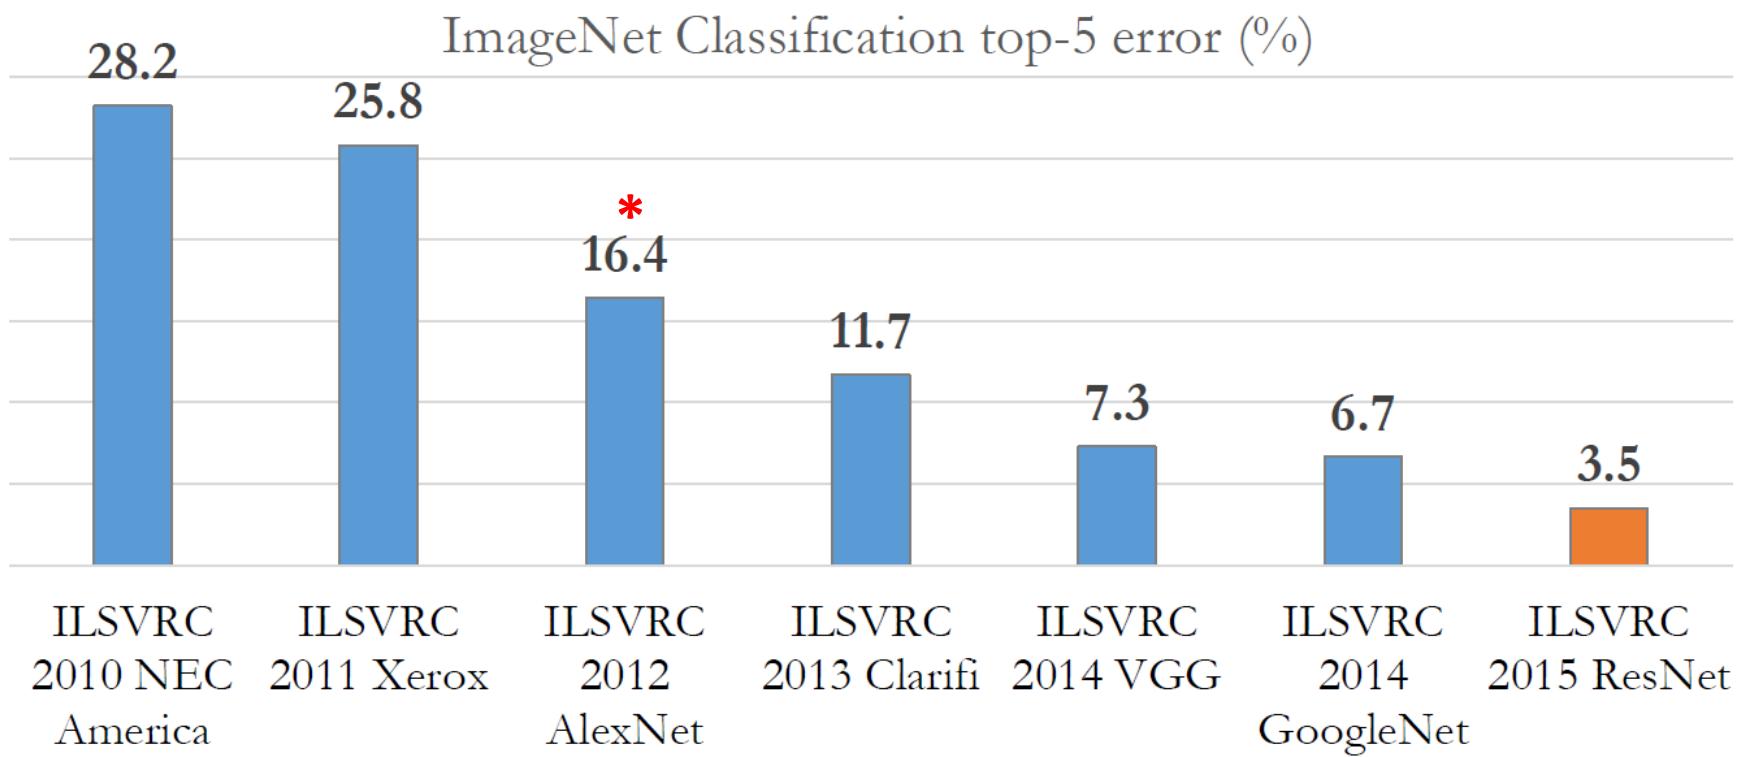
Data + Compute + Algorithms = Deep Learning



ImageNet



Alex Krizhevsky



- Alex Krizhevsky wins the ImageNet competition in 2012
- AlexNet
- Heralds a new era of AI
- Same story in language processing, speech recognition etc

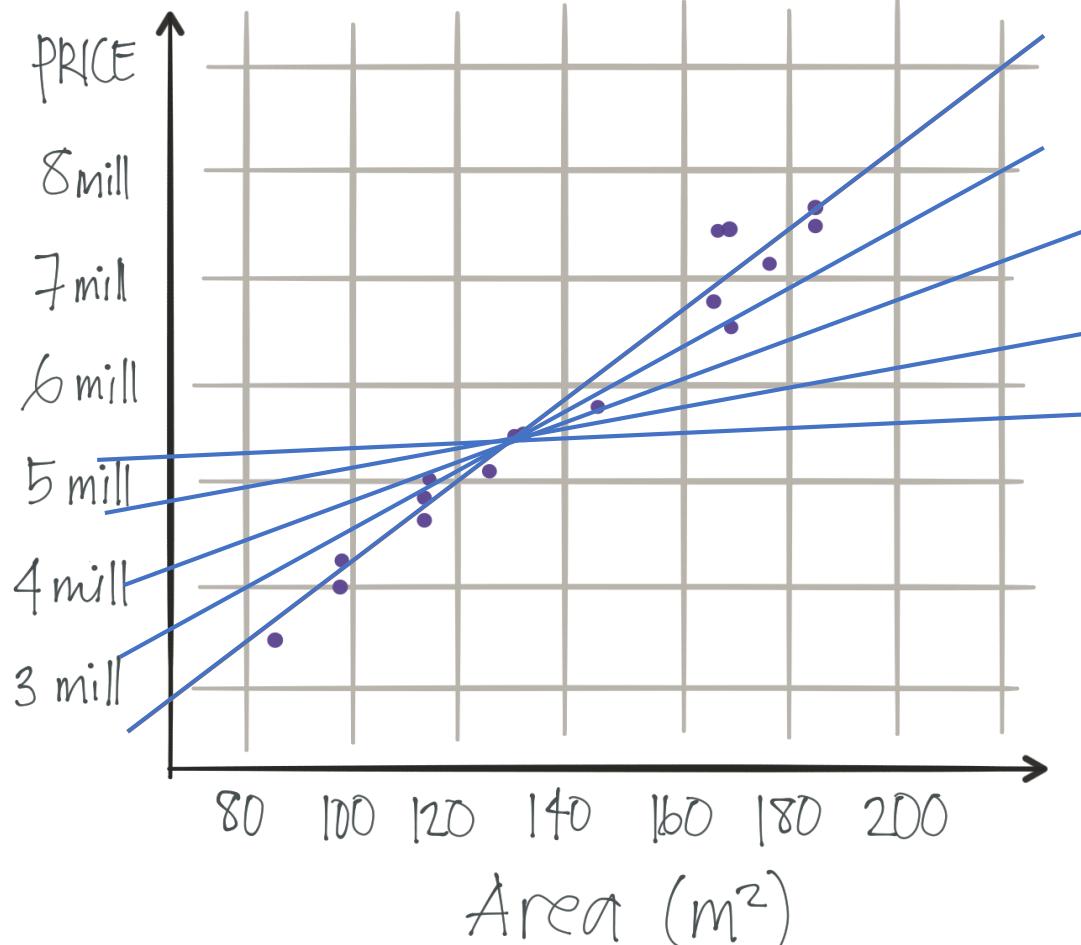


AREA (m ²)	TYPE	ZIP-CODE	PRICE (kr)
134	HOUSE	90210	5 495
115	TOWNHOME	11436	4700
167	HOUSE	90210	7500
185	HOUSE	11436	7775
84	HOUSE	79021	3500
98	TOWNHOME	11436	4000

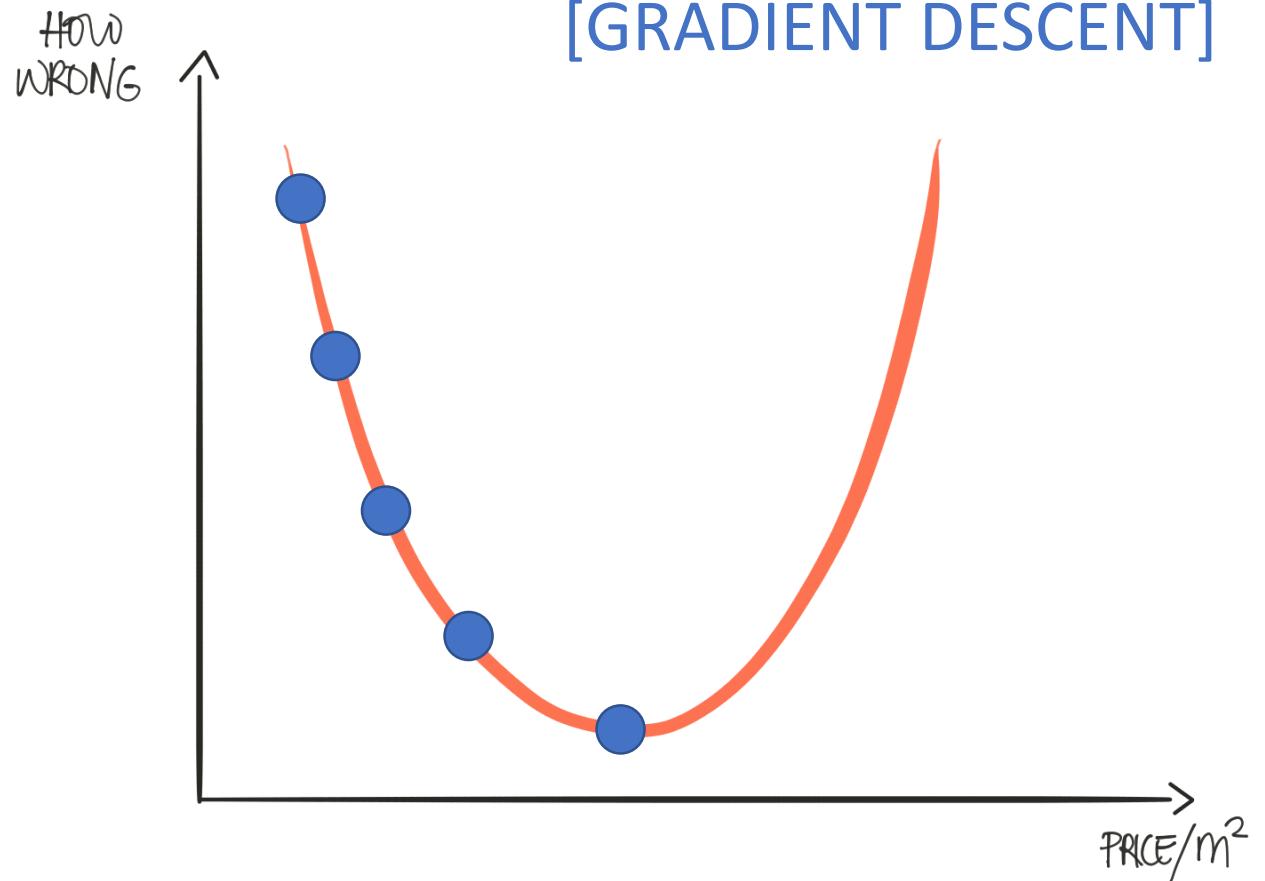
```
int EstimatePrice(...){  
    price = 10000 +  
            6700 * area_in_sqm +  
            20000 * has_pool +  
            10000 * new_kitchen +  
            5000 * neighborhood_quality;  
    return price;  
}
```

Price = b + w1*area_in_sqm + w2*has_pool + ...

[LINEAR REGRESSION]
[GRADIENT DESCENT]



$$\text{Price} = b + w_1 * \text{area_in_sqm}$$

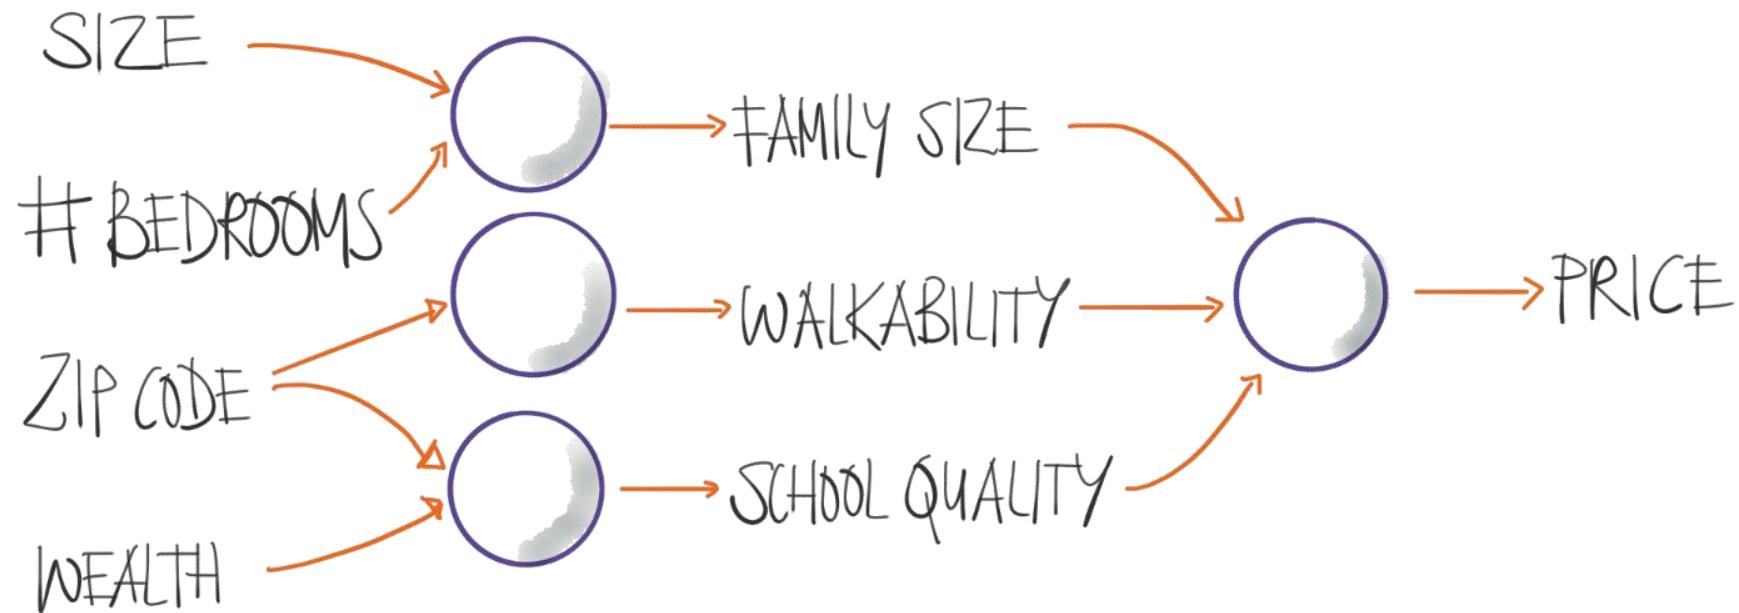


[LINEAR REGRESSION]

```
int EstimatePrice(...){  
    price = 10000 +  
            6700 * area_in_sqm +  
            20000 * has_pool +  
            10000 * new_kitchen +  
            5000 * neighborhood_quality;  
    return price;  
}
```

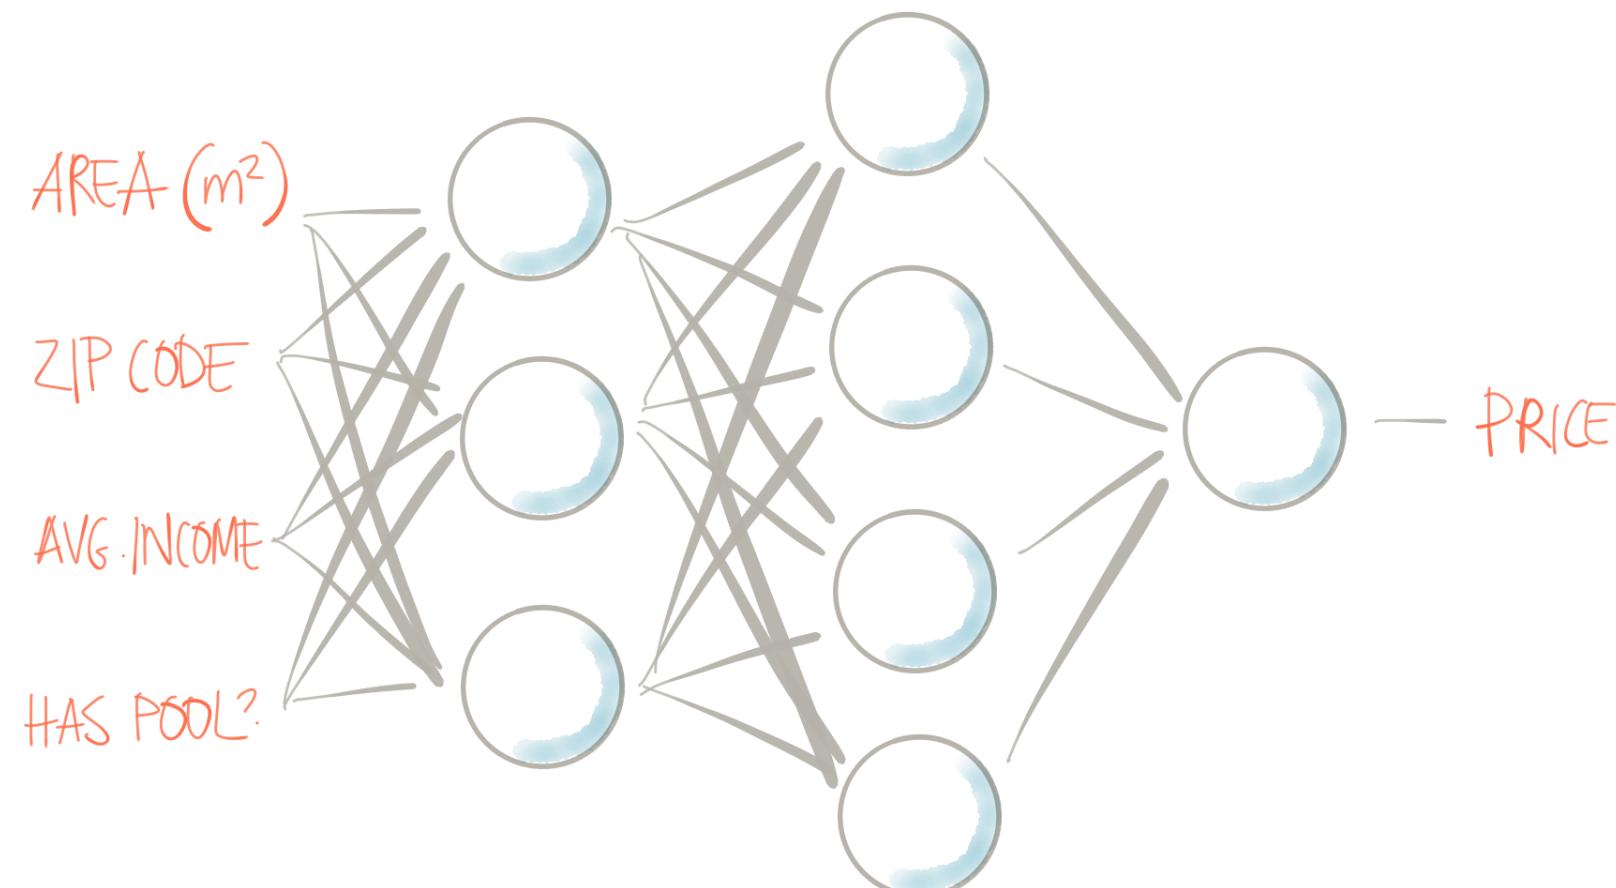
Price = $b + w_1 \cdot \text{area_in_sqm} + w_2 \cdot \text{has_pool} + \dots$

HOUSE PRICES



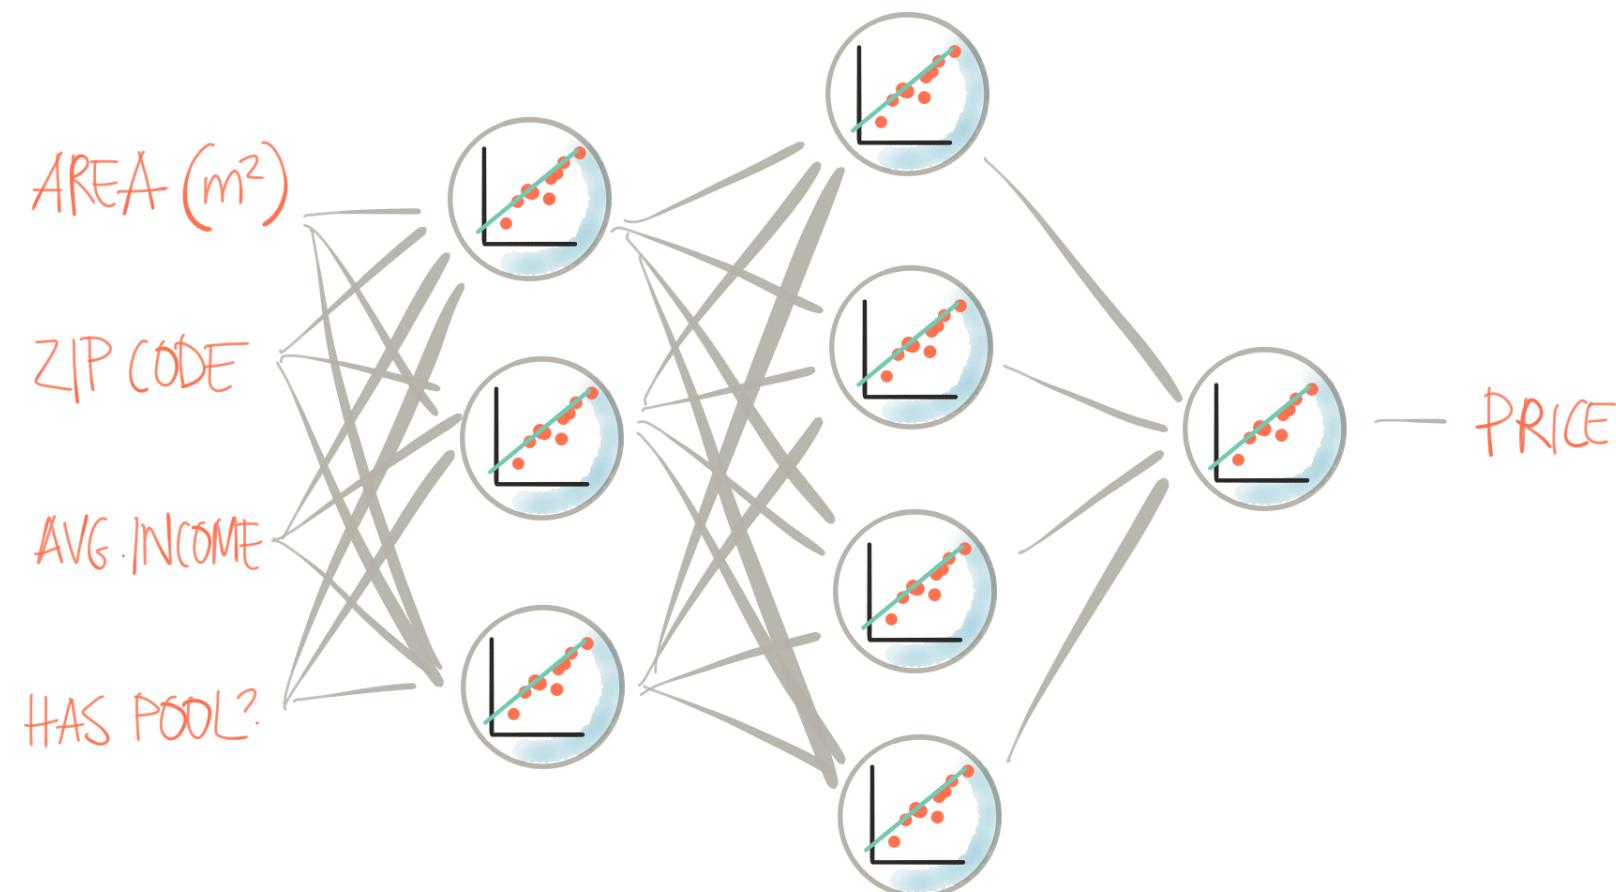
[NEURAL NET]

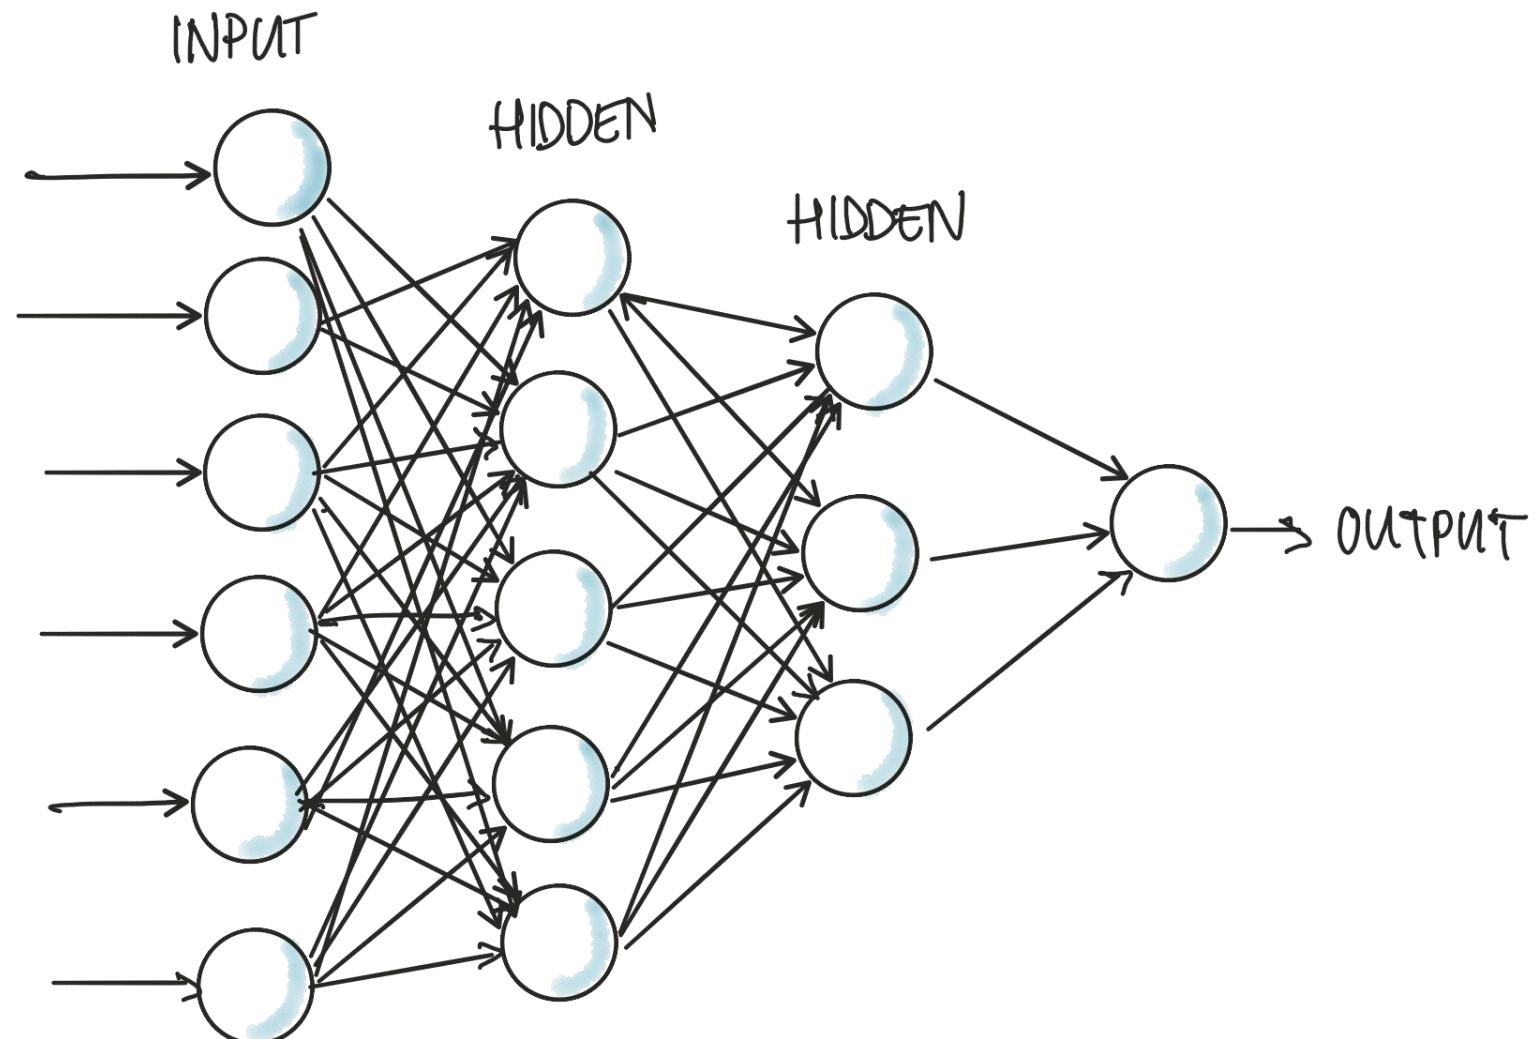
HOUSE PRICES

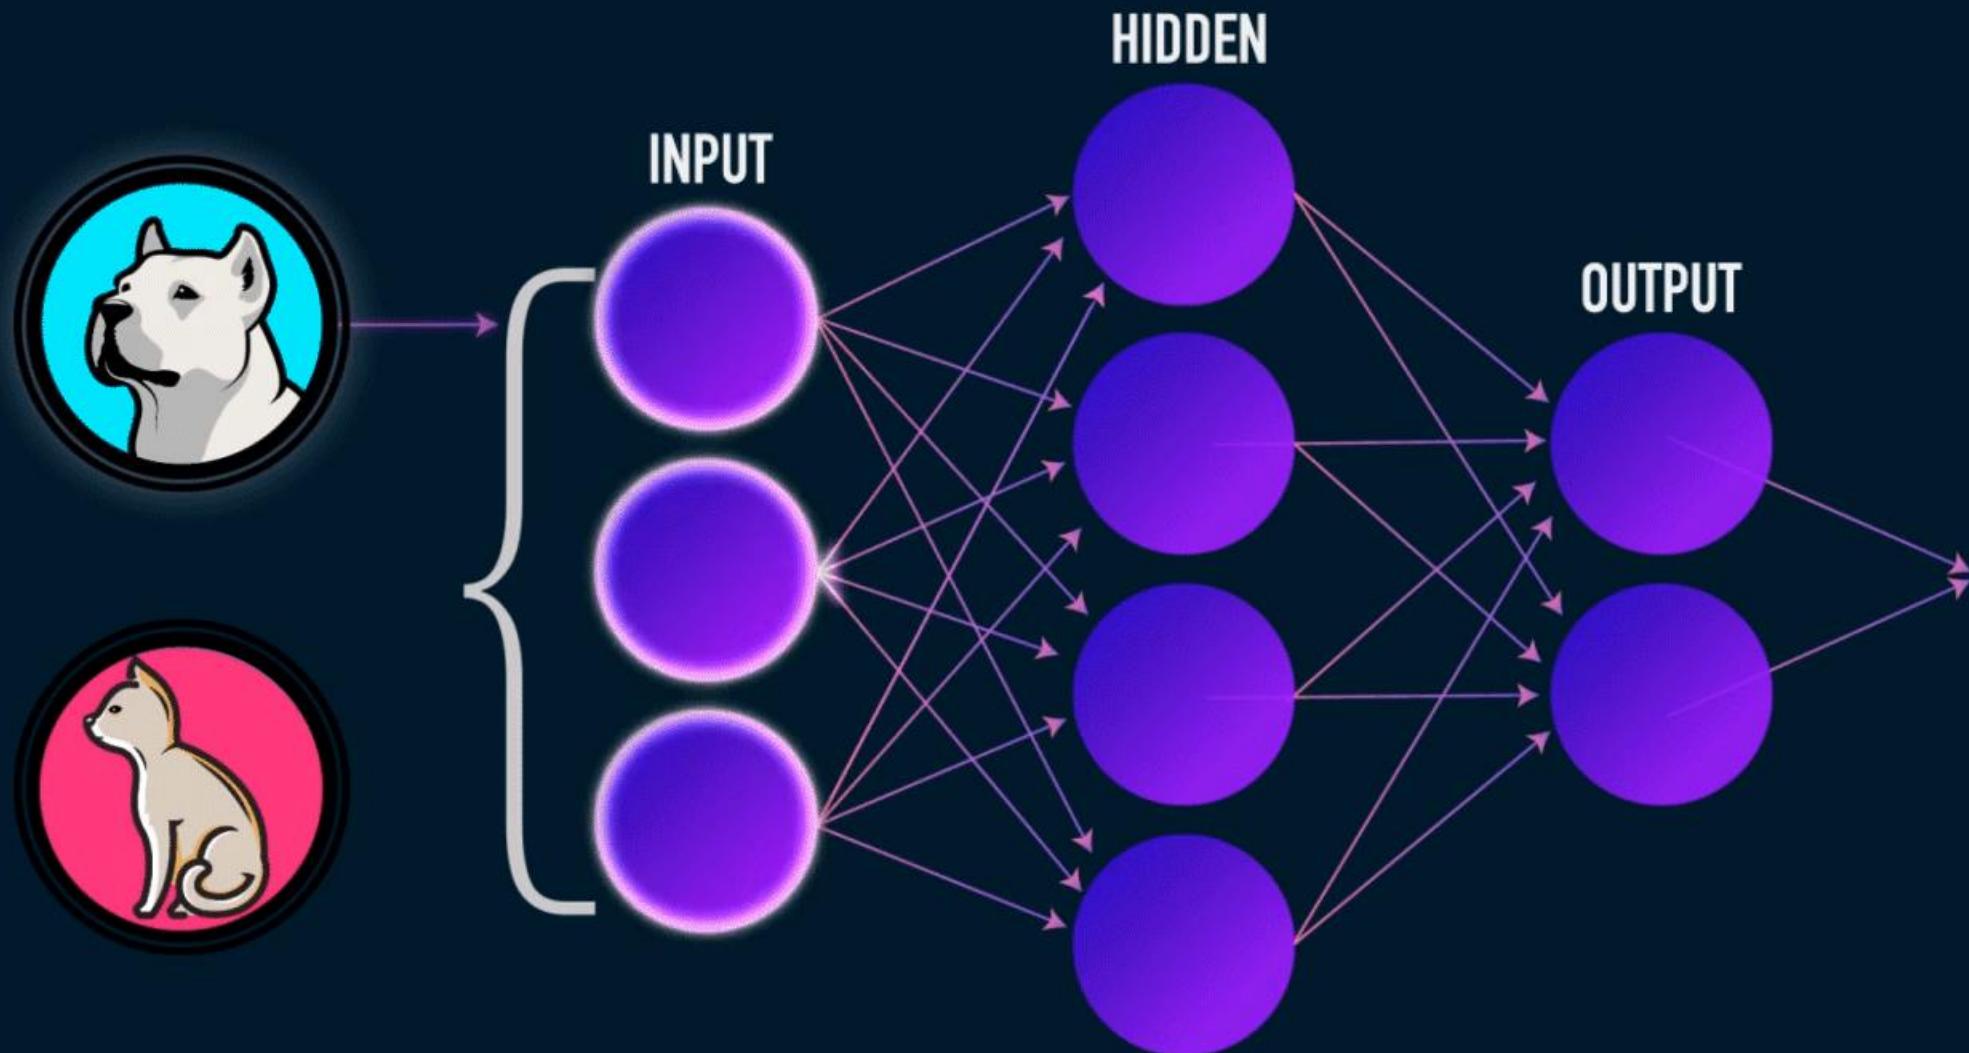


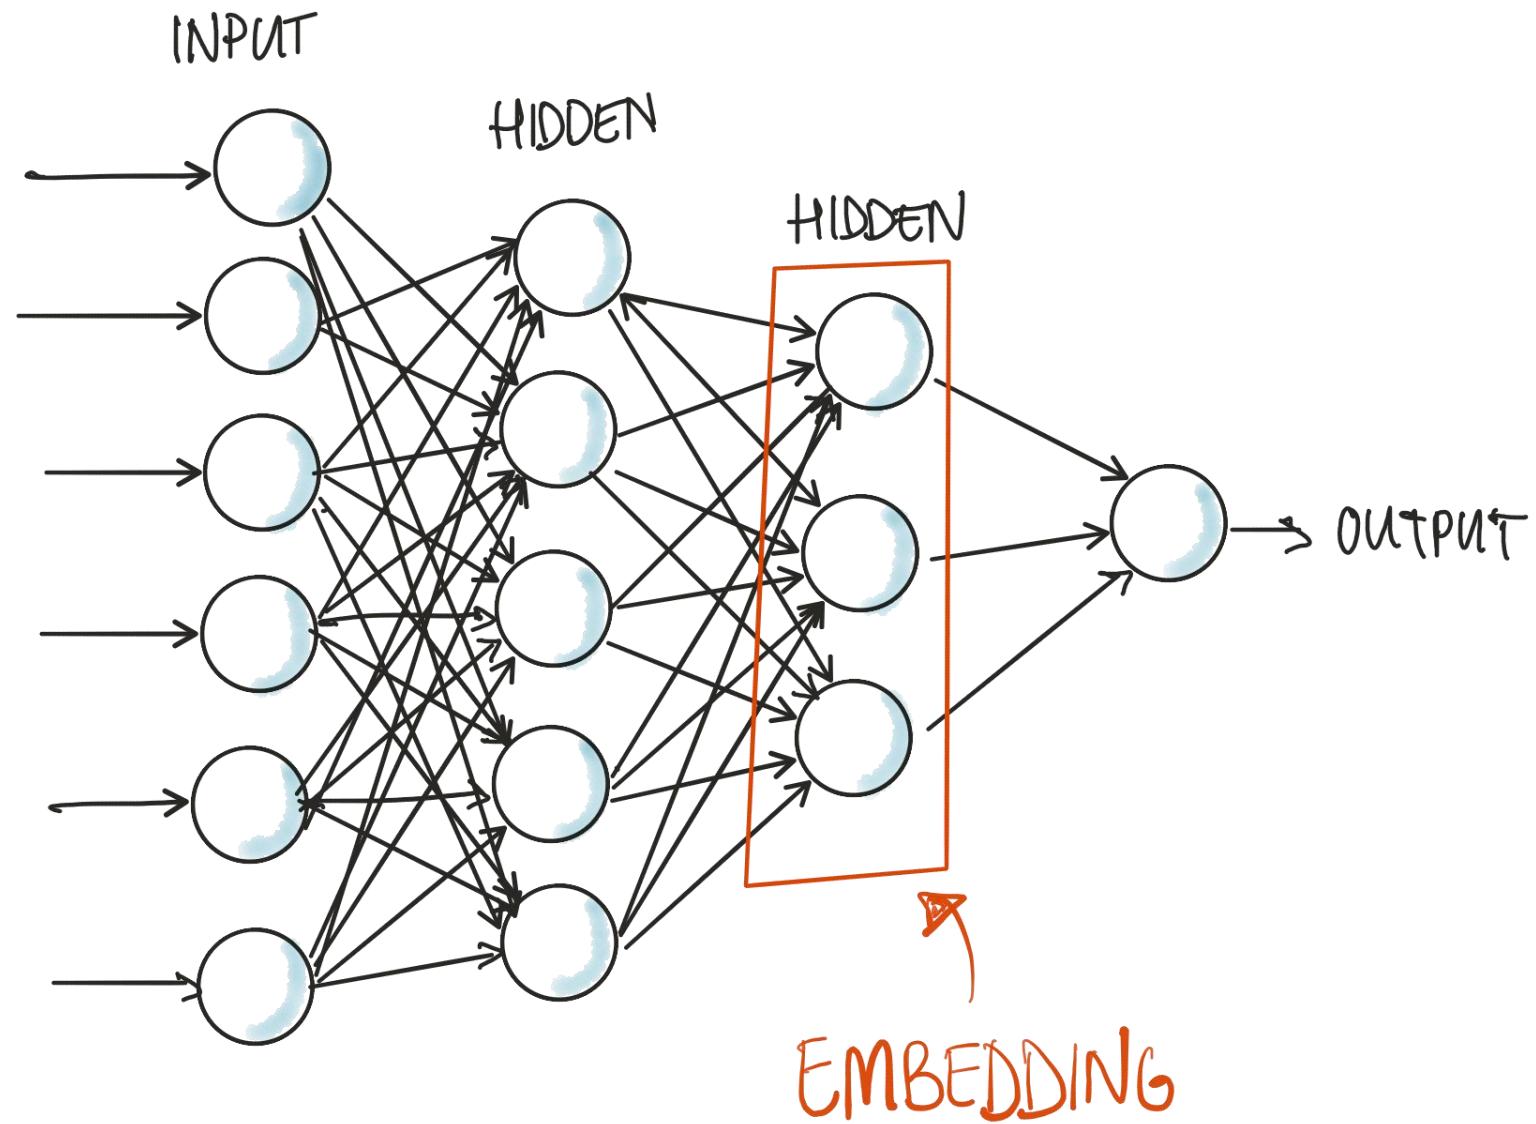
[NEURAL NET]

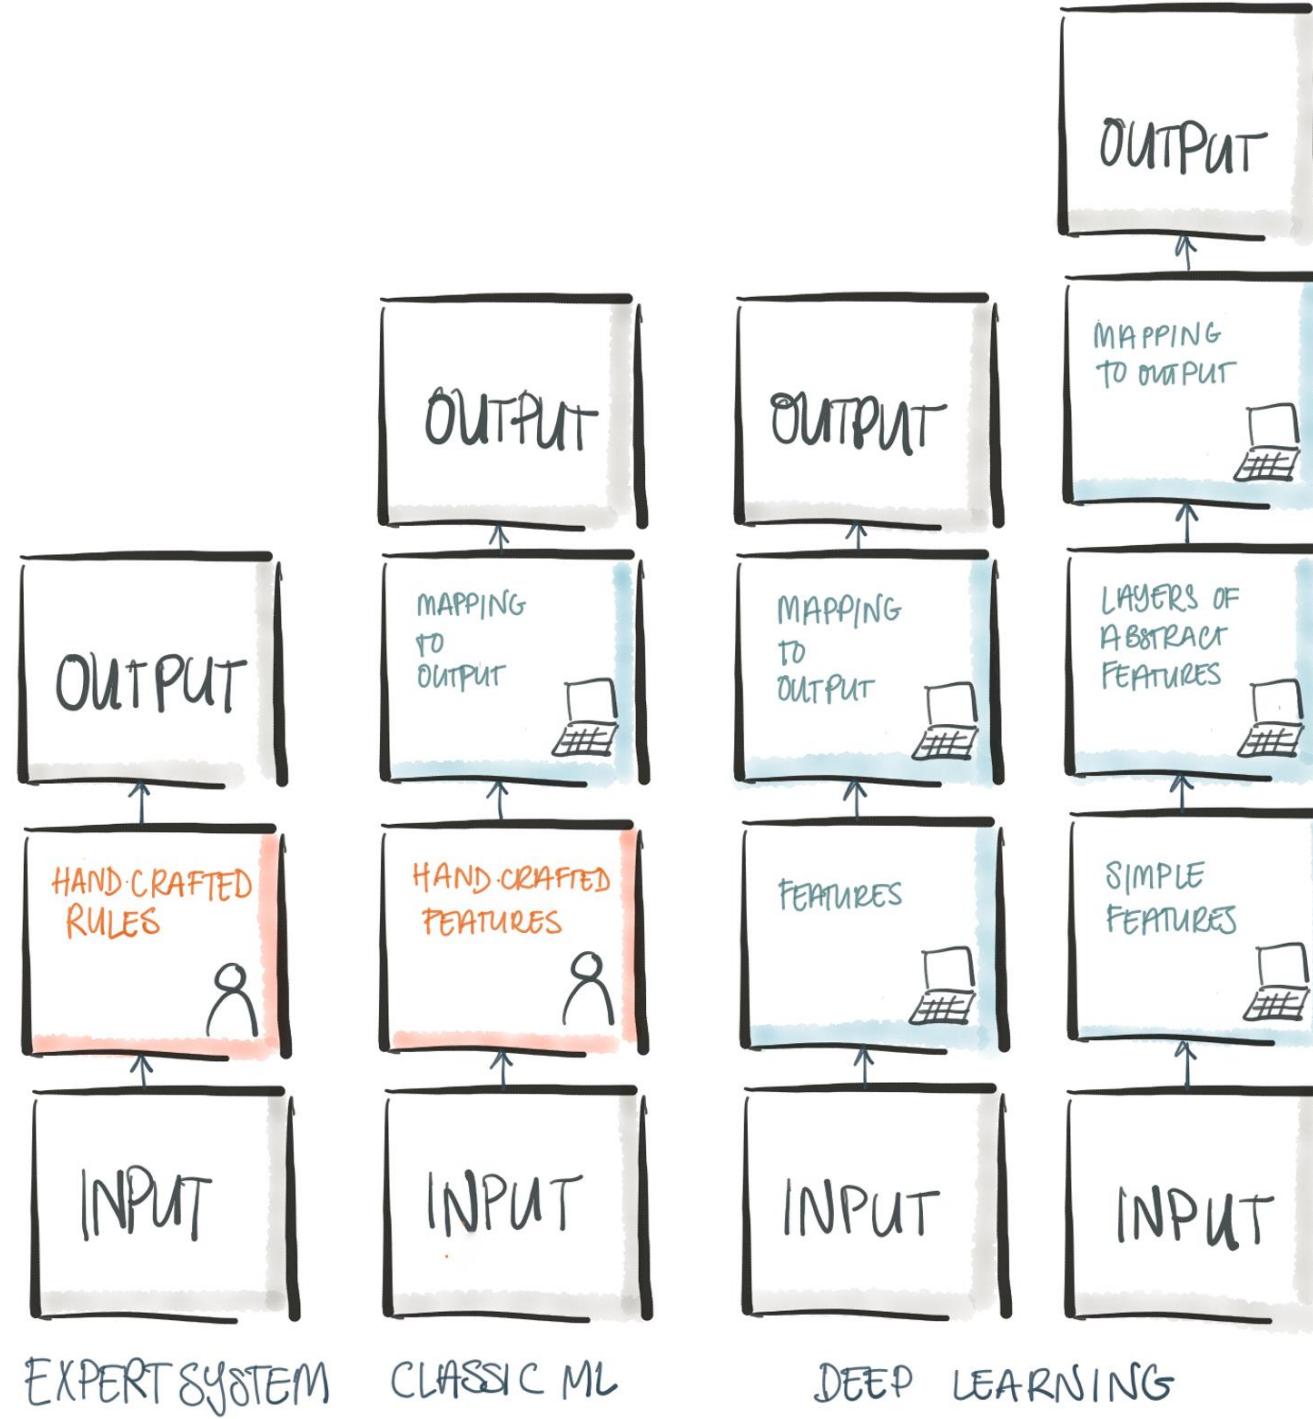
HOUSE PRICES





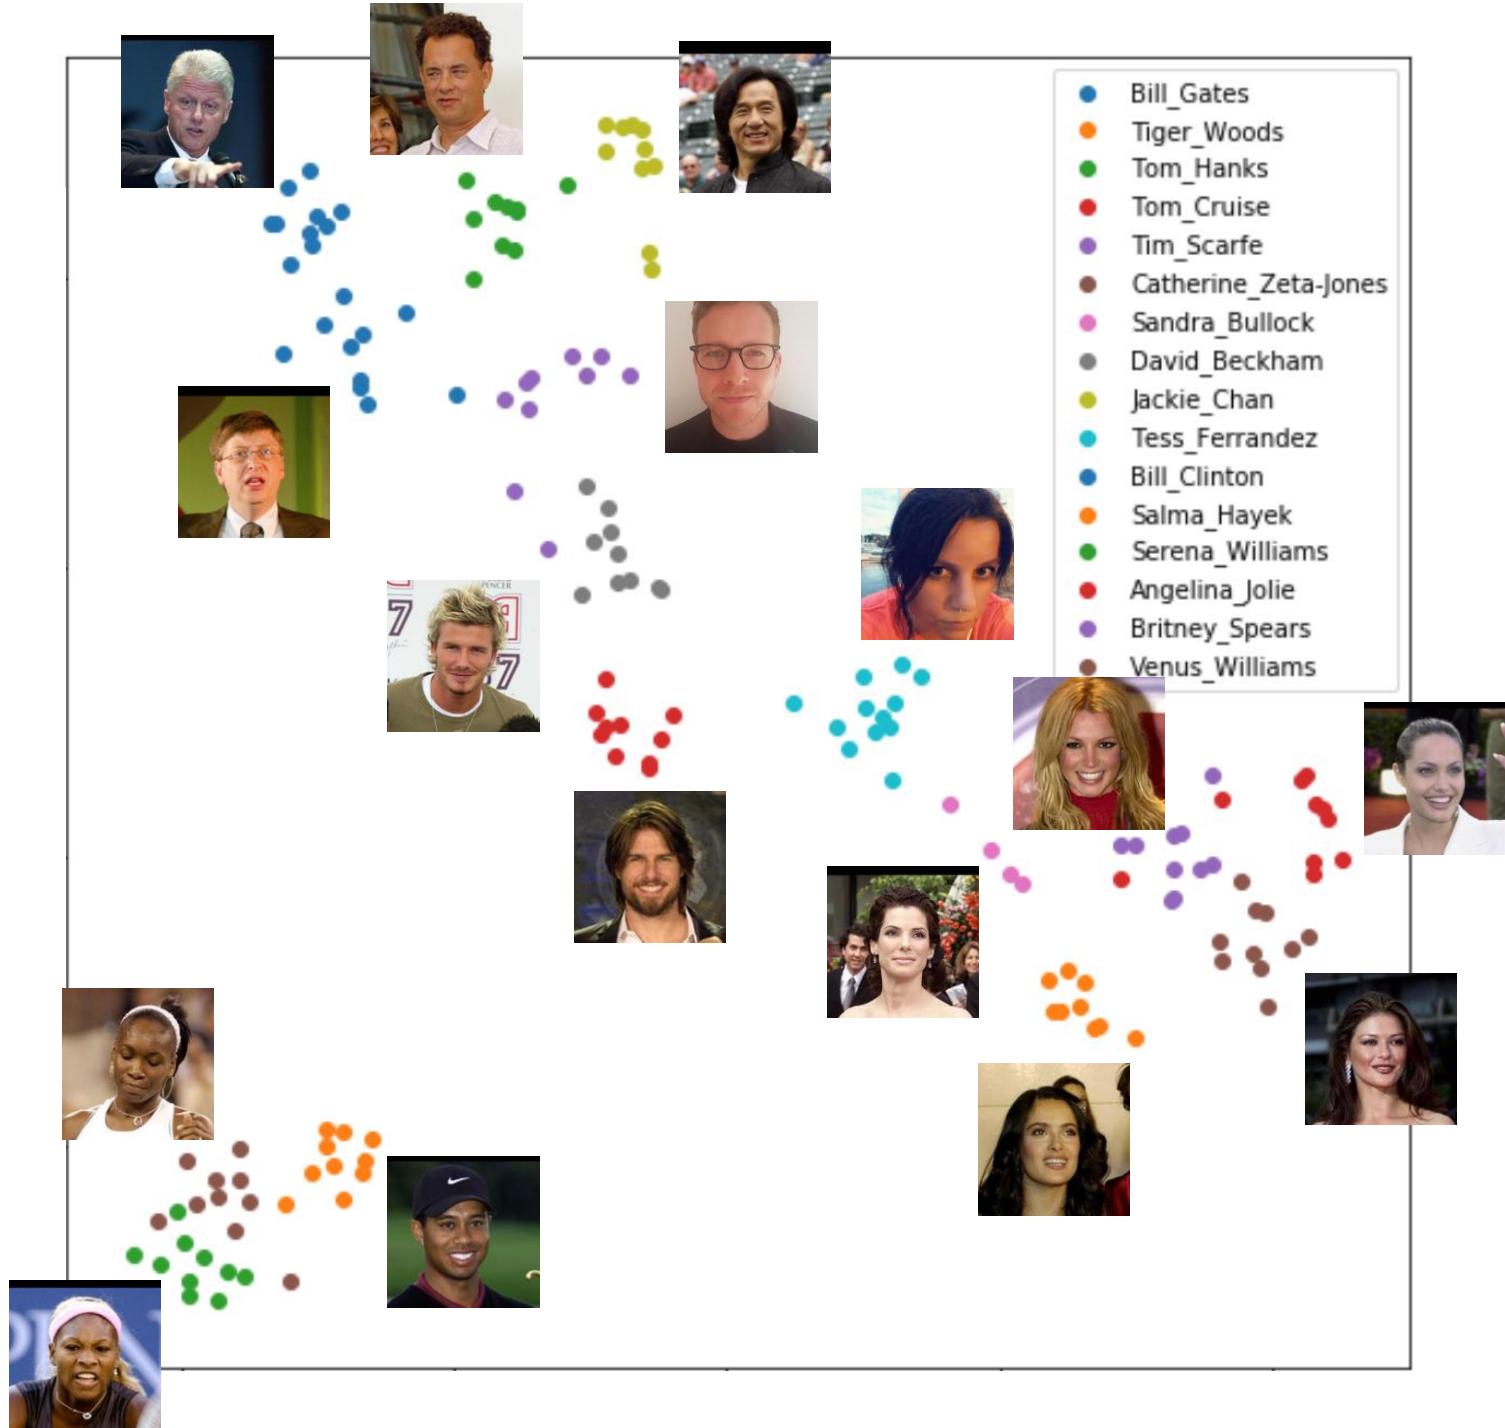






[FACENET]

T-SNE
Projection of
128D to 2D



Word2Vec: Word vectors are cool!



<https://projector.tensorflow.org>

Word2Vec: First names





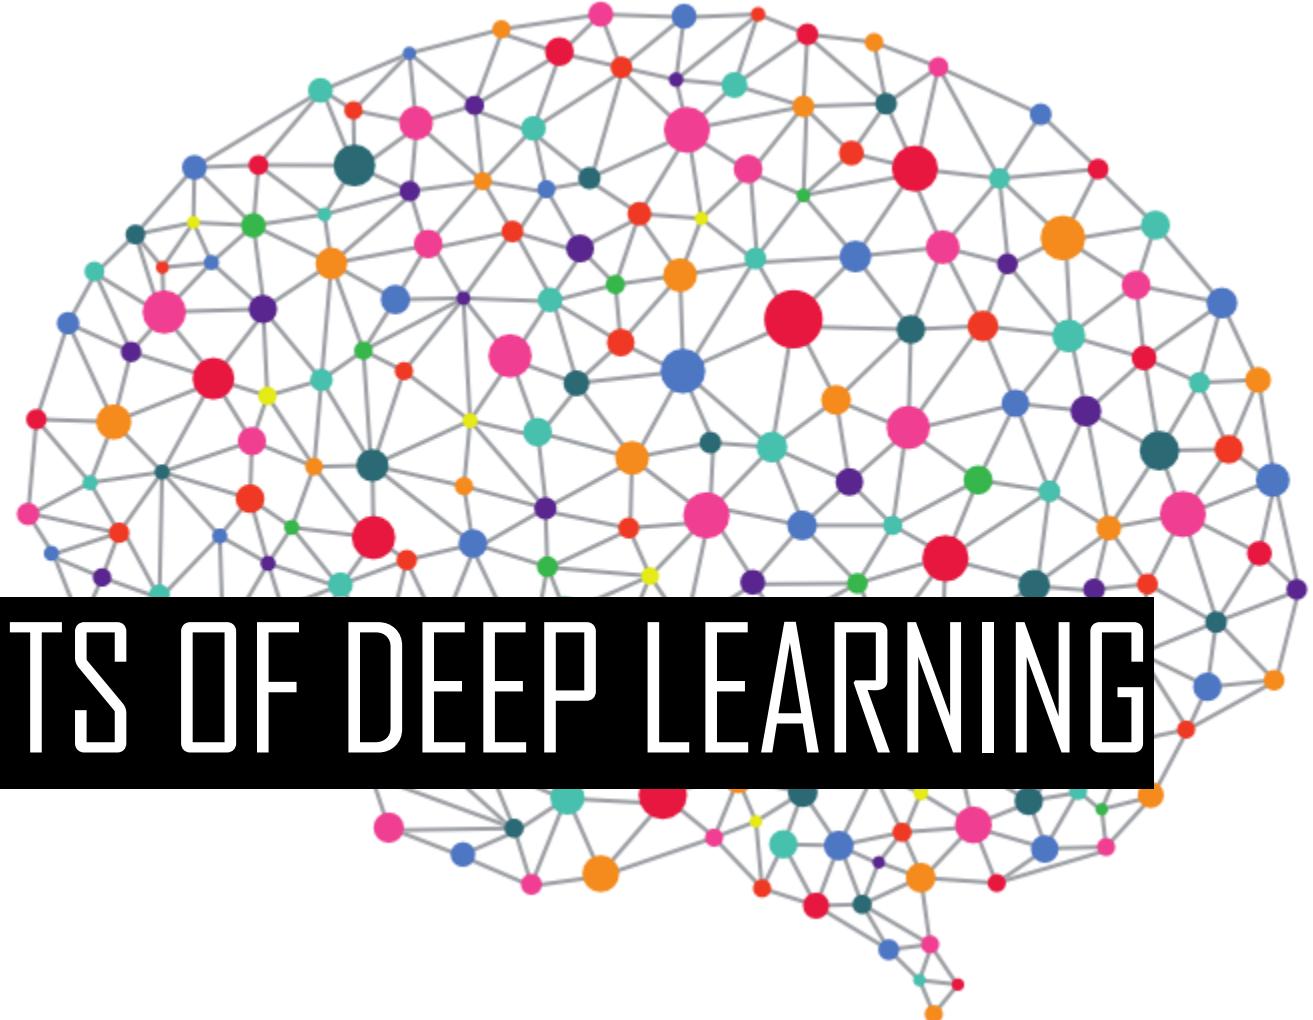
A HOUSE WITH
A NICE FENCE AND
A BEAUTIFUL
DOOR

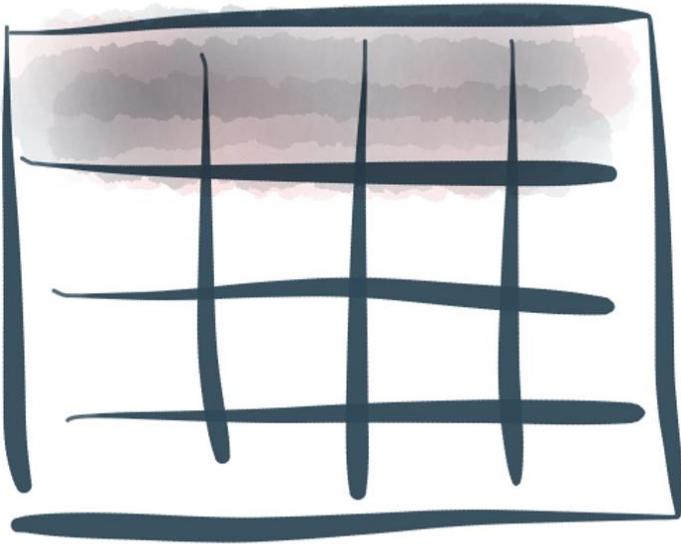


A Style-Based Generator Architecture for Generative
Adversarial Networks
Tero Karras, Samuli Laine, Timo Aila



DISTILLED CONCEPTS OF DEEP LEARNING





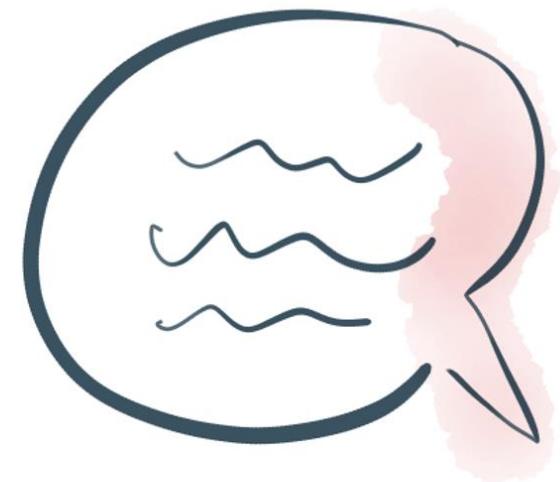
Dense
Neural Network



Convolutional
Neural Network

[SPACE]

[SEQUENCE]



Recurrent
Neural Network



Traditional ML

Manual Features



Trainable Classifier



Deep Learning

Representations are hierarchical and trained automatically

Low Level Features



Mid Level Features

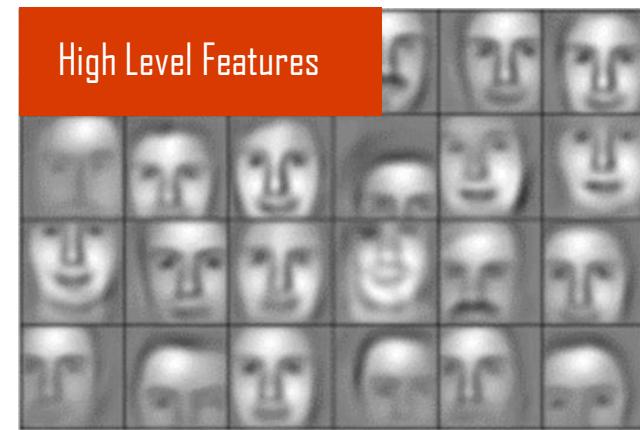
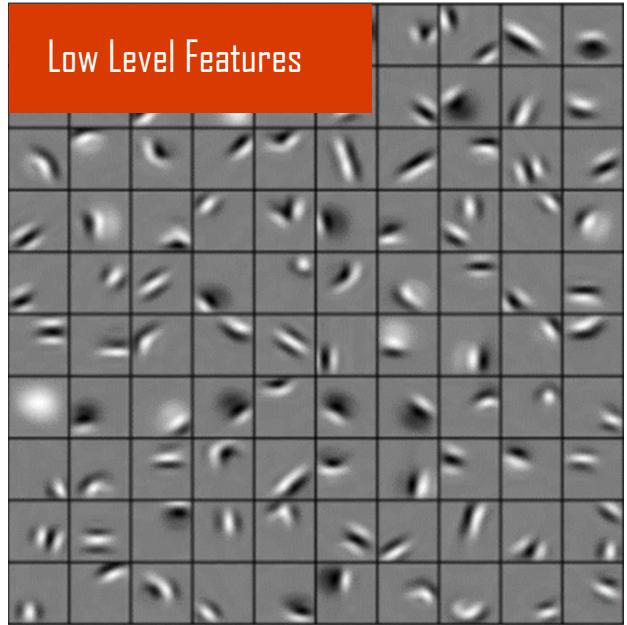


High Level Features



Trainable Classifier

Deep Learning = learning a representation



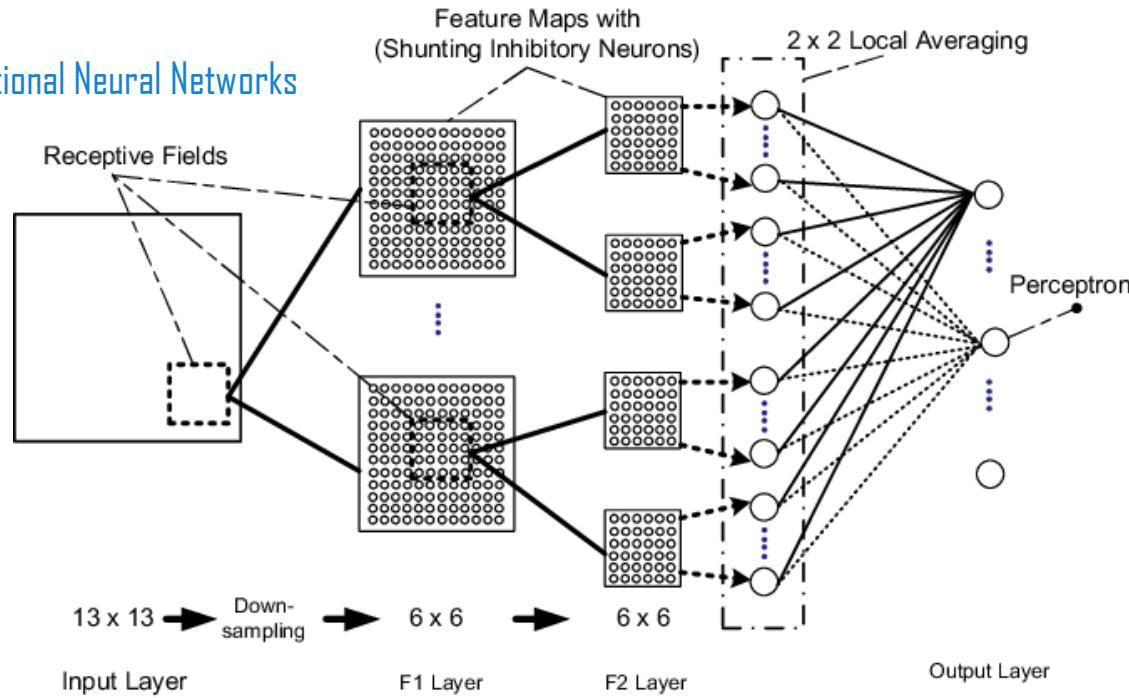
Intermediate representations are learned automatically



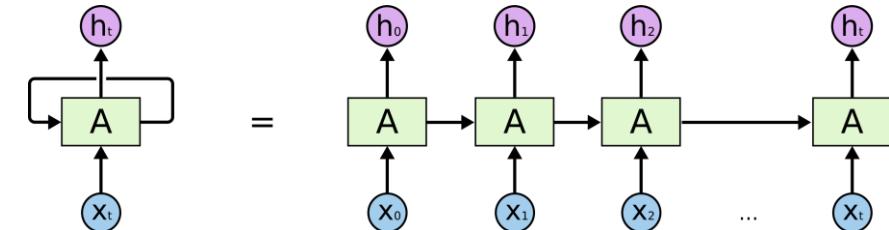
Unlike other algorithms, NNs can natively encode useful and obvious relationships in the data domain

- Local spatial dependencies (vision) i.e. CNNs
- Time dependencies (language, speech) i.e. RNNs

Convolutional Neural Networks



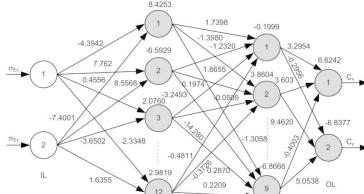
Recurrent Neural Networks



NATIVE DATA-DOMAIN FEATURES

Universal functional approximators

STRUCTURED
SPEECH
LANGUAGE
IMAGES
VIDEO

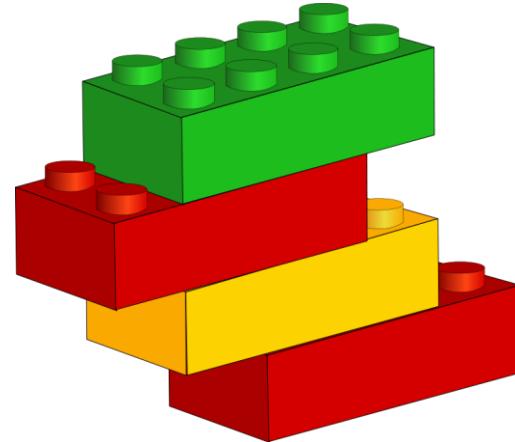


STRUCTURED
SPEECH
LANGUAGE
IMAGES
VIDEO

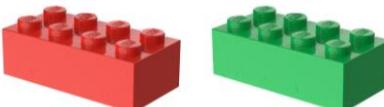
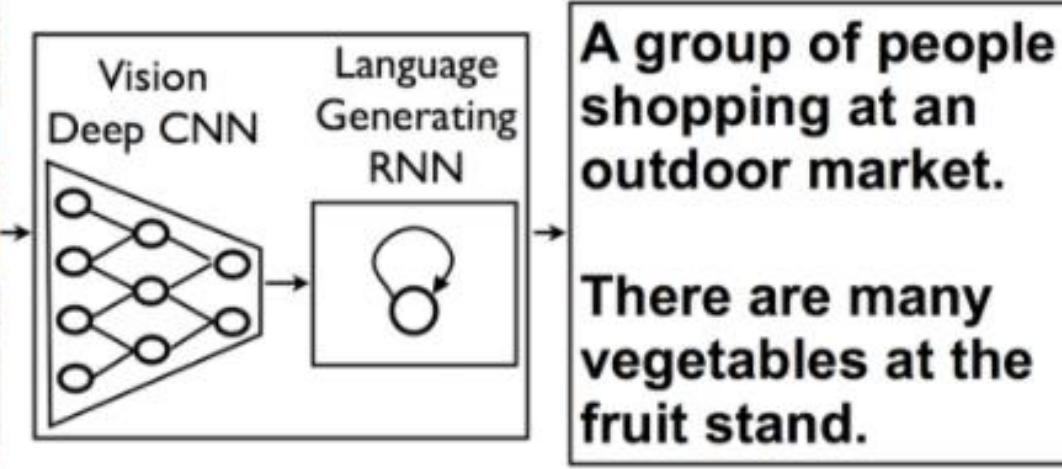
- Map between domains
- Temporal & spatial modelling
- Generative models possible!
- Multiple inputs and outputs!

Composability

- Machine learning is becoming a form of software development
- Machine learning models are like software
- There is a dichotomy between the software engineering process and the data science process



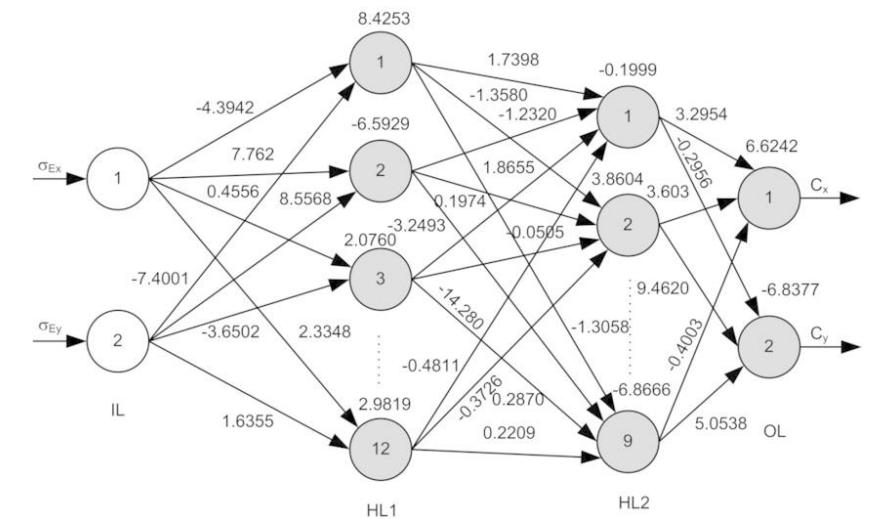
Building predictive architectures like lego blocks



(Vinyals, et al. (2014))

We call this “Software 2.0”

- Software 2.0 is a confection of data generated code and handwritten code
- Some features are **impossible** to write with traditional code i.e. speech synthesis, machine translation, computer vision use cases, game AI
- Computationally homogeneous
- Simple to bake into silicon.
- It is highly portable.
- Composable
- Better than you



When to use deep learning vs classical machine learning

- Large volumes of data
- Unstructured data i.e. images, video, sound
- Modelling temporal or spatial dependencies
- Novel prediction architectures
- Composability, reusability, transferability of models
- Consistency of approach is more conducive to software engineering

When not to use deep learning

- When interpretability is important
- When the data is very small deep learning behaves like classical machine learning i.e. no representation learning
- When you need statistical guarantees
- Structured Data (sometimes)

Start with logistic regression



Epoch
003,928

Learning rate
0.03

Activation
ReLU

Regularization
None

Regularization rate
0

Problem type
Classification

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%

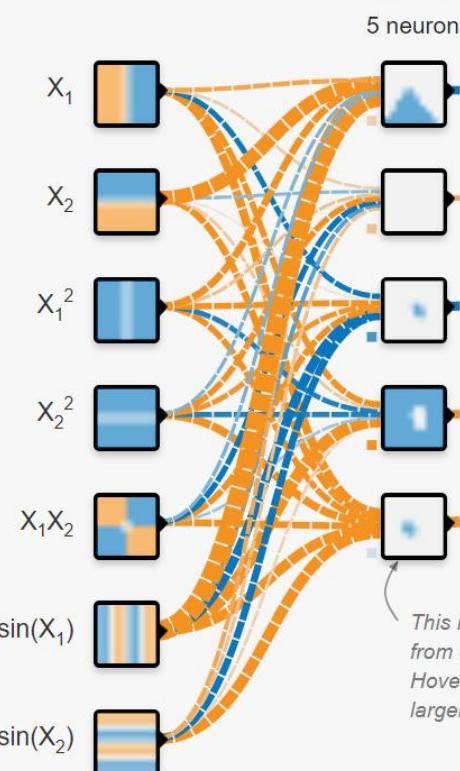
Noise: 0

Batch size: 10

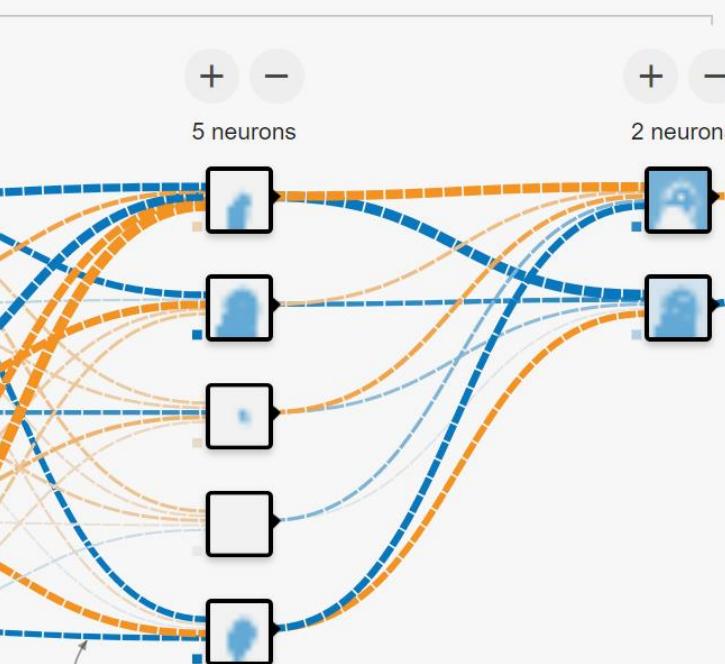
REGENERATE

FEATURES

Which properties do you want to feed in?

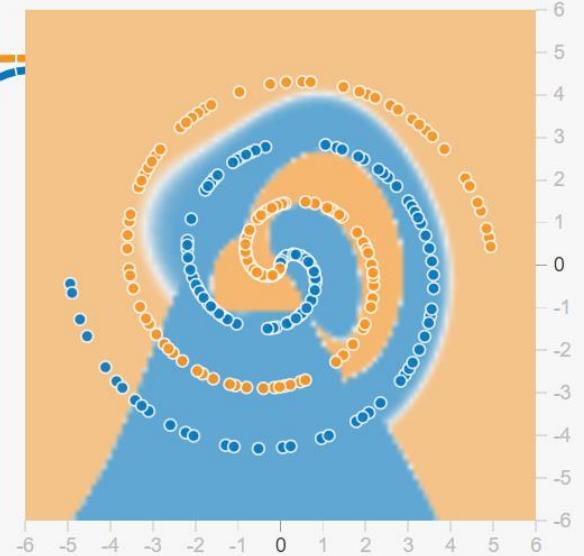


+ - 3 HIDDEN LAYERS



OUTPUT

Test loss 0.211
Training loss 0.185

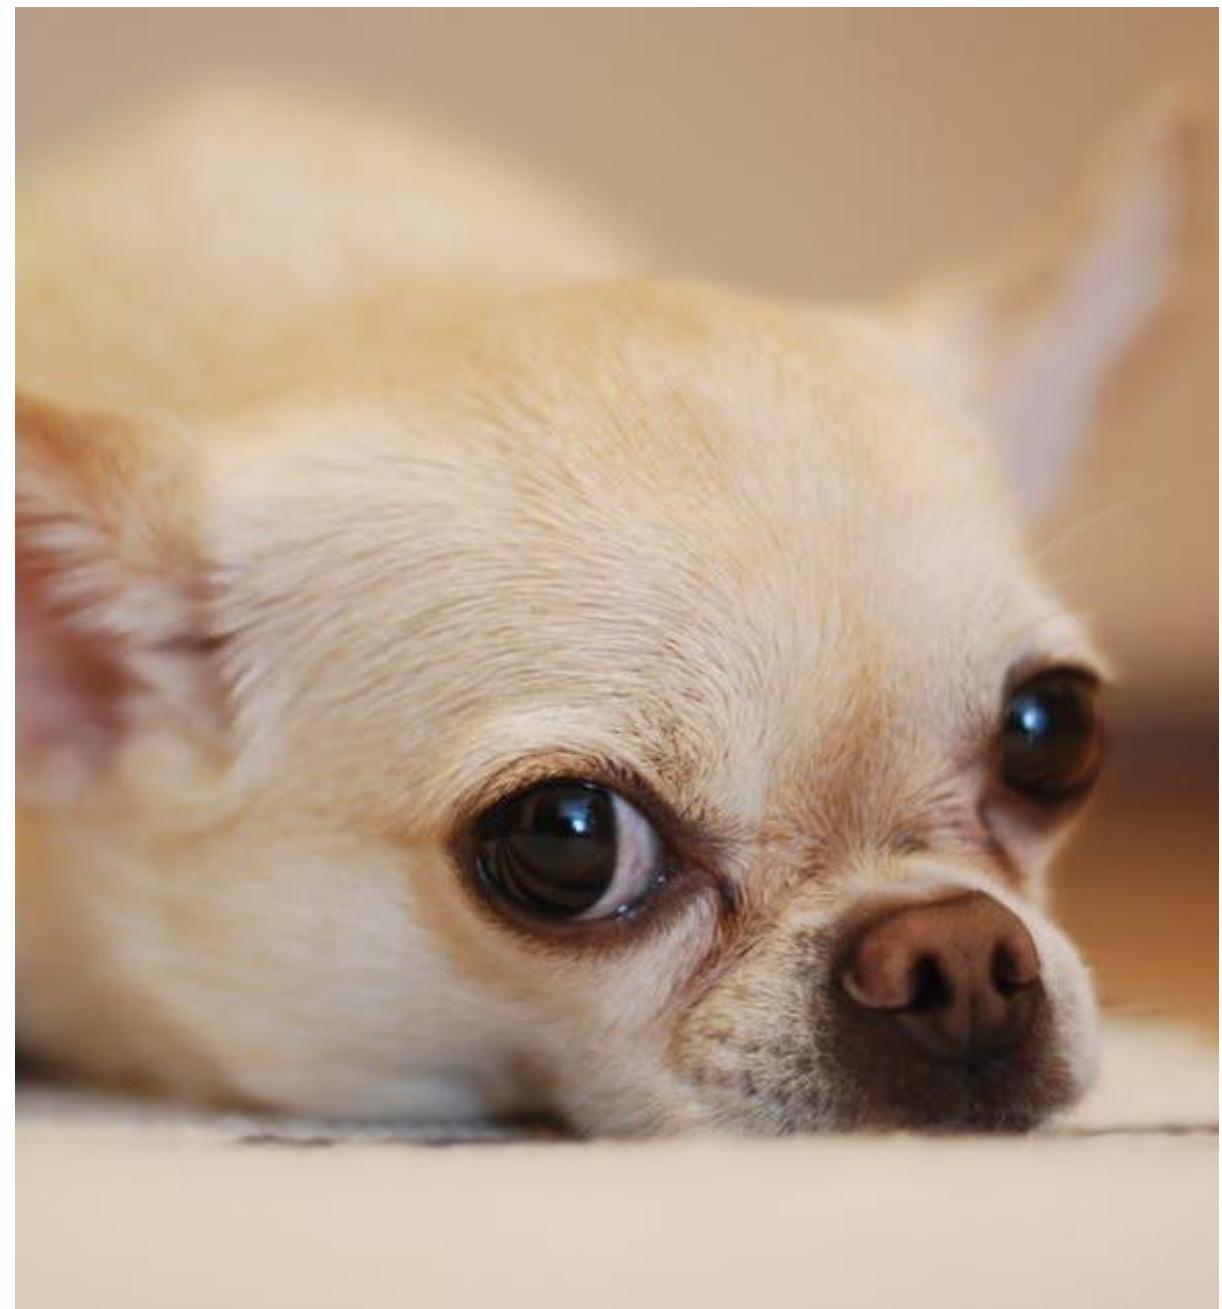


Colors show data, neuron and weight values.



Show test data

Discretize output



```
IF PICTURE.CONTAINS(2 EYES & NOSE)
    RETURN CHIUAWA
ELSE IF PICTURE.CONTAINS(PAPERCUP)
    RETURN MUFFIN
ELSE
    RETURN I HAVE NO CLUE
```



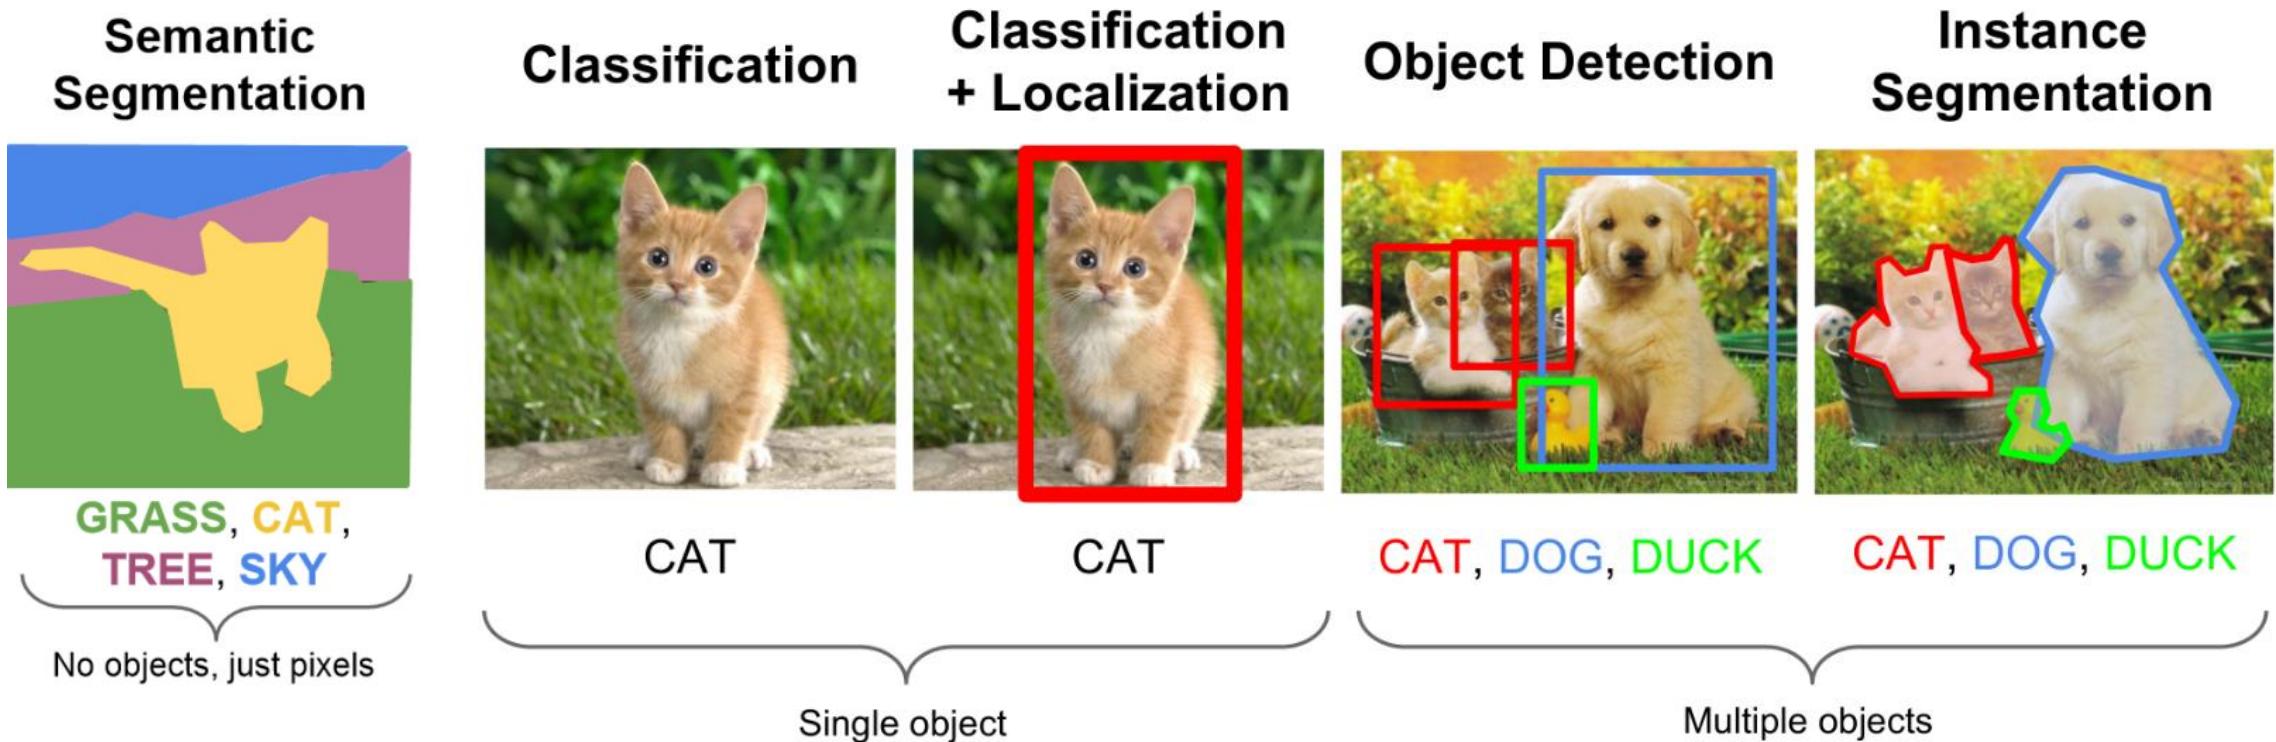
from @teenybiscuit



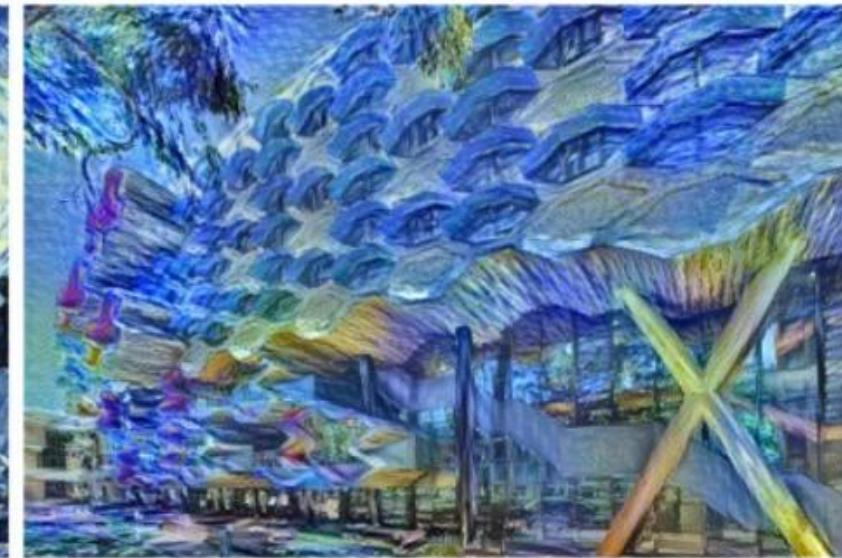
What can you do with computer vision?

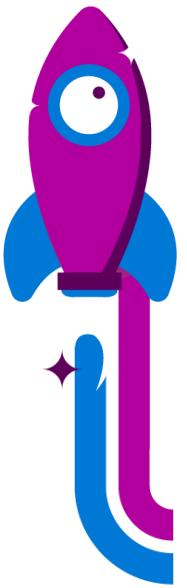


DISCRIMINATIVE ARCHITECTURES



GENERATIVE ARCHITECTURES



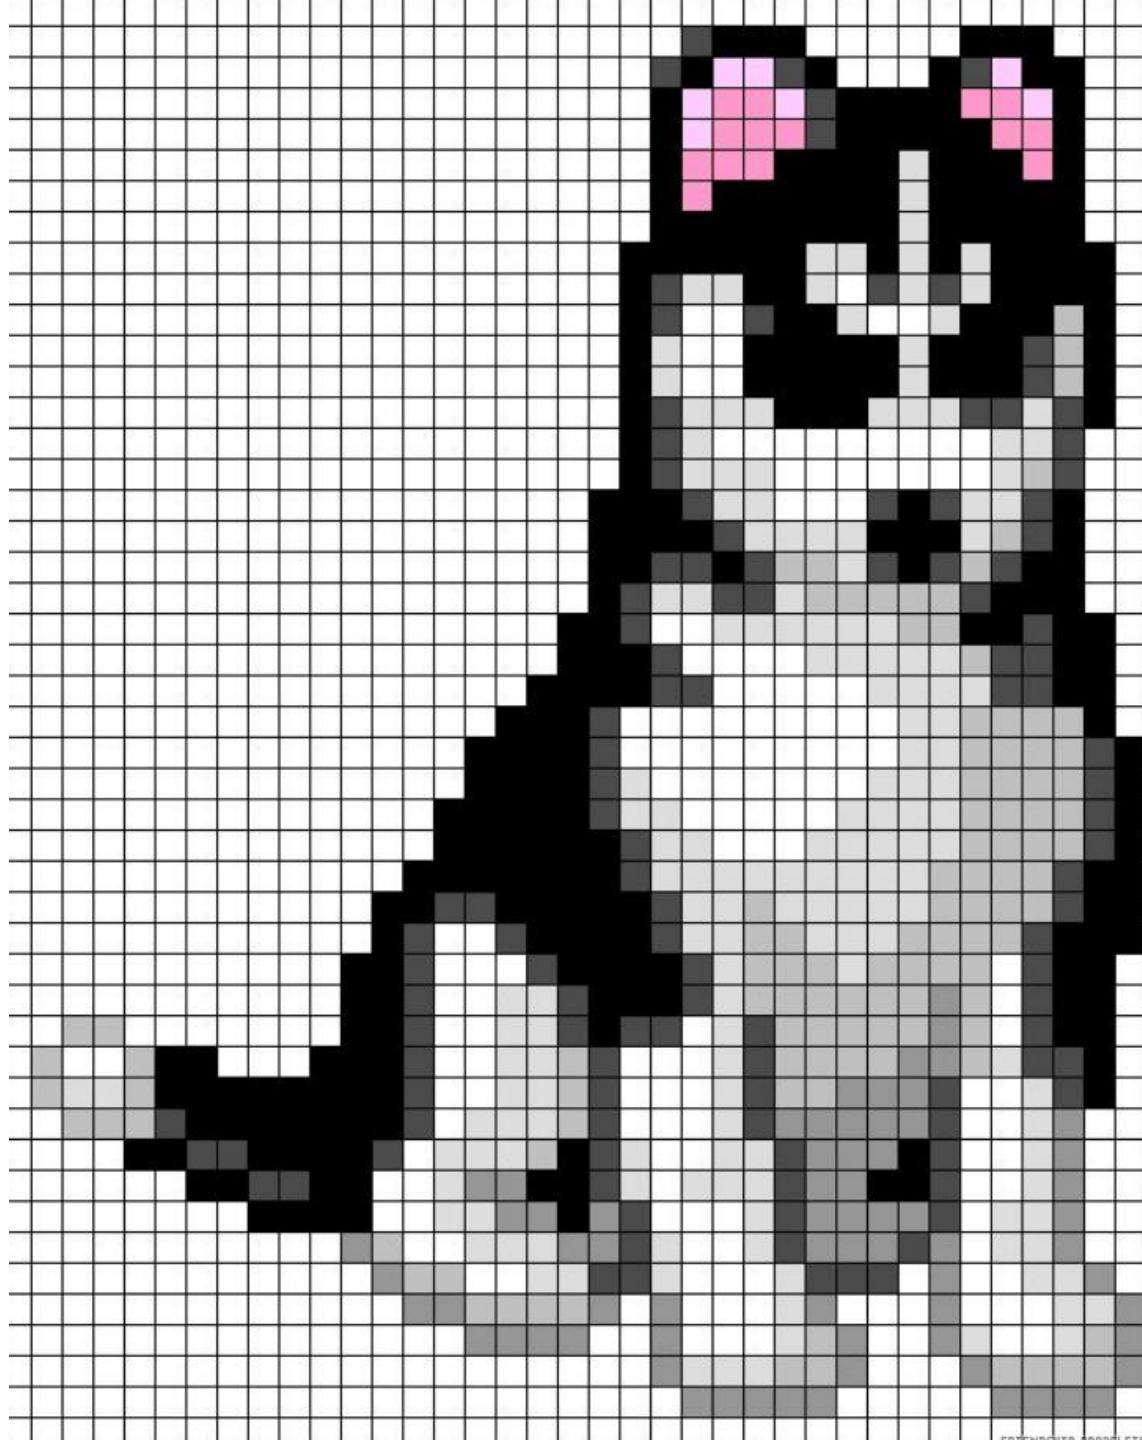
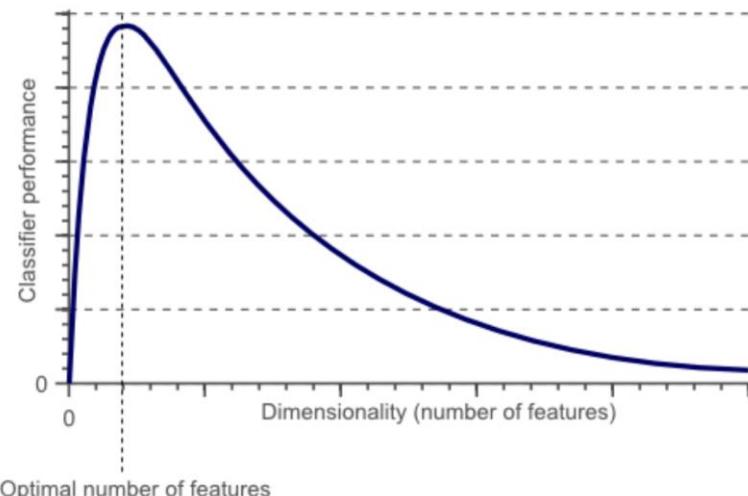


Vision before CNNs?



PIXEL-WISE LEARNING

- Not “translation” invariant
- Not scale invariant
- Not rotational invariant
- Curse of dimensionality



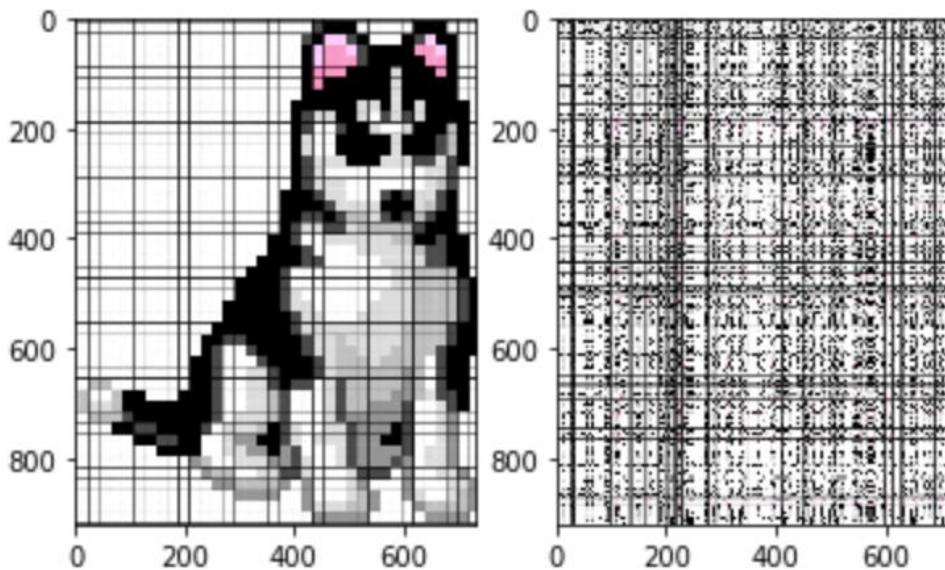
WHAT A CLASSICAL ALGORITHM SEES (I)

```
img=mpimg.imread('dog.png')

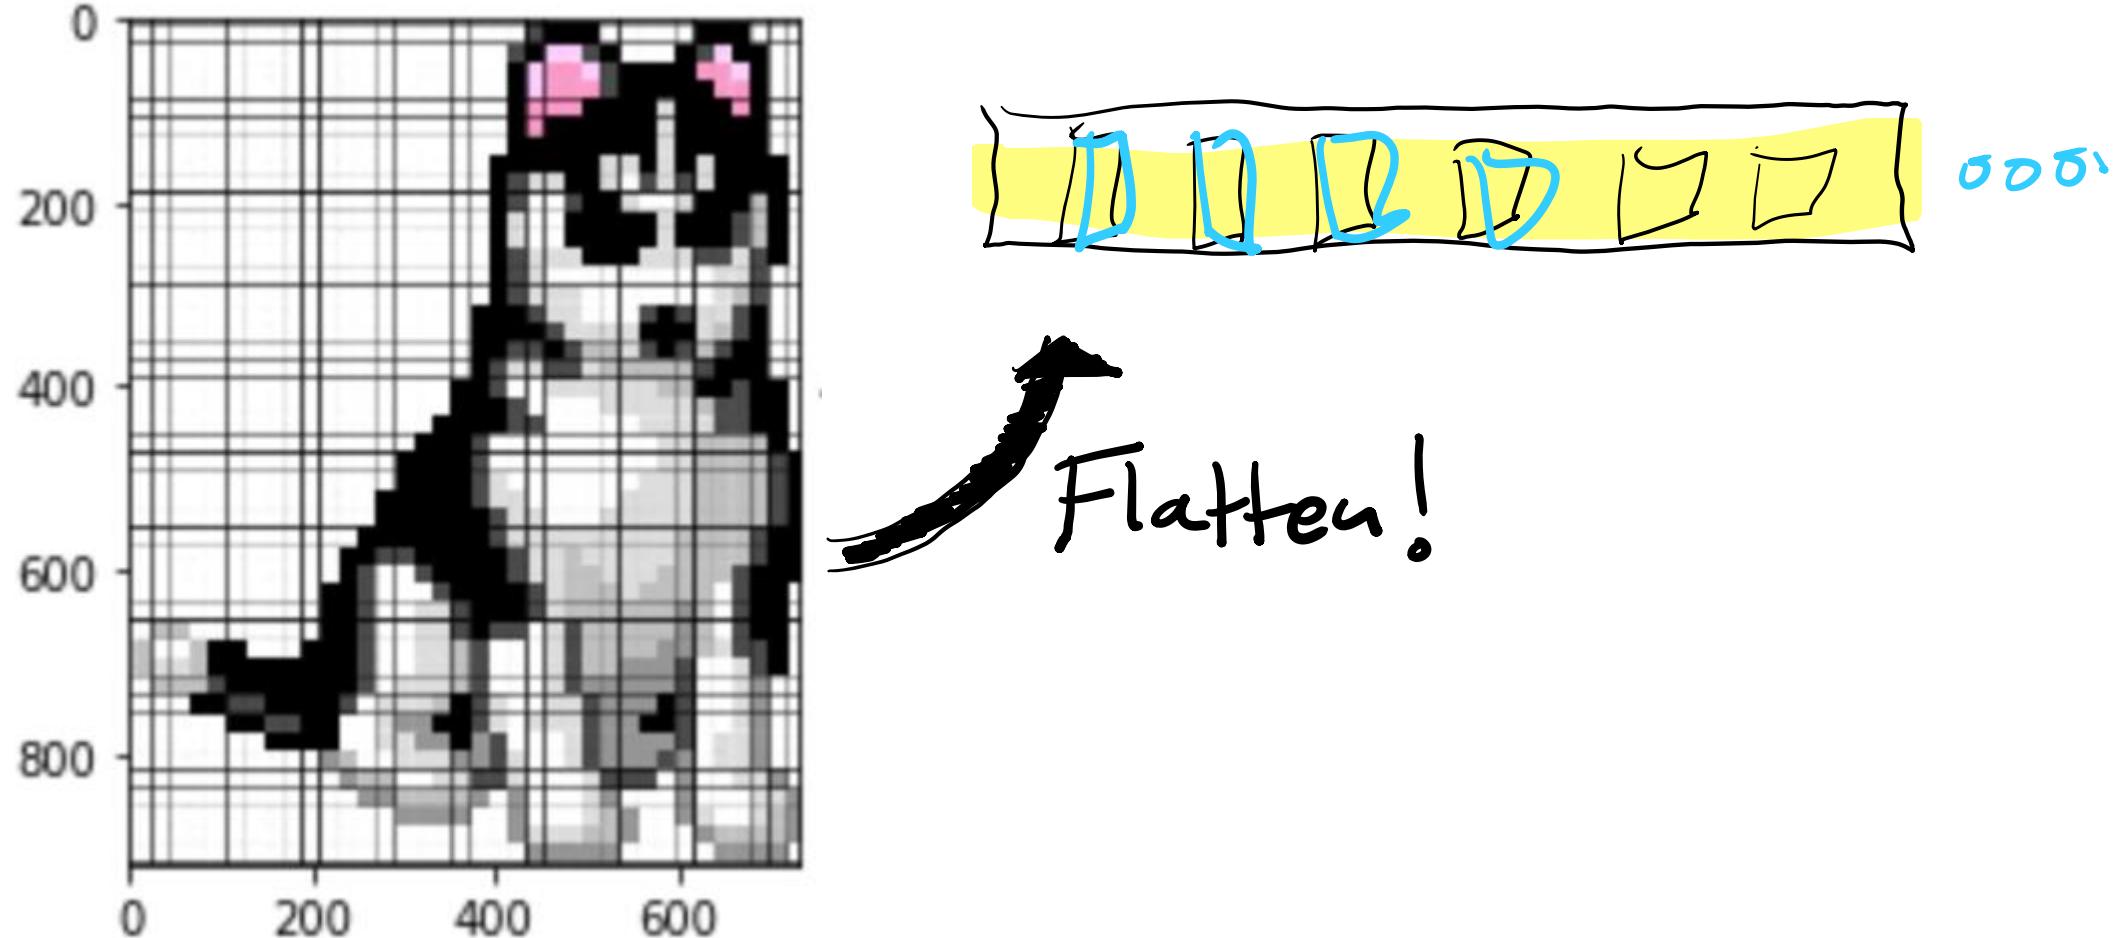
plt.subplot(121)
imgplot = plt.imshow(img)

imgshuff = np.take(img,np.random.permutation(img.shape[1]),axis=1,out=img)
imgshuff = np.take(imgshuff,np.random.permutation(img.shape[0]),axis=0,out=img)
plt.subplot(122)
plt.imshow(imgshuff)
```

Out[32]: <matplotlib.image.AxesImage at 0x25384e0ff98>

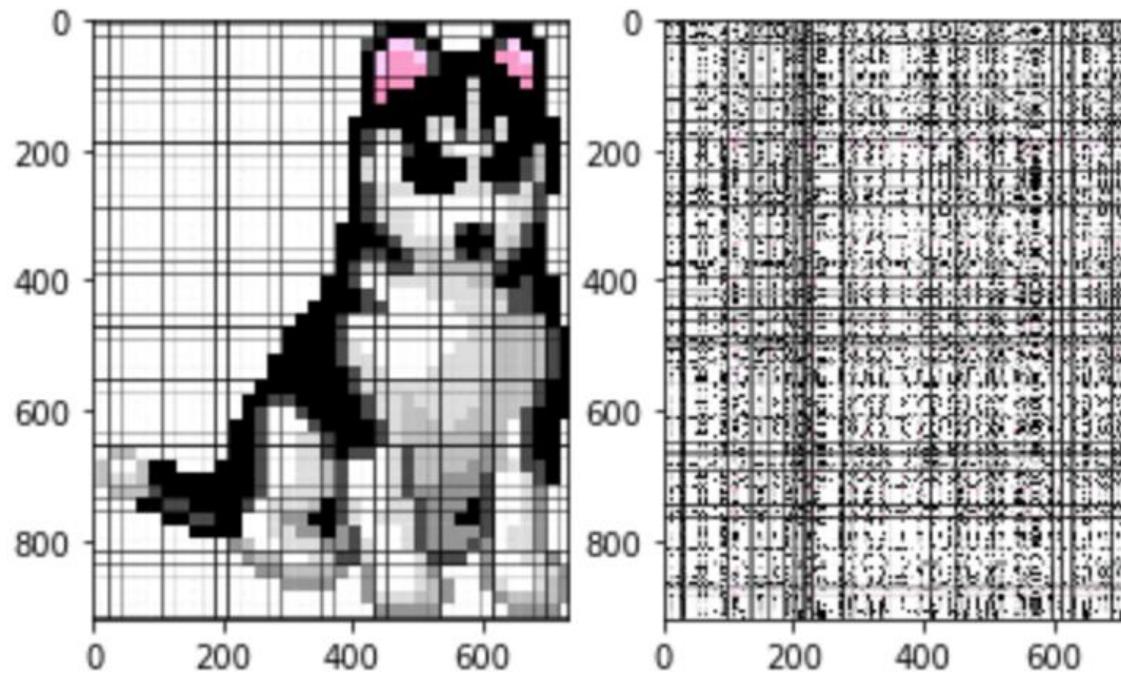


WHAT A CLASSICAL ALGORITHM SEES (2)



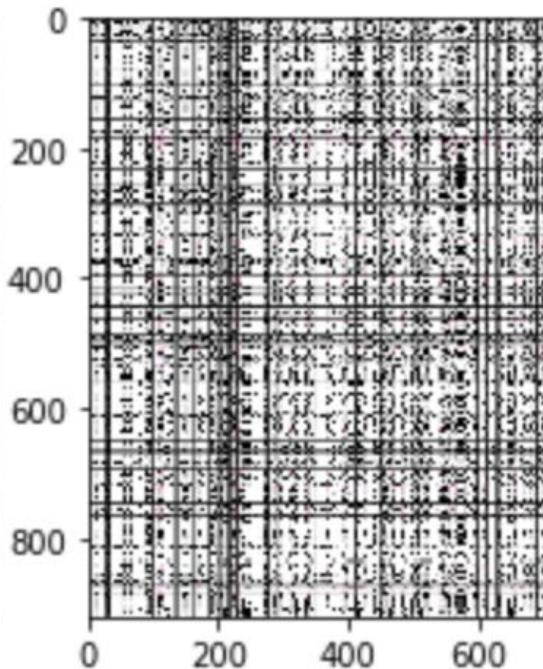
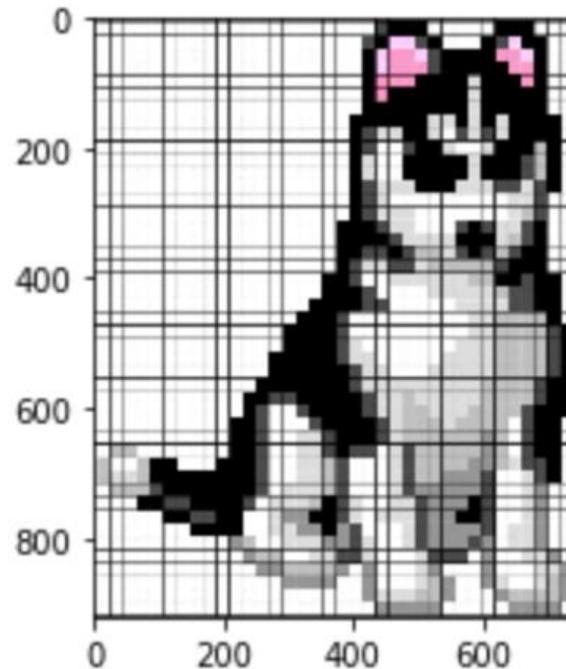
WHAT A CLASSICAL ALGORITHM SEES (3)

- There are two key types of information in an image;
 - the pixel values themselves and ...
 - **their spatial relationships/context**



WHAT A CLASSICAL ALGORITHM SEES (3)

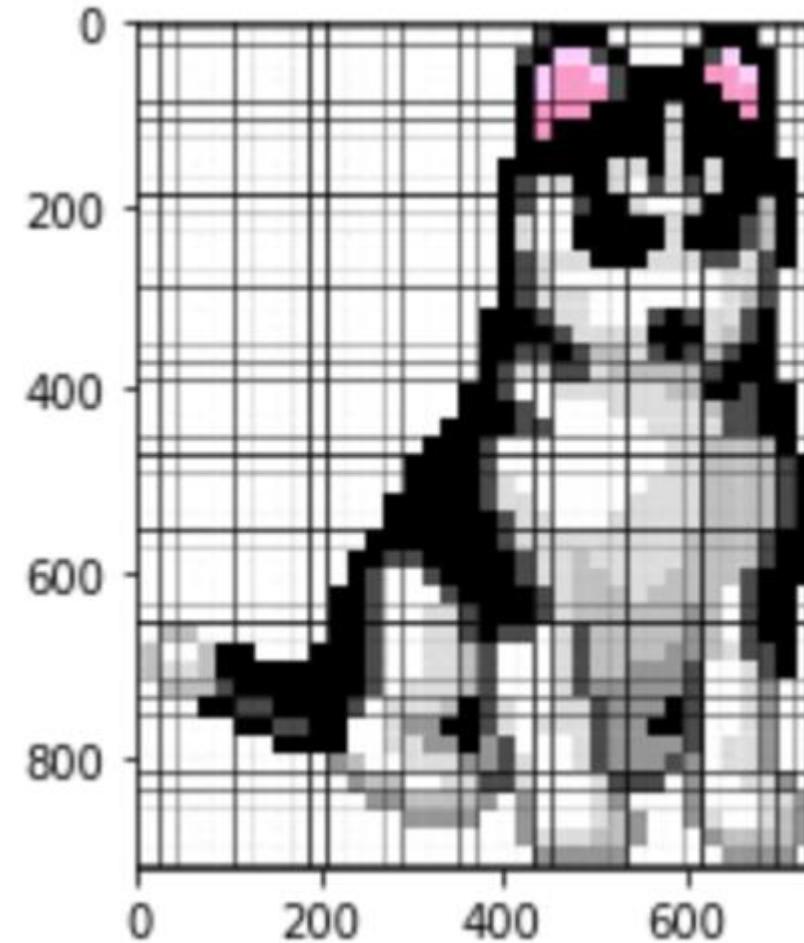
- When you shuffle the pixels of an image, **all the meaning is lost** but the statistical **distribution remains the same**.



IMAGES ARE...

- Statistically undescriptive
- Sparse
- Huge

So we need a way to transform the information **“out of the space domain”**



600*800 ~ .5M dimensions!

PCA

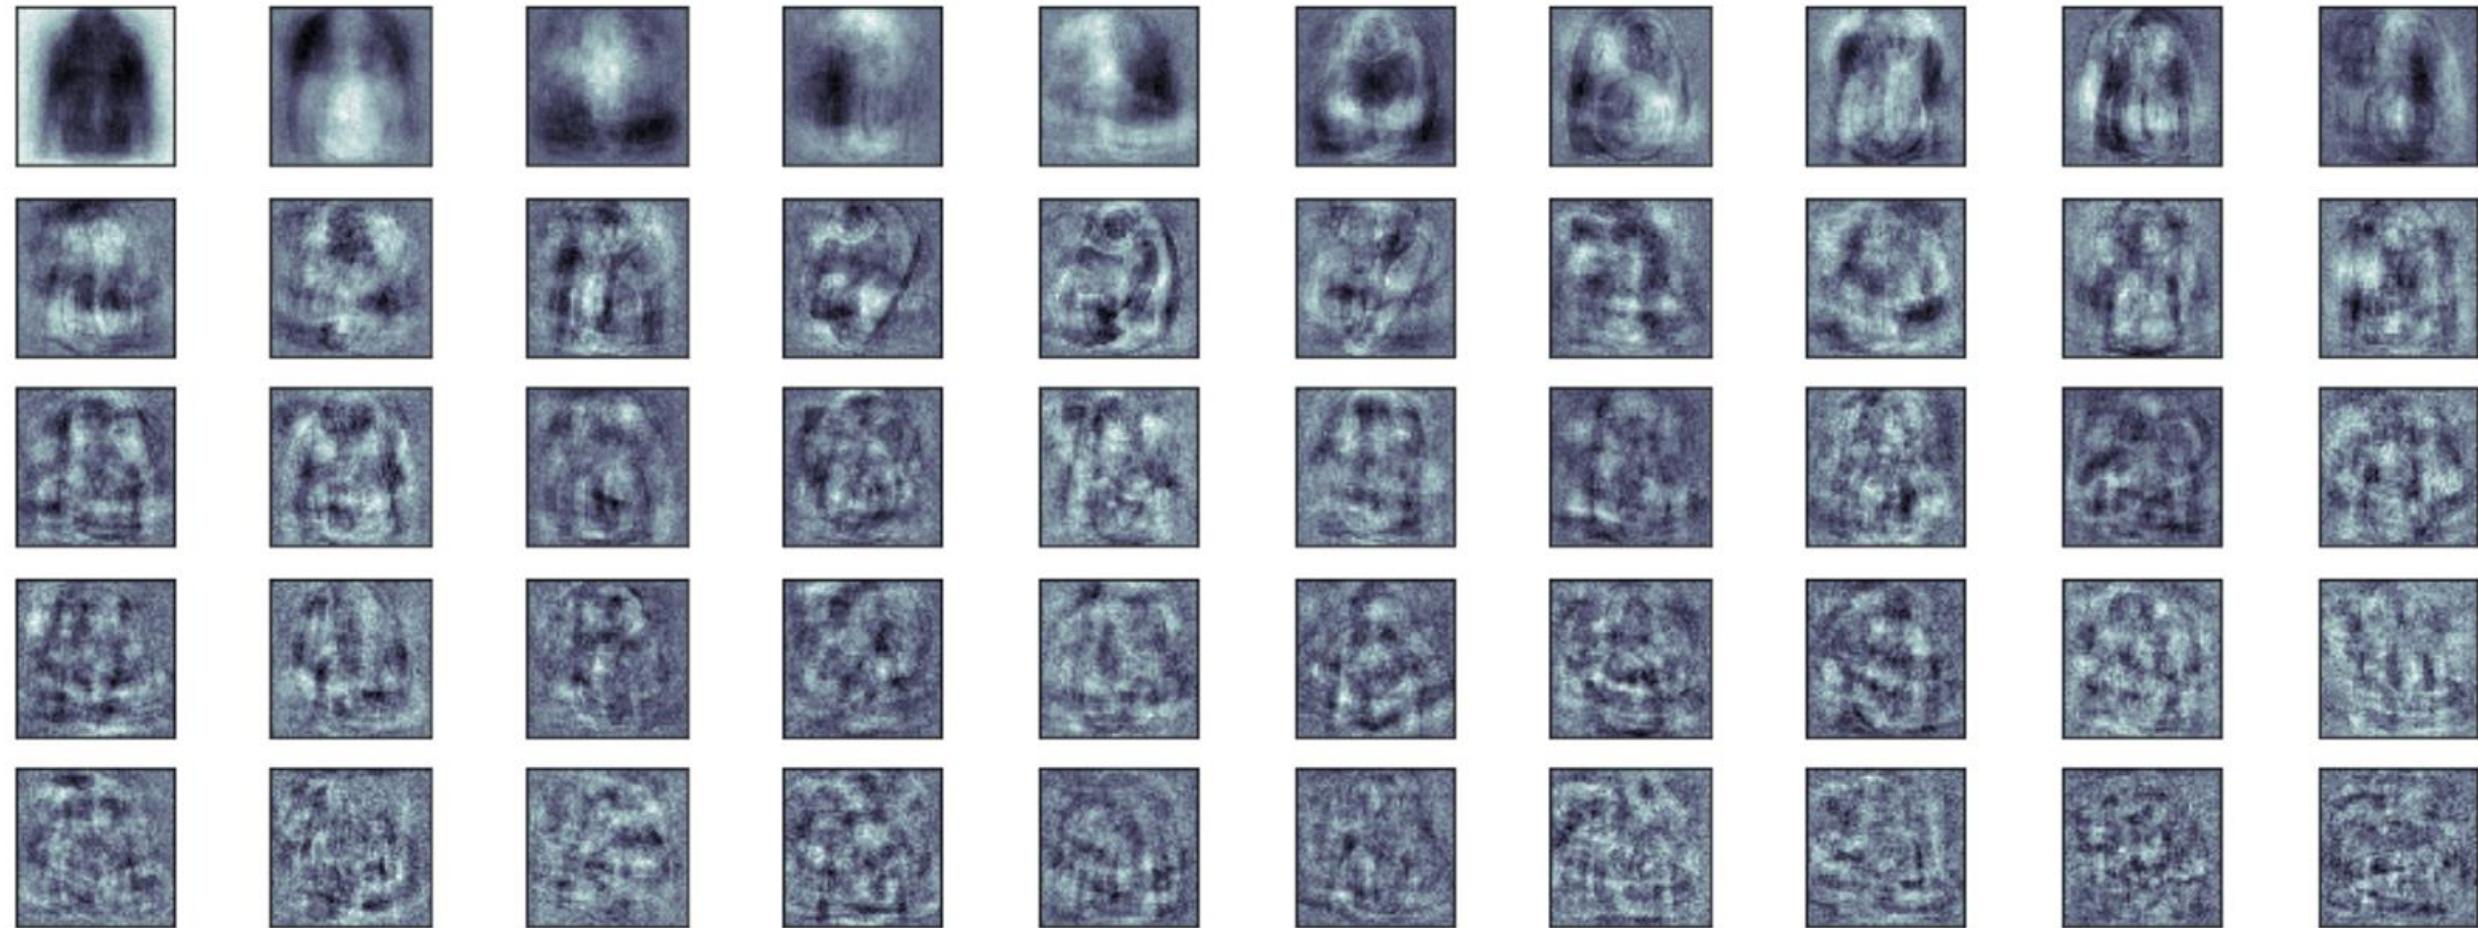
```
In [45]: from sklearn import decomposition

def PCA(xtr, xte, n_components=50, flat_size = 128*128):

    pca = decomposition.PCA(n_components=n_components)
    pca.fit(xtr.reshape(xtr.shape[0],flat_size))

    xtr_pca = pca.transform(xtr.reshape(xtr.shape[0],flat_size))
    xte_pca = pca.transform(xte.reshape(xte.shape[0],flat_size))

    return xtr_pca, xte_pca, pca
```



HISTOGRAMS & FEATURES

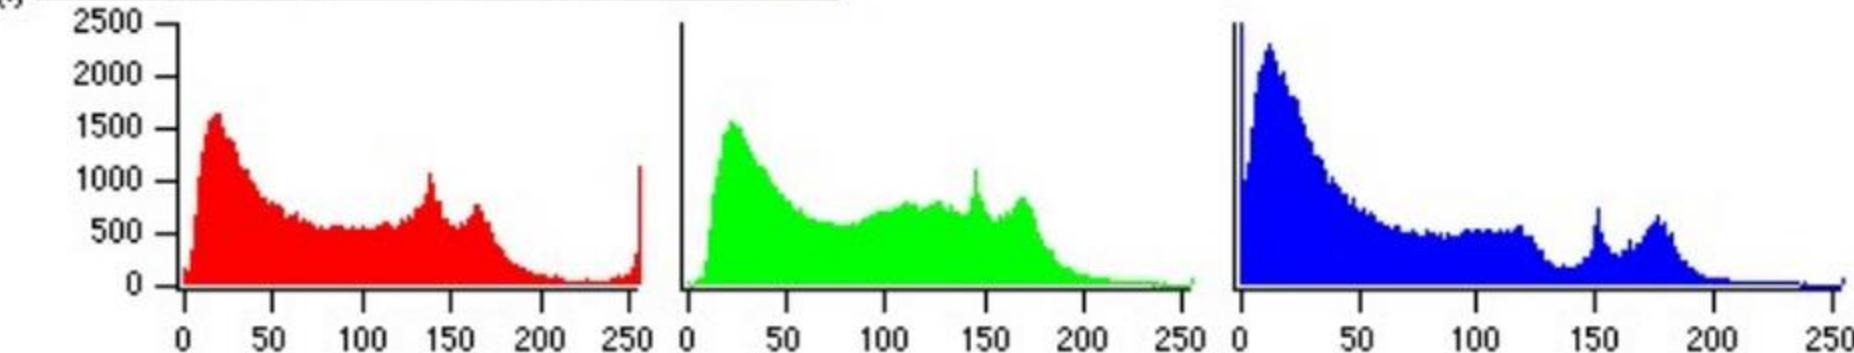
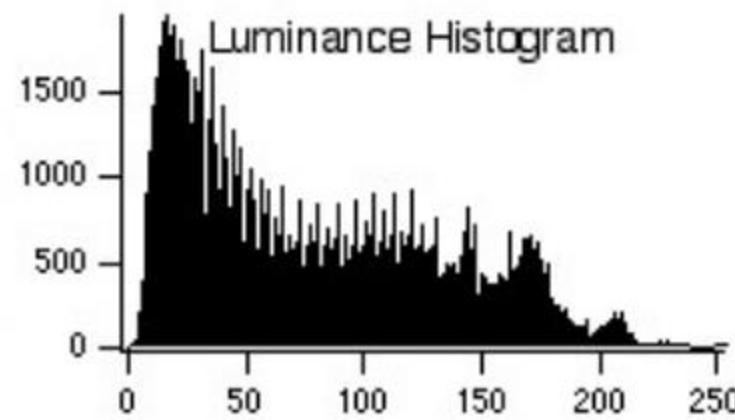
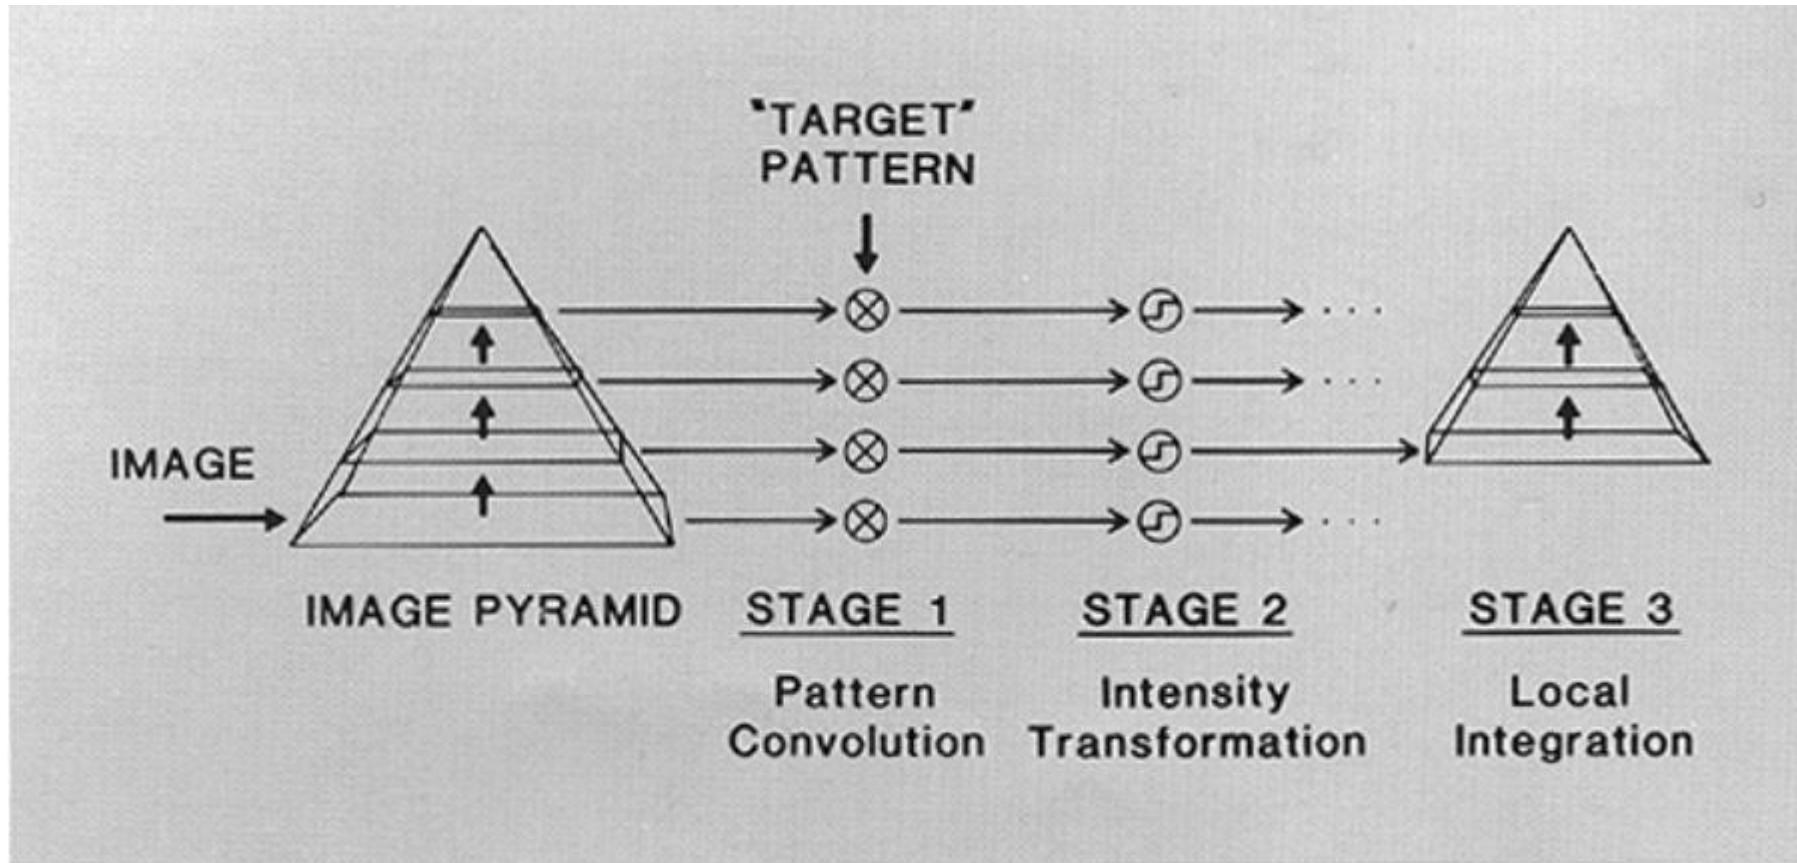
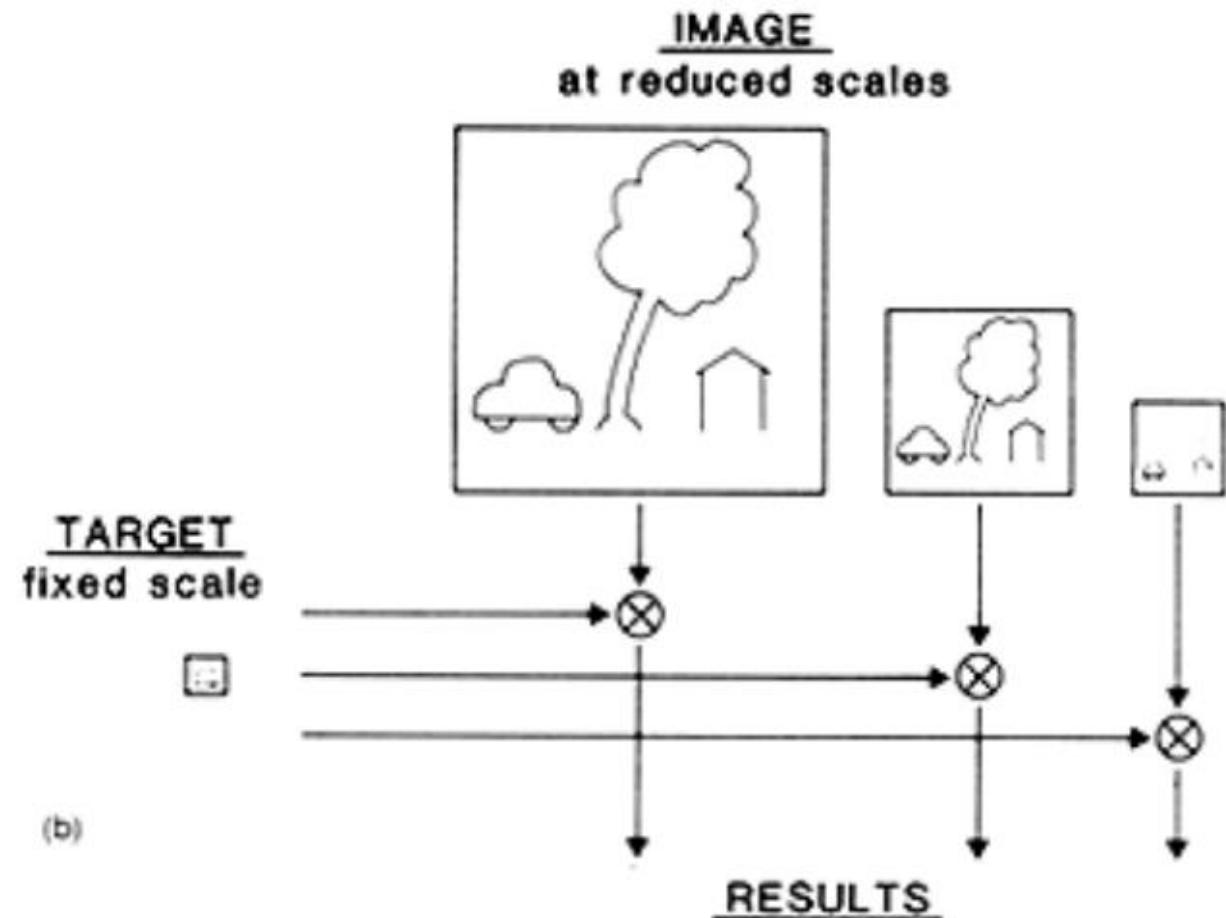
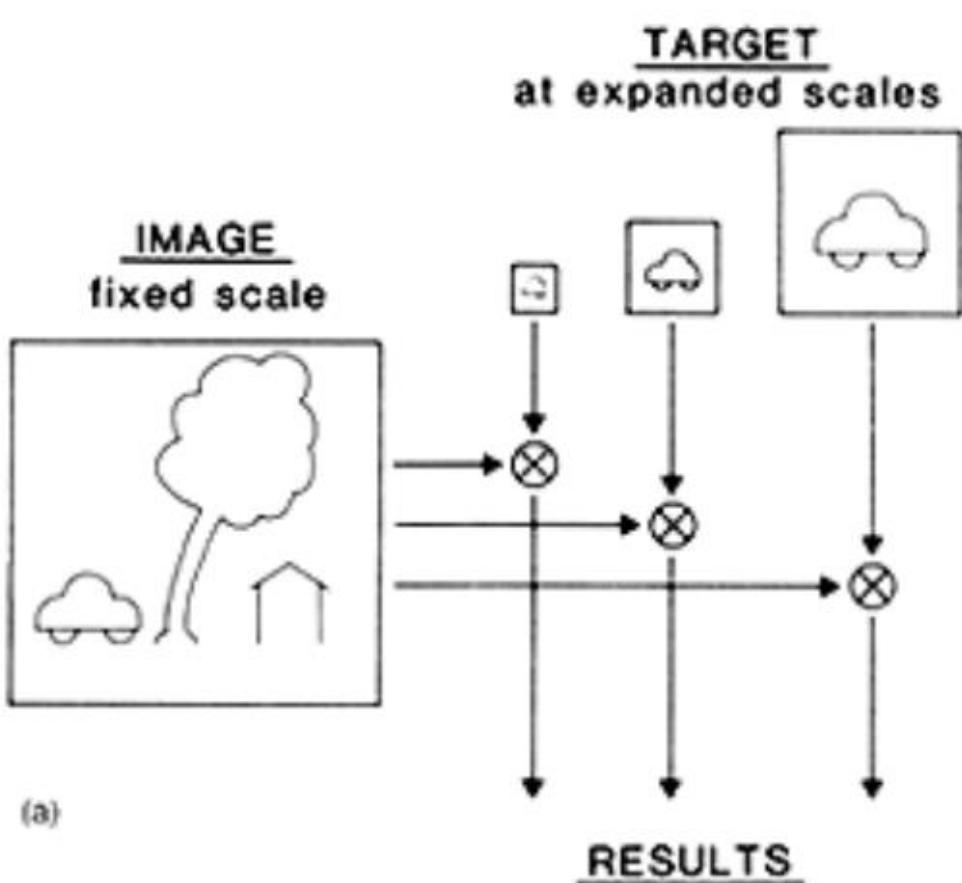


IMAGE PYRAMID (I)



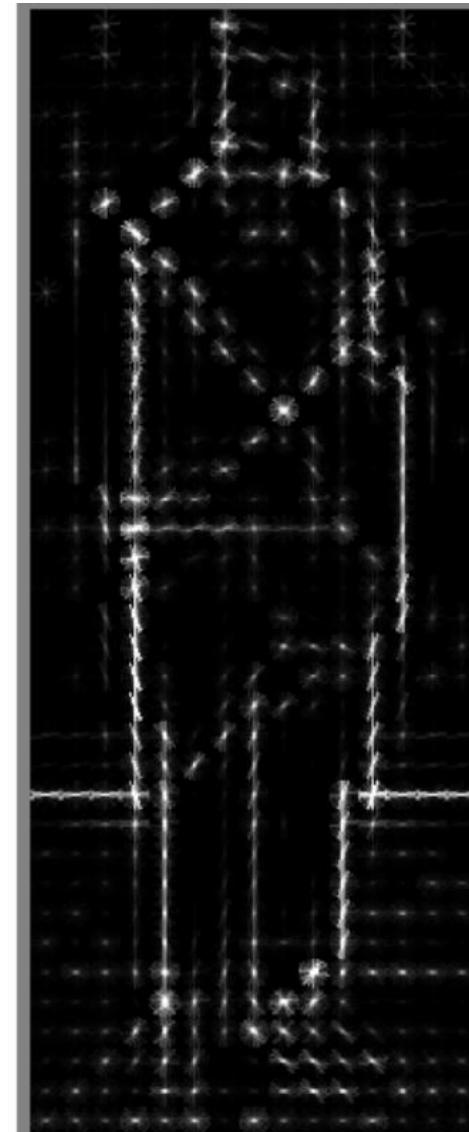
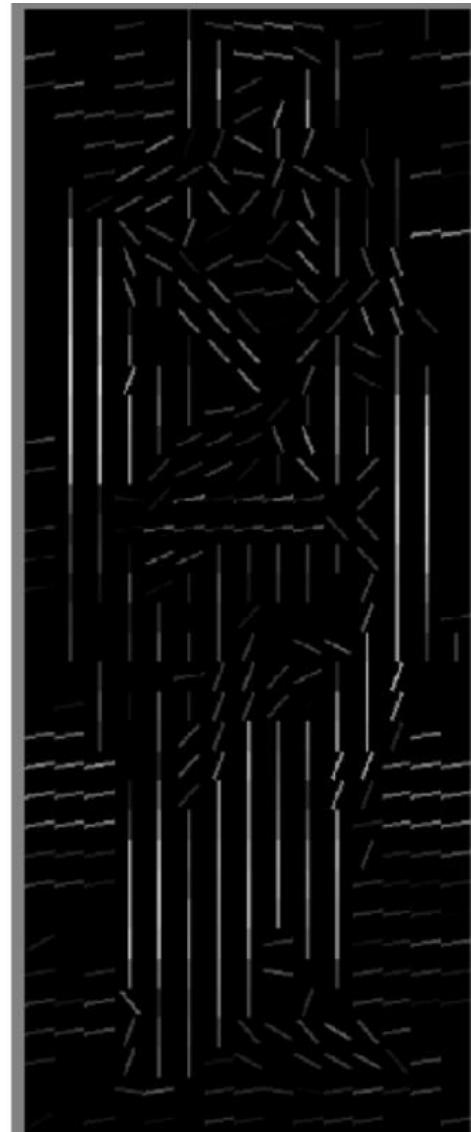
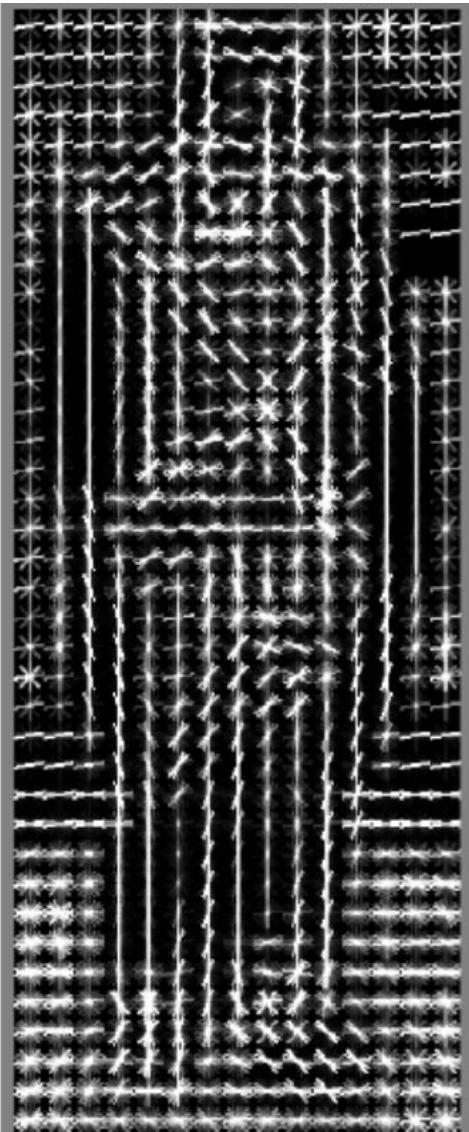
Adelson et al '84

IMAGE PYRAMID (2)



Adelson et al '84

HISTOGRAMS OF GRADIENTS (HOG)



EDGE DETECTION (WITH CONVOLUTION)

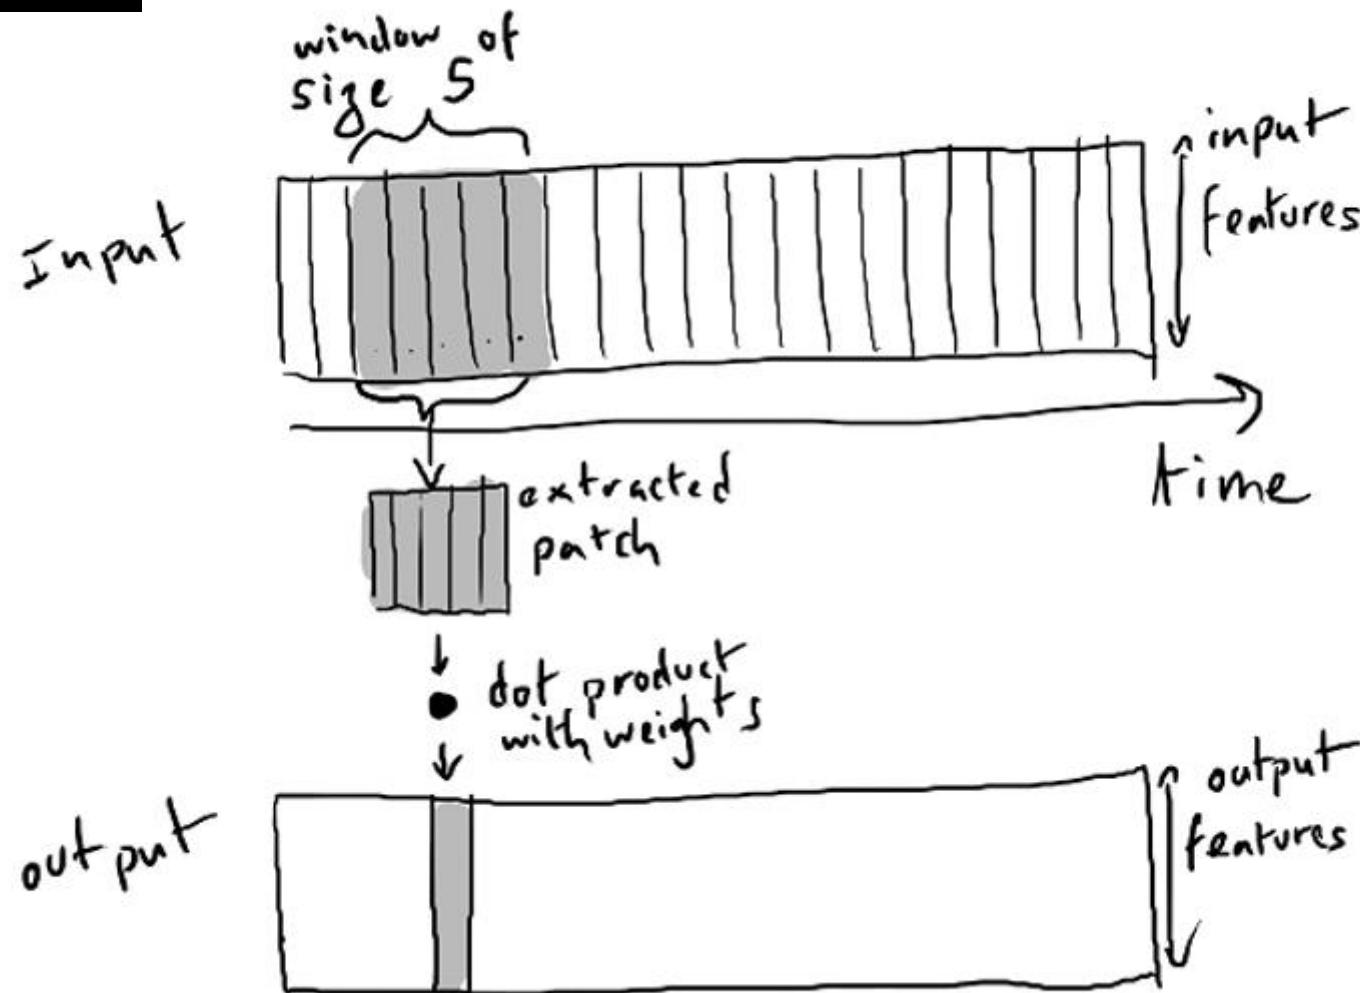




How do convolutional neural networks work?

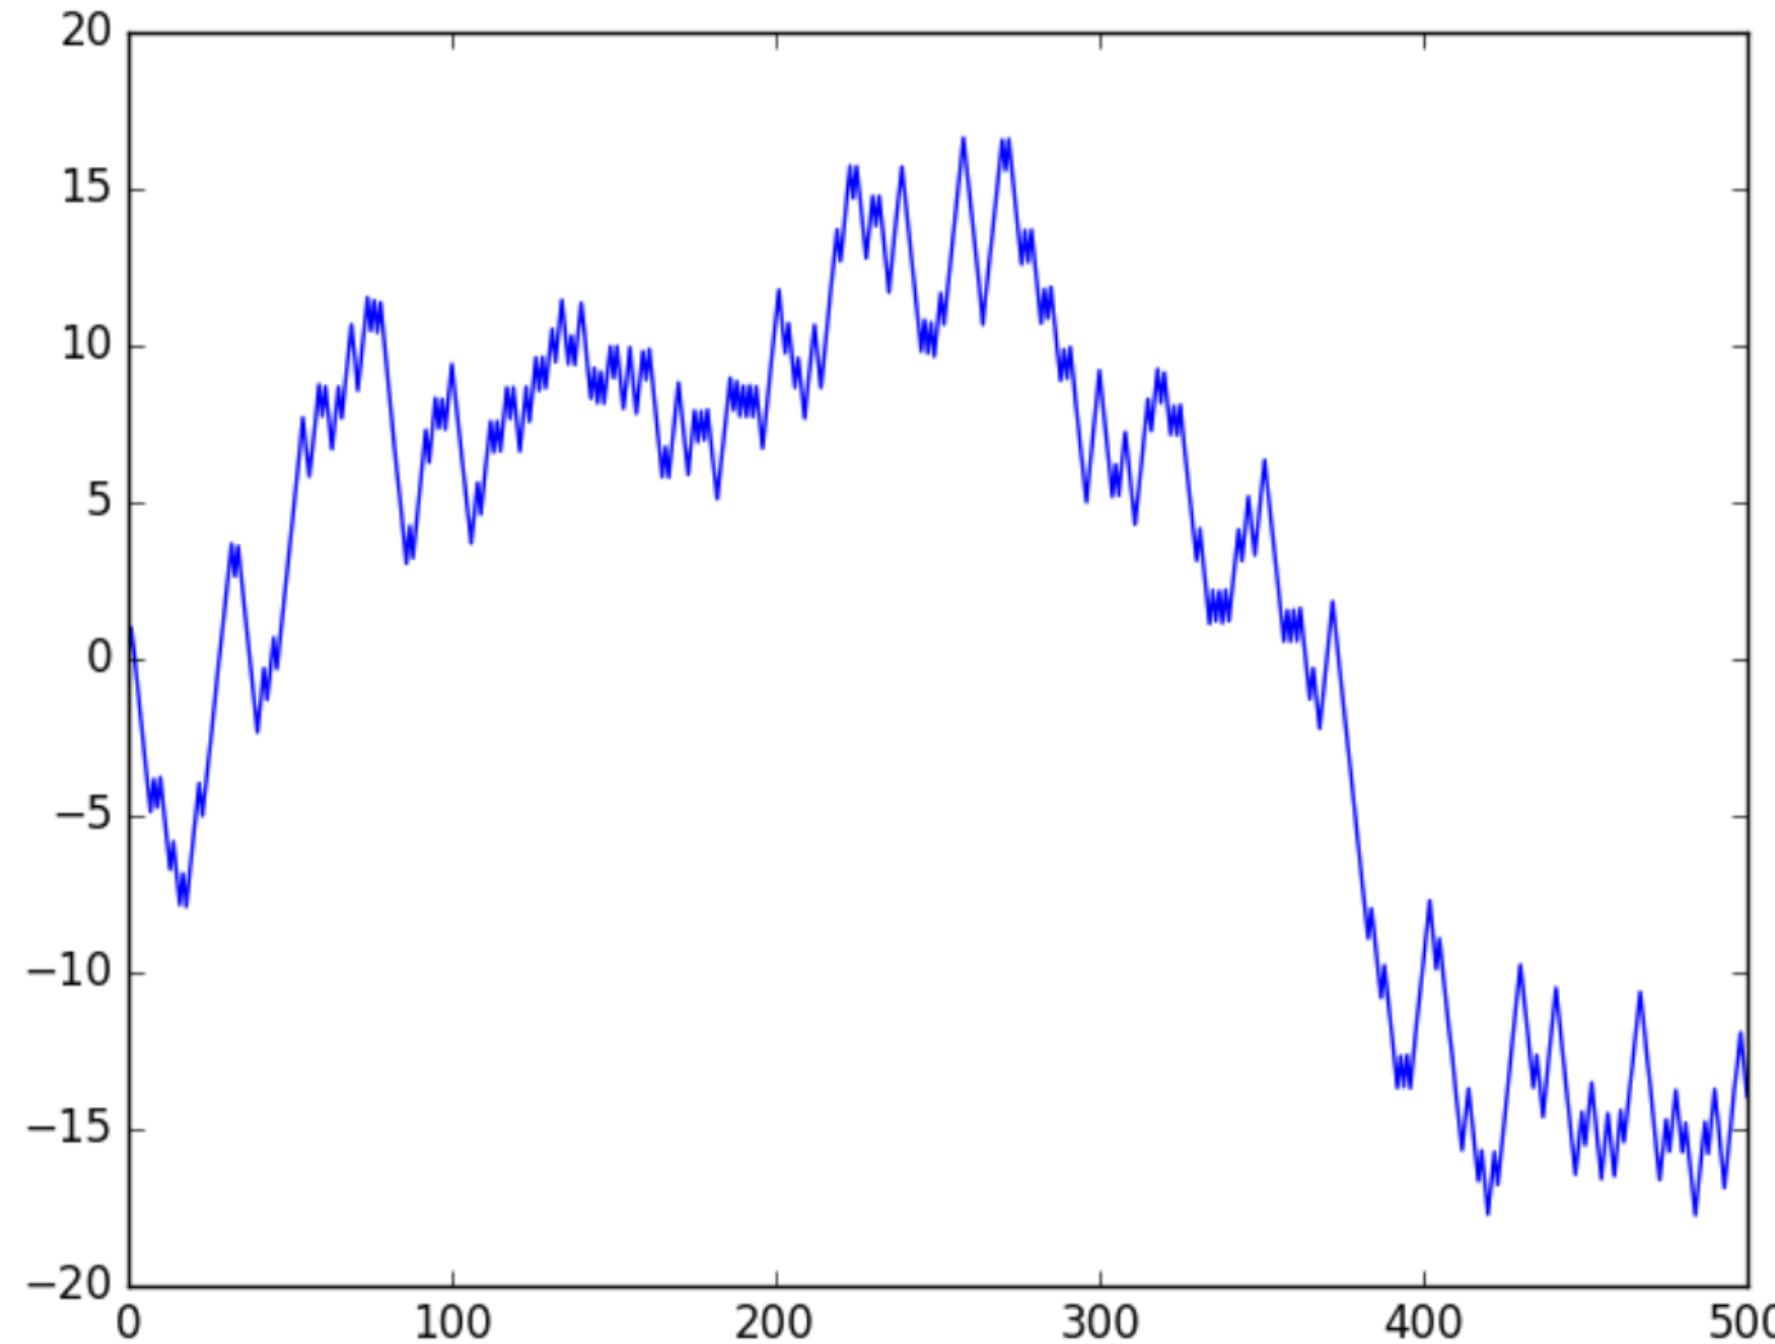


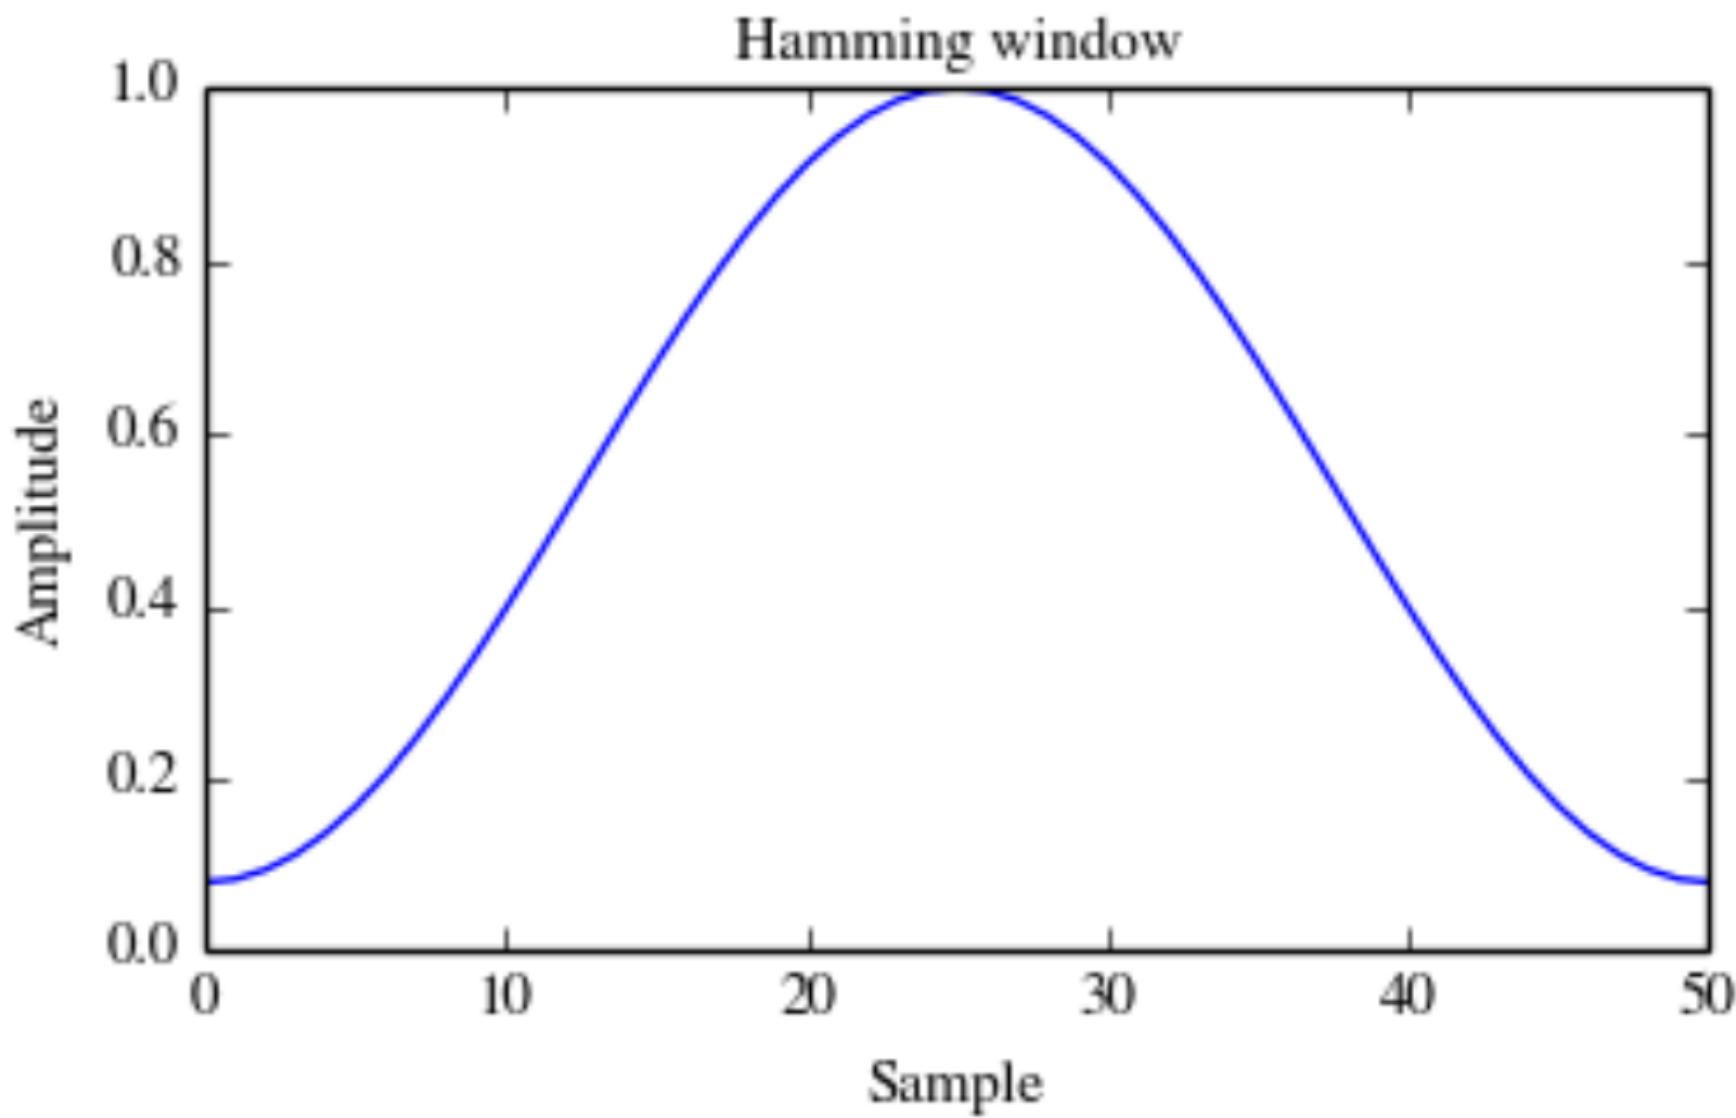
CONVOLUTION



© Francois Chollet

```
1 x = [0]
2
3 for j in range(500):
4     step_x = random.randint(0,1)
5     if step_x == 1:
6         x.append(x[j] + 1 + 0.05*np.random.normal())
7     else:
8         x.append(x[j] - 1 + 0.05*np.random.normal())
```

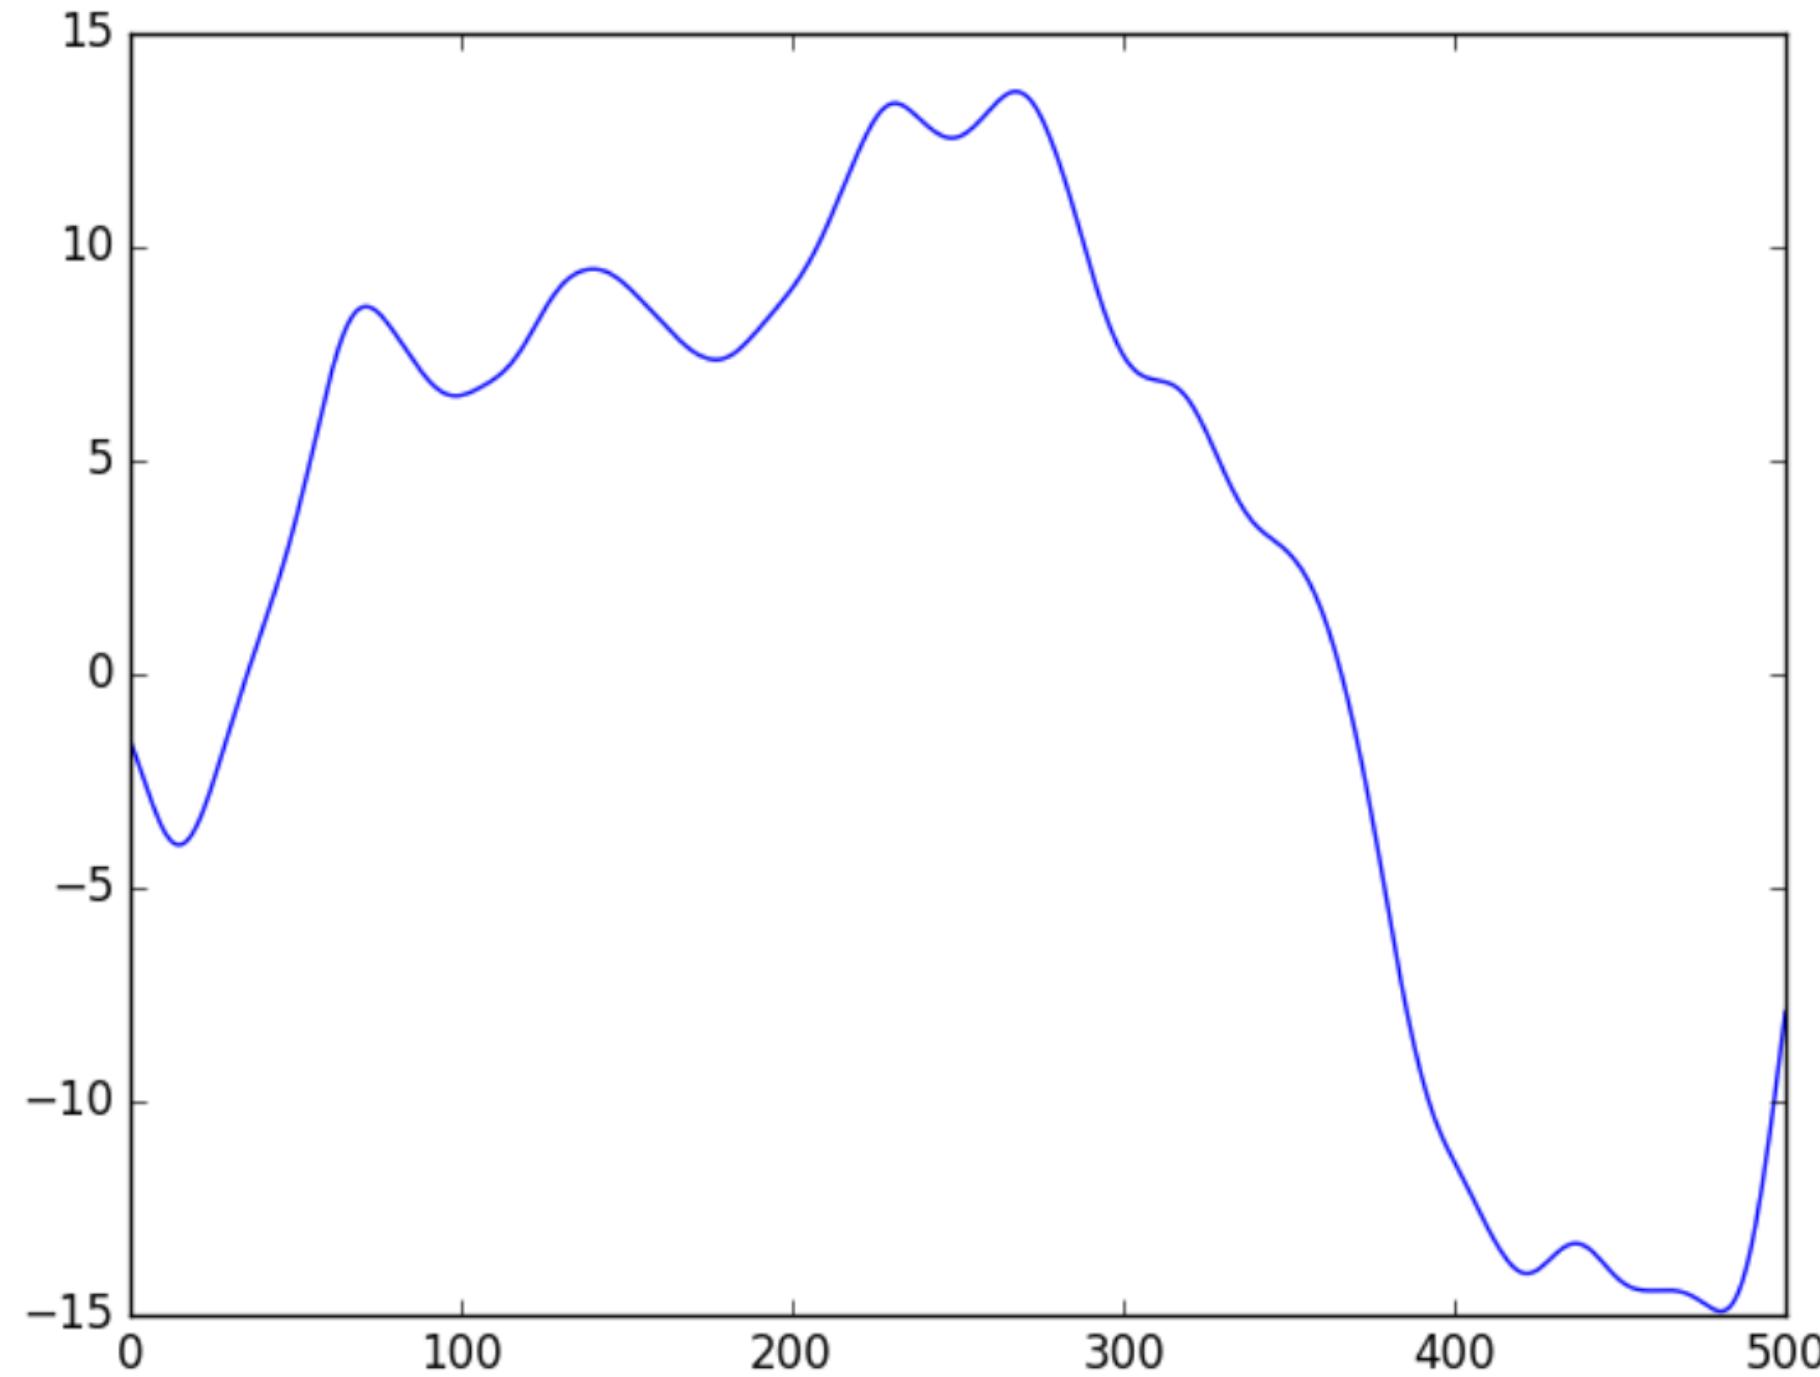


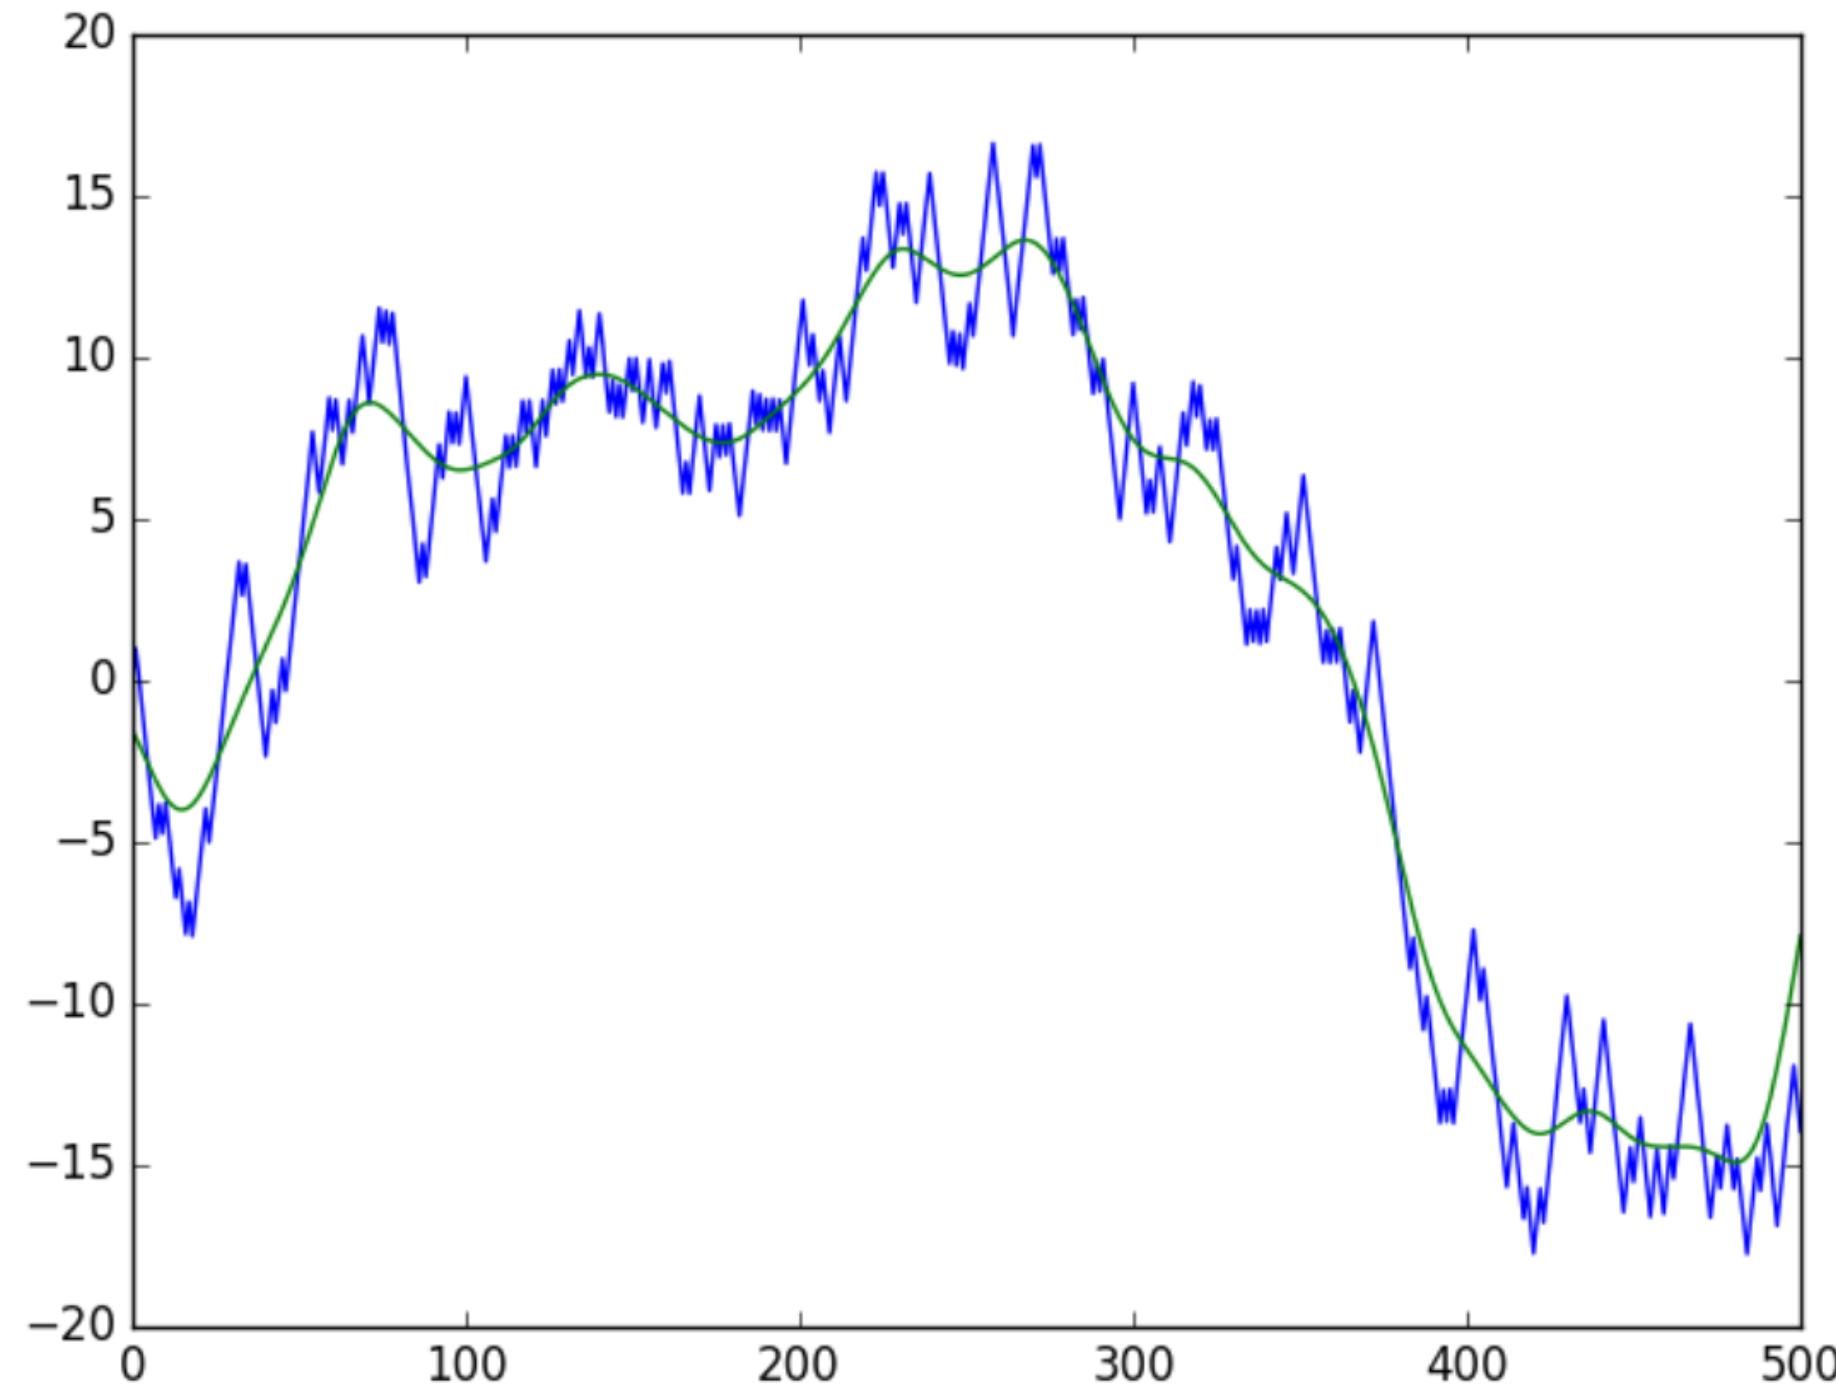


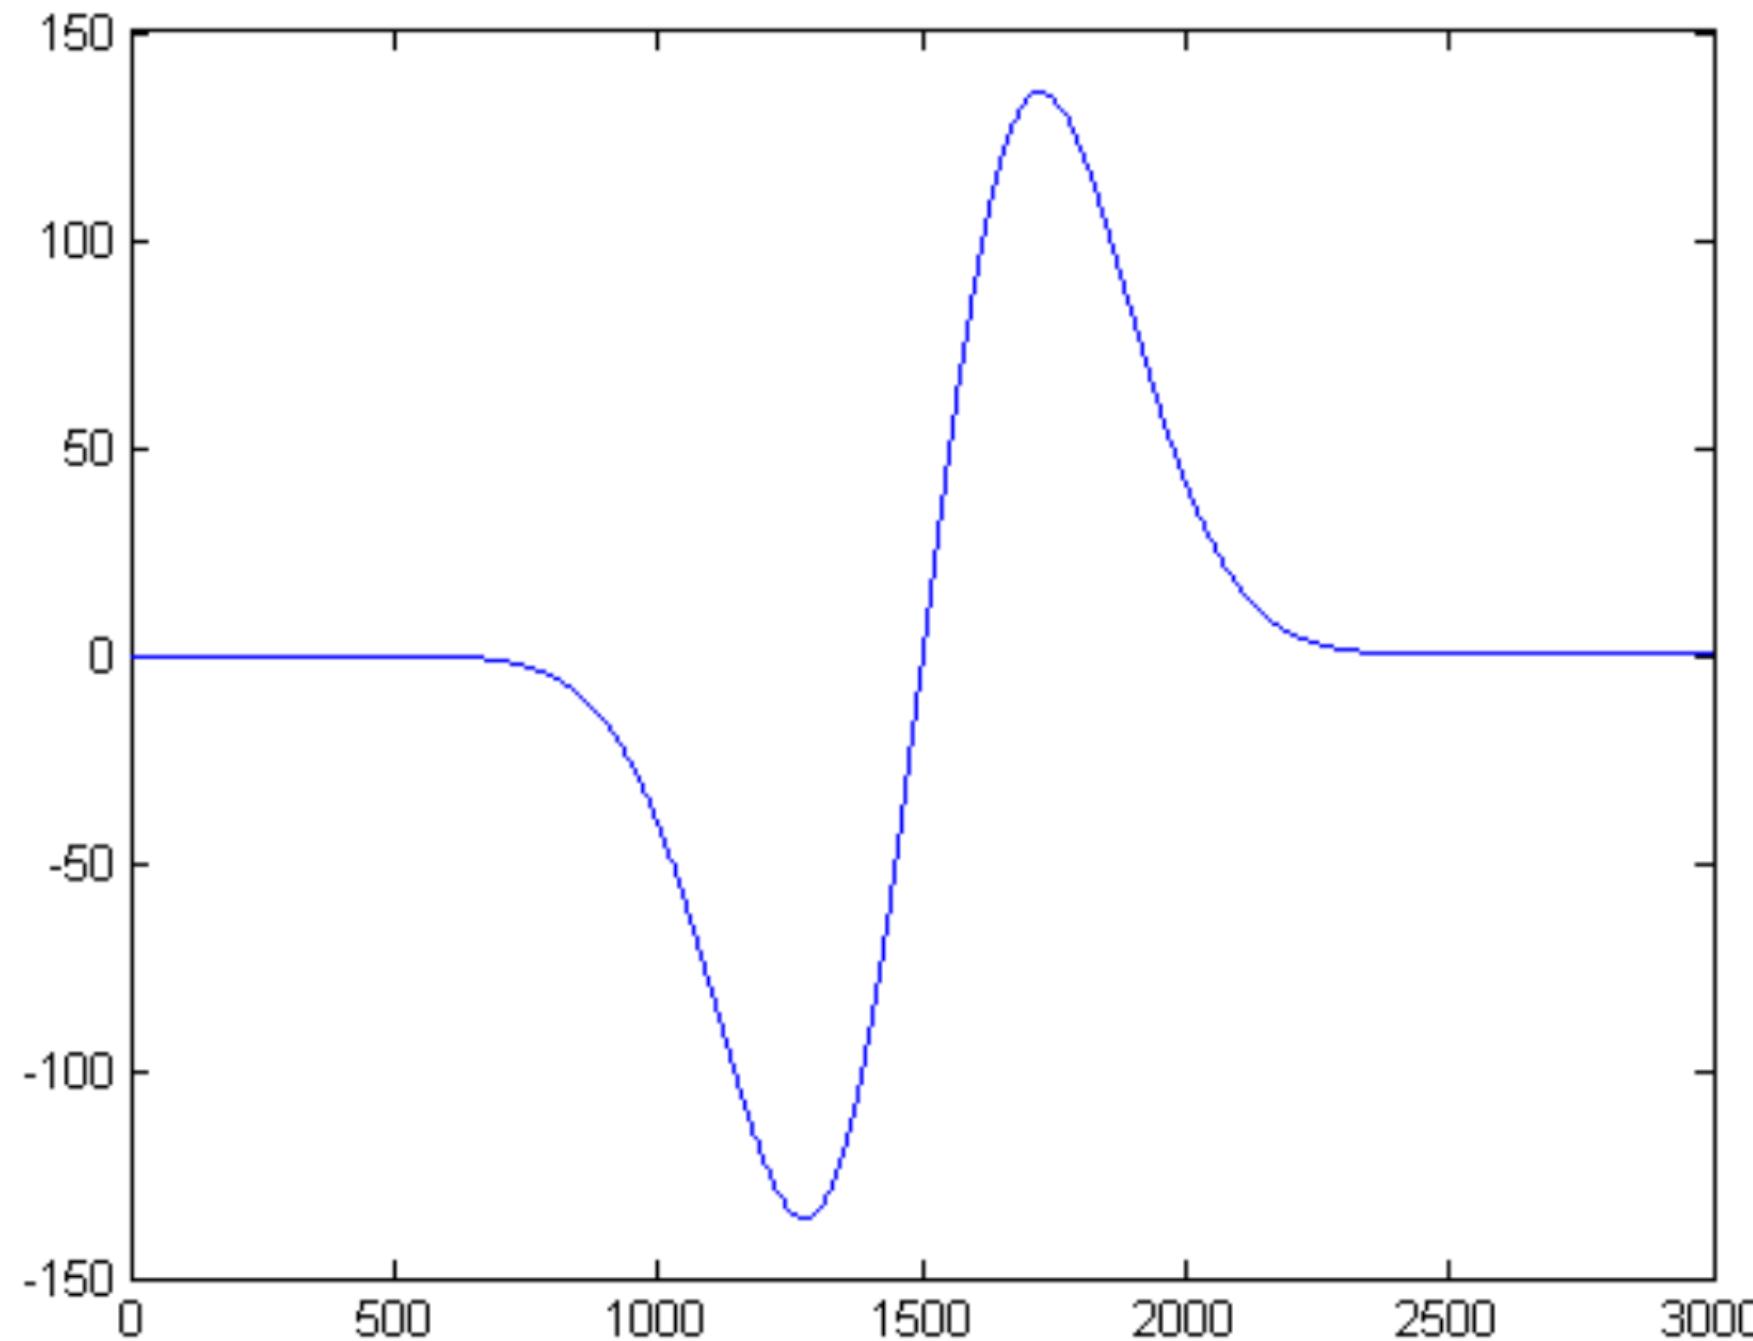
$$w(n) = 0.54 - 0.46 \cos\left(2\pi \frac{n}{N}\right), \quad 0 \leq n \leq N.$$

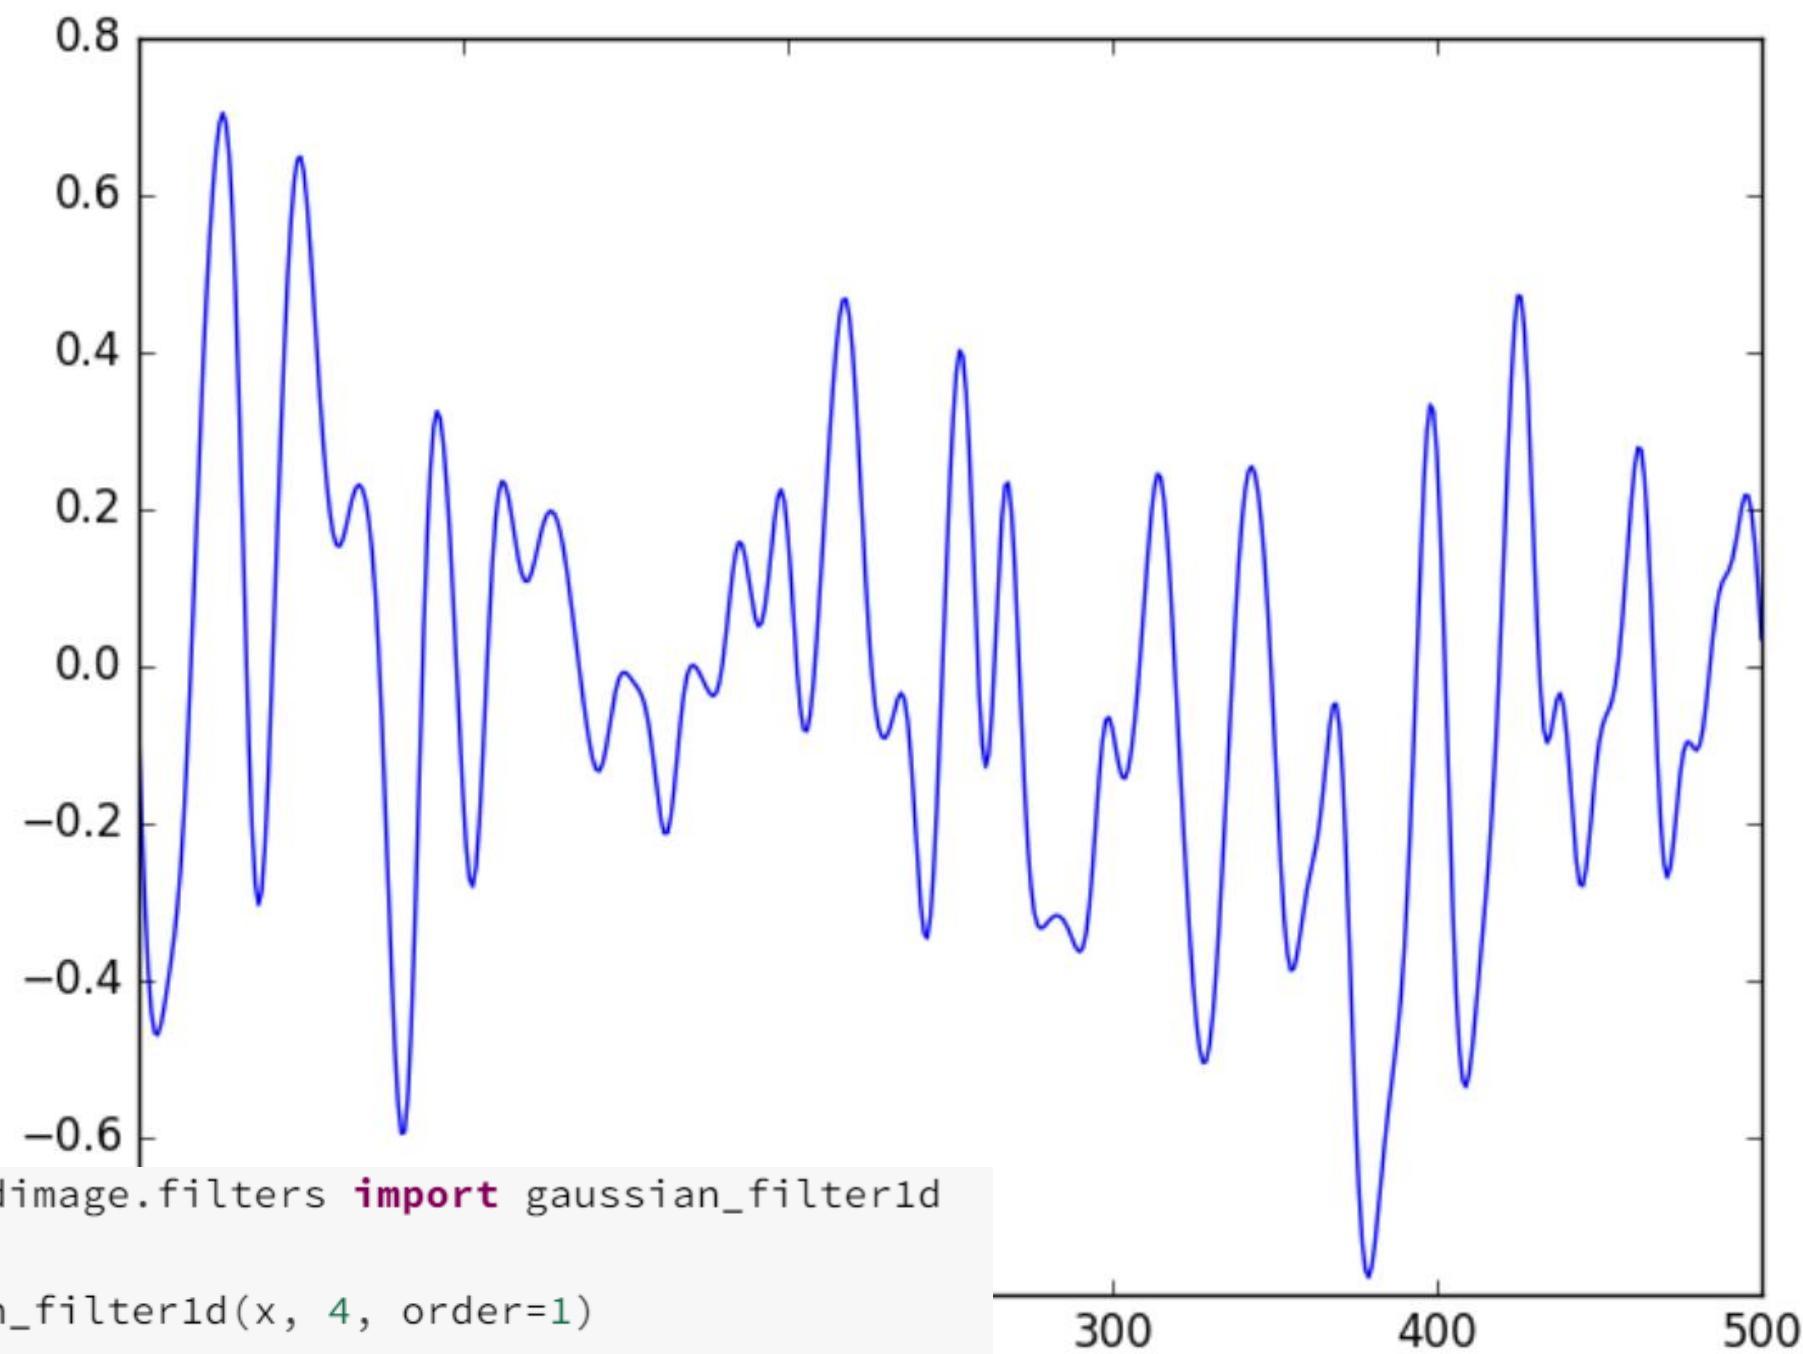
```
from scipy import signal

win = signal.hann(50)
filtered = signal.convolve(x, win, mode='same') / sum(win)
```

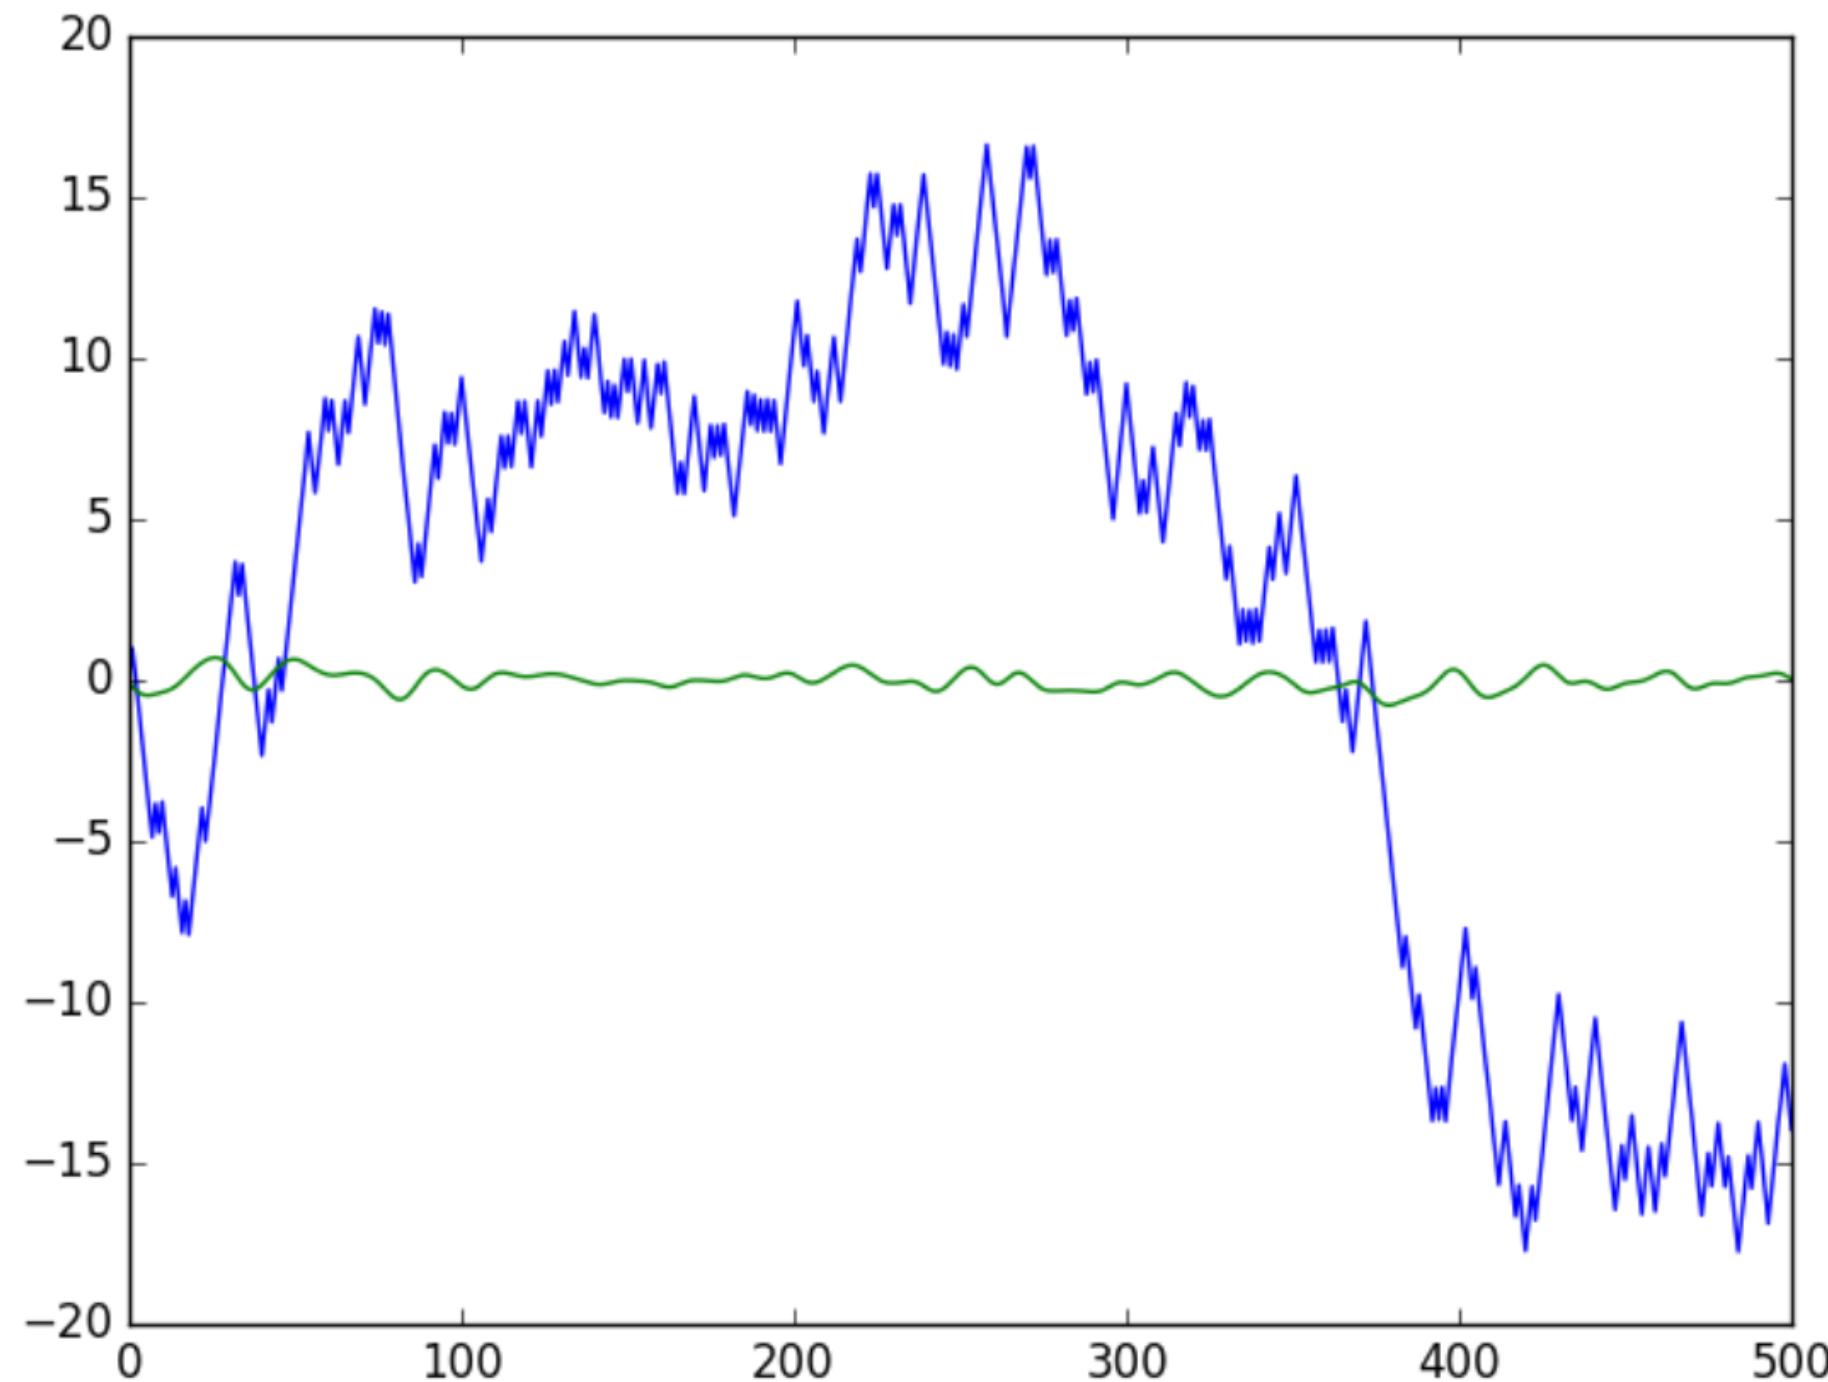






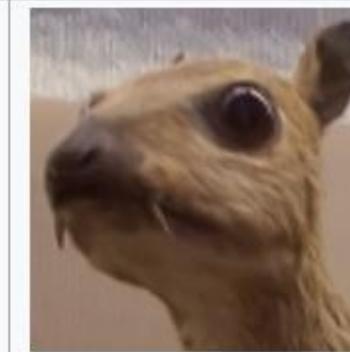


```
from scipy.ndimage.filters import gaussian_filter1d  
  
x2 = gaussian_filter1d(x, 4, order=1)
```

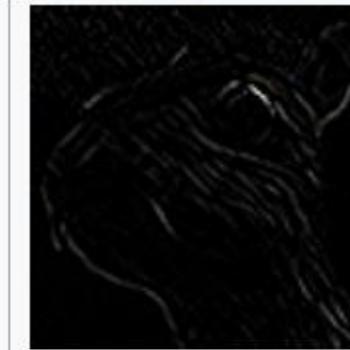


Identity

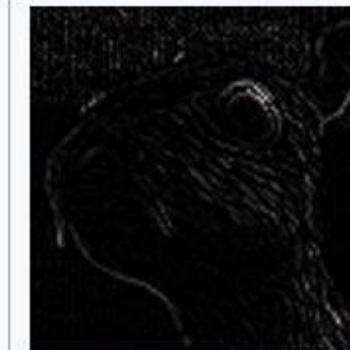
$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

**Edge detection**

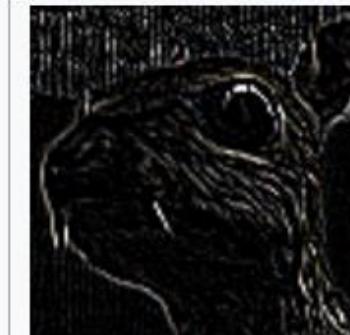
$$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

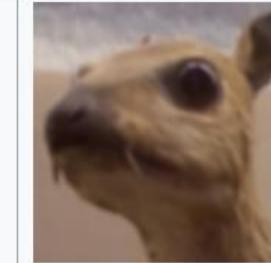
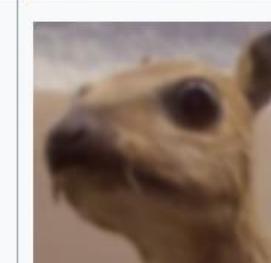


$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

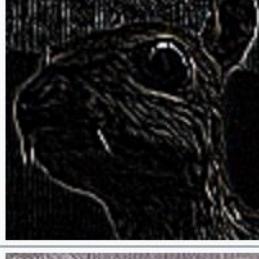


$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur 3×3 (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	
Gaussian blur 5×5 (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	
Unsharp masking 5×5 Based on Gaussian blur with amount as 1 and threshold as 0 (with no image mask)	$\frac{-1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & -476 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	

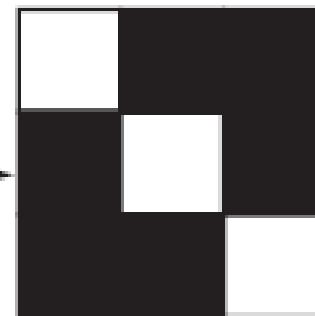
DEMO

Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	

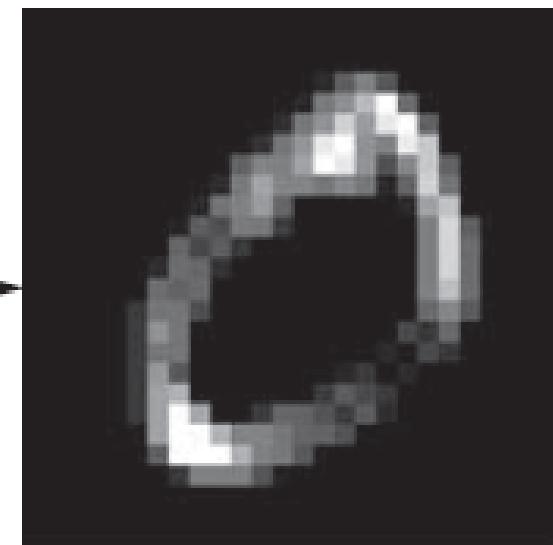
Original input



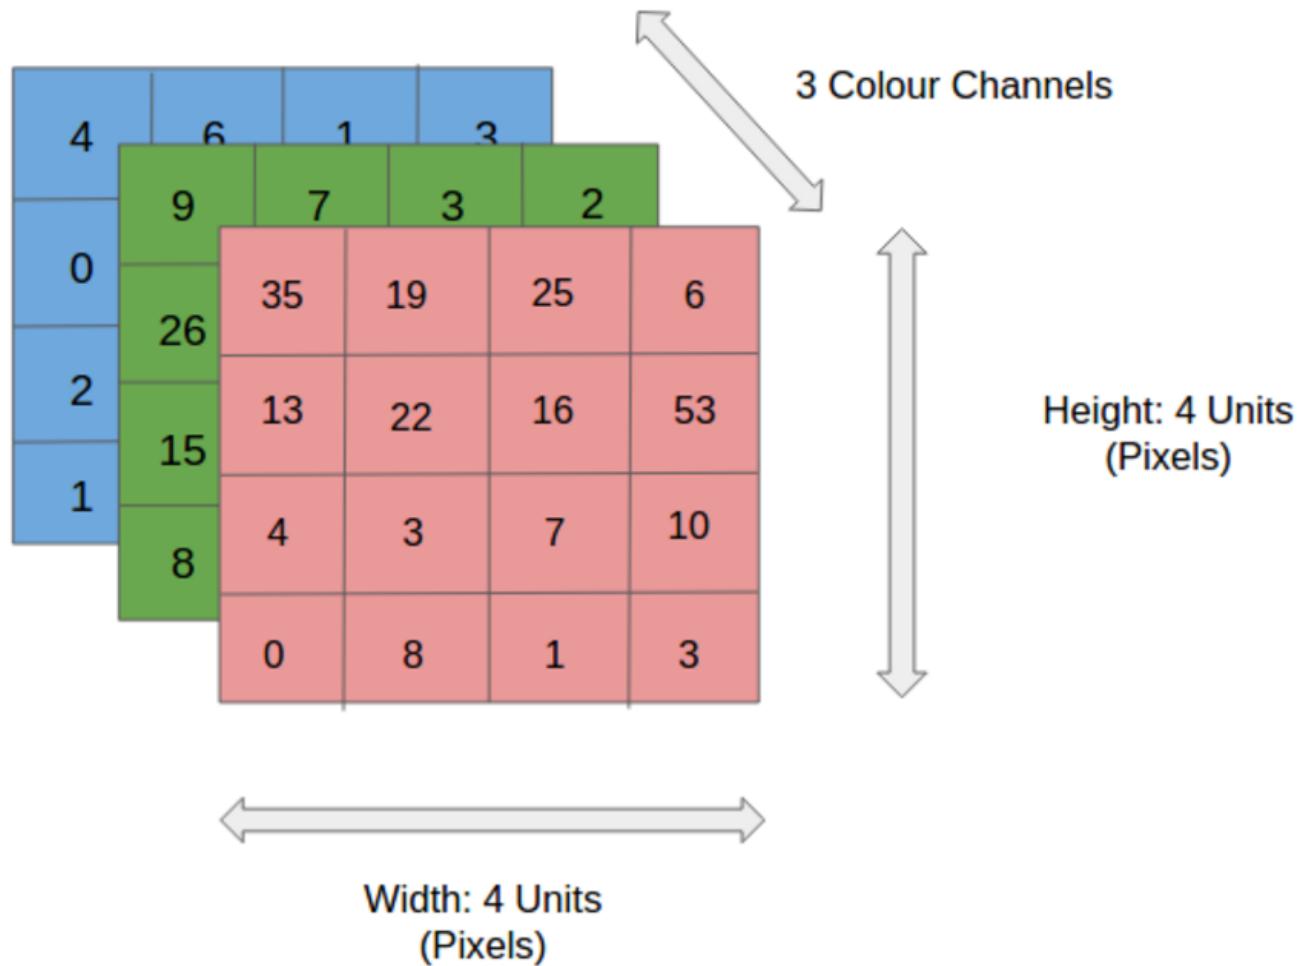
Single filter



Response map,
quantifying the presence
of the filter's pattern at
different locations

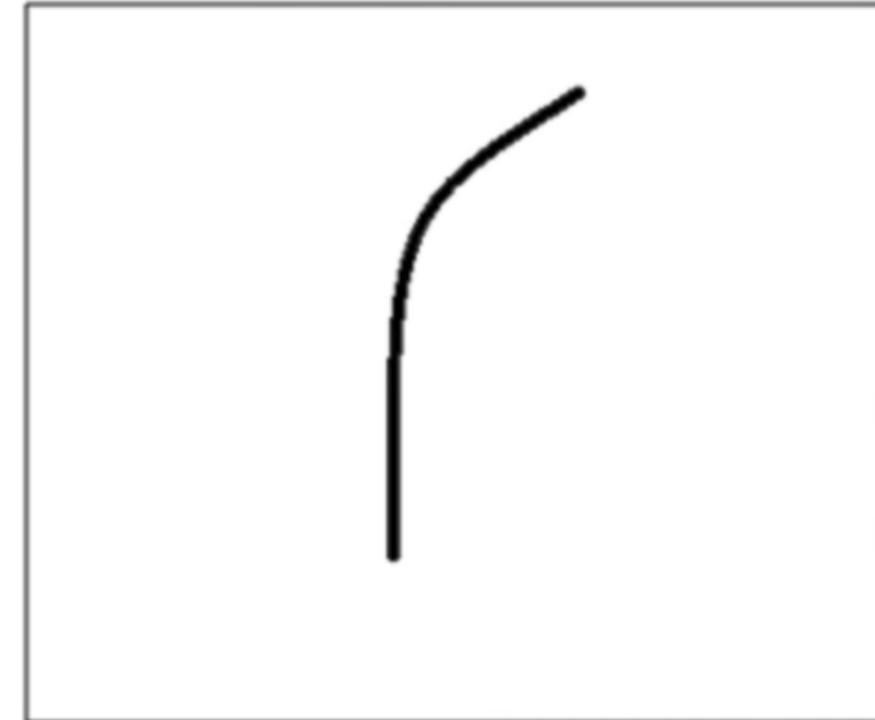


PREPARE DATASET OF IMAGES



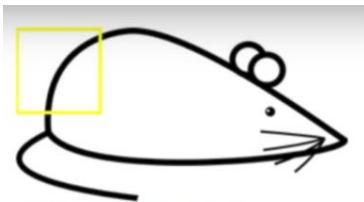
0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter



Visualization of a curve detector filter

CONVOLUTION FILTER



Visualization of the filter on the image



Visualization of the receptive field

$$(50 \times 30) + (50 \times 30) + (50 \times 30) + (20 \times 30) + (50 \times 30) = 6600$$

0	0	0	0	0	0	30
0	0	0	0	50	50	50
0	0	0	20	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0
0	0	0	50	50	0	0

Pixel representation of the receptive field

*

0	0	0	0	0	0	30	0
0	0	0	0	0	30	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	0	0	0	0	0

Pixel representation of filter

CONVOLUTION FILTER MATCH



Visualization of the filter on the image

MULTIPLY AND SUMMATION = 0

0	0	0	0	0	0	0
0	40	0	0	0	0	0
40	0	40	0	0	0	0
40	20	0	0	0	0	0
0	50	0	0	0	0	0
0	0	50	0	0	0	0
25	25	0	50	0	0	0

Pixel representation of receptive field

*

0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

CONVOLUTION FILTER NO MATCH

1 x1	1 x0	1 x1	0	0
0 x0	1 x1	1 x0	1	0
0 x1	0 x0	1 x1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

1	1	1	0	0
0	1	1 x1	1 x0	0 x1
0	0	1 x0	1 x1	1 x0
0	0	1 x1	1 x0	0 x1
0	1	1	0	0

Image

4	3	4
2	4	3

Convolved Feature

1	1	1	0	0
0	1	1	1	0
0	0	1 x1	1 x0	1 x1
0	0	1 x0	1 x1	0 x0
0	1	1 x1	0 x0	0 x1

Image

4	3	4
2	4	3
2	3	4

Convolved Feature



CONVOLUTION

4	6	1	3
0	8	12	9
2	3	16	100
1	46	74	27



35	19	25	6
13	22	16	63
4	3	7	10
9	8	1	3



(i)

(iii)

9	7	3	2
26	37	14	1
15	29	16	0
8	6	54	2



(ii)

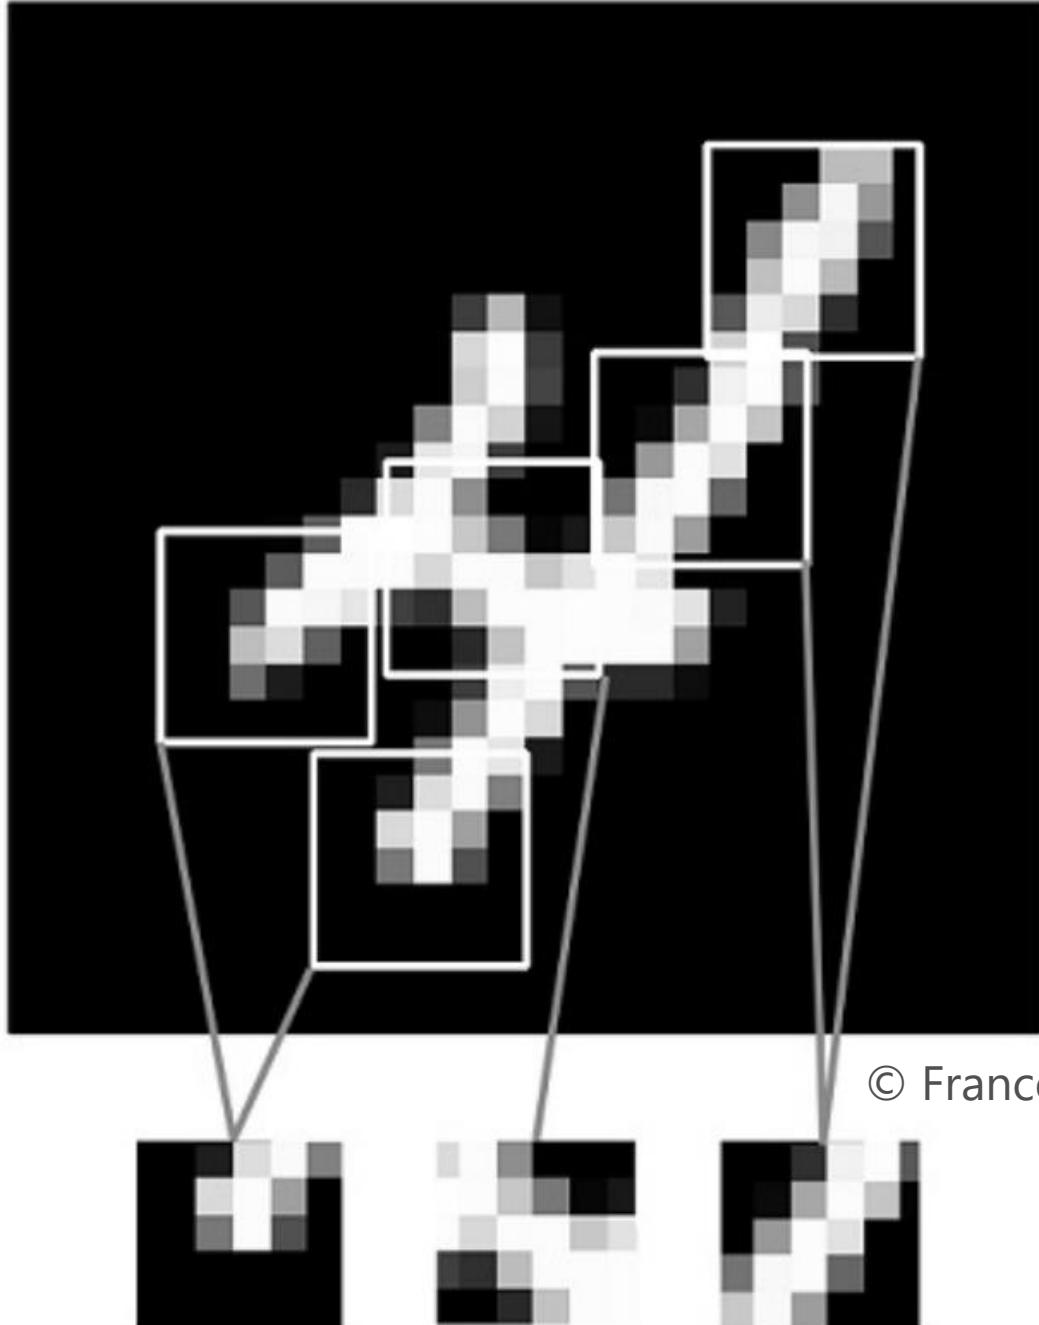
35	19	25	6
13	22	16	63
4	3	7	10
9	8	1	3



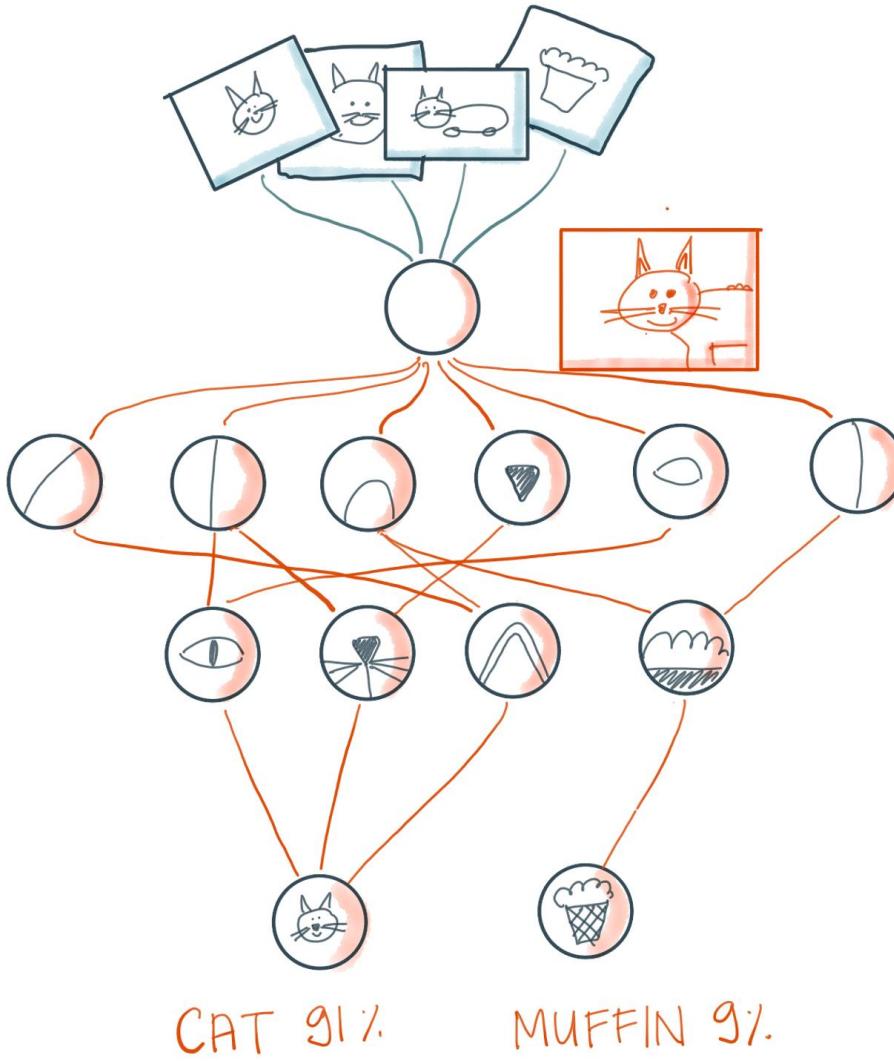
(iv)

POOLING

- Translation Invariant
- Spatial hierarchies of patterns
- Sort of scale invariant
- Not rotationally invariant!

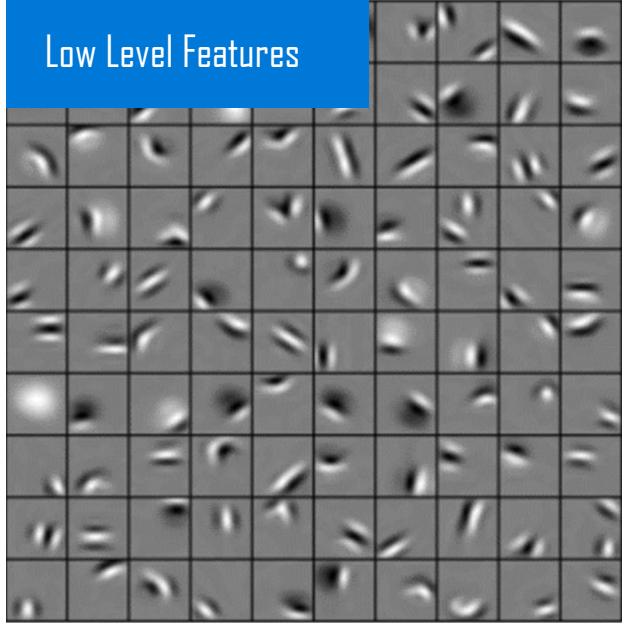


© Francois Chollet

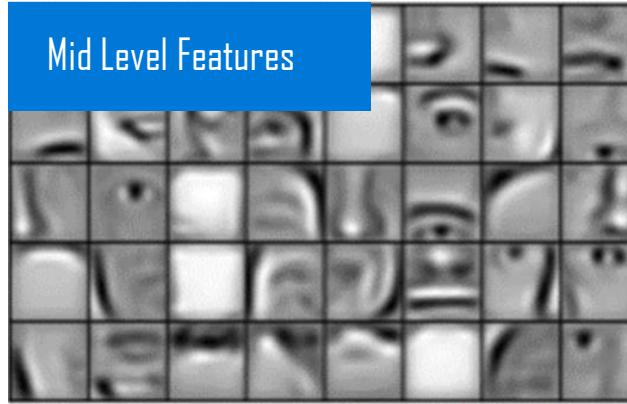


SPATIAL HIERARCHY OF PATTERNS

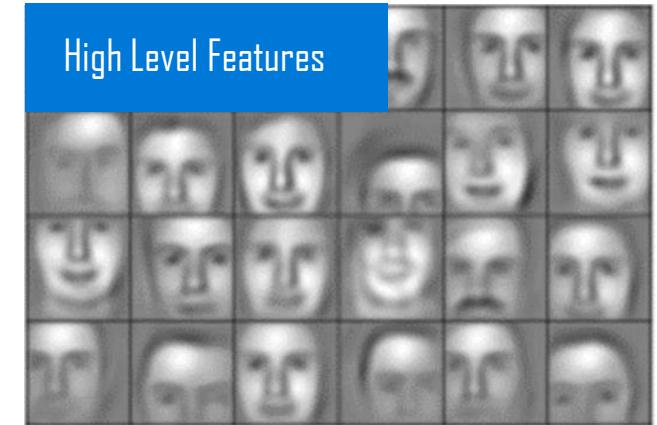
Low Level Features



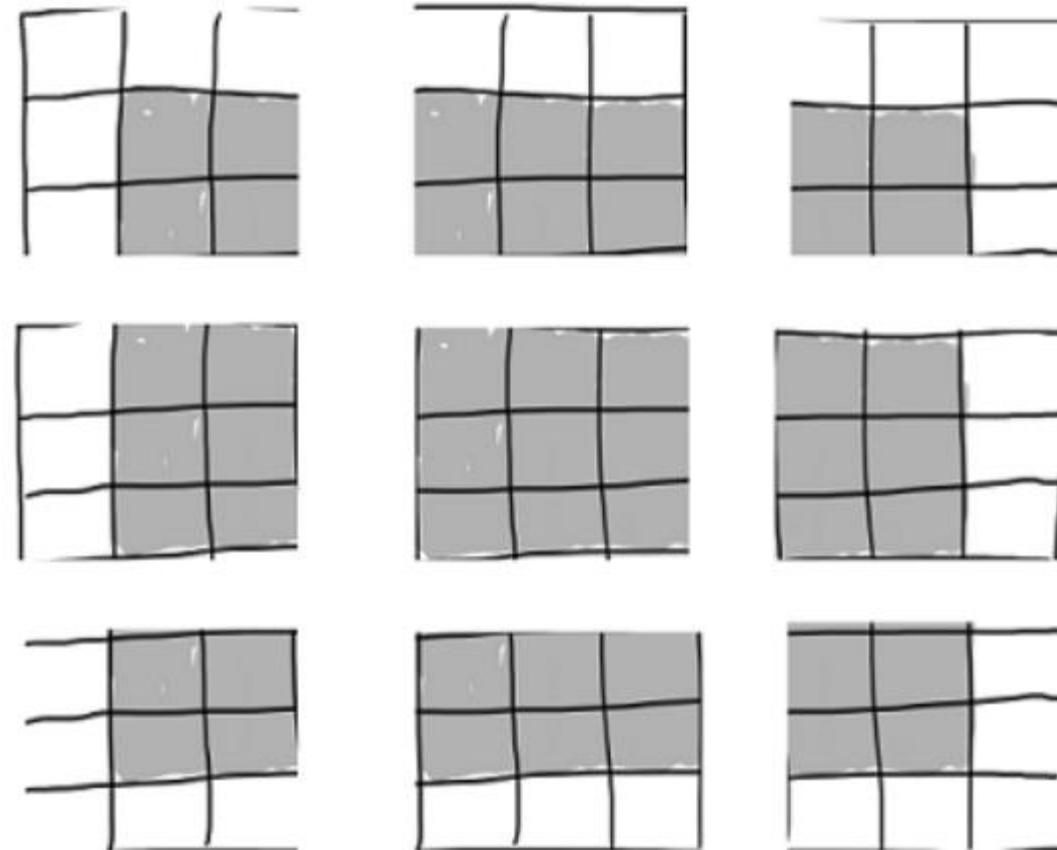
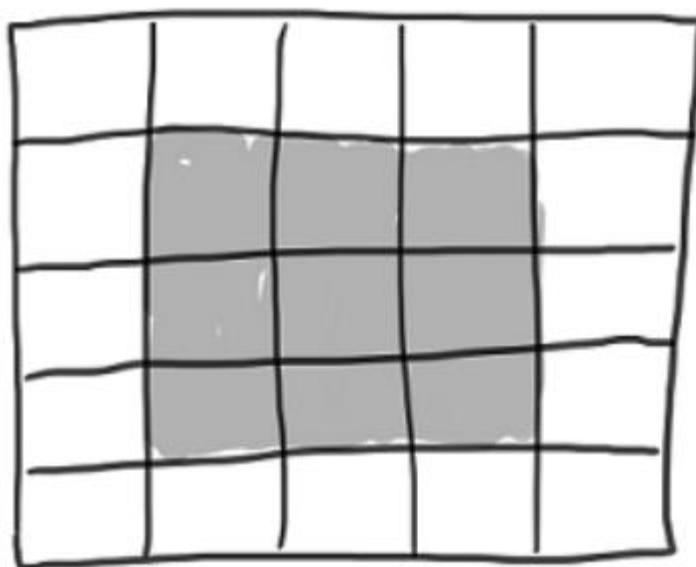
Mid Level Features



High Level Features



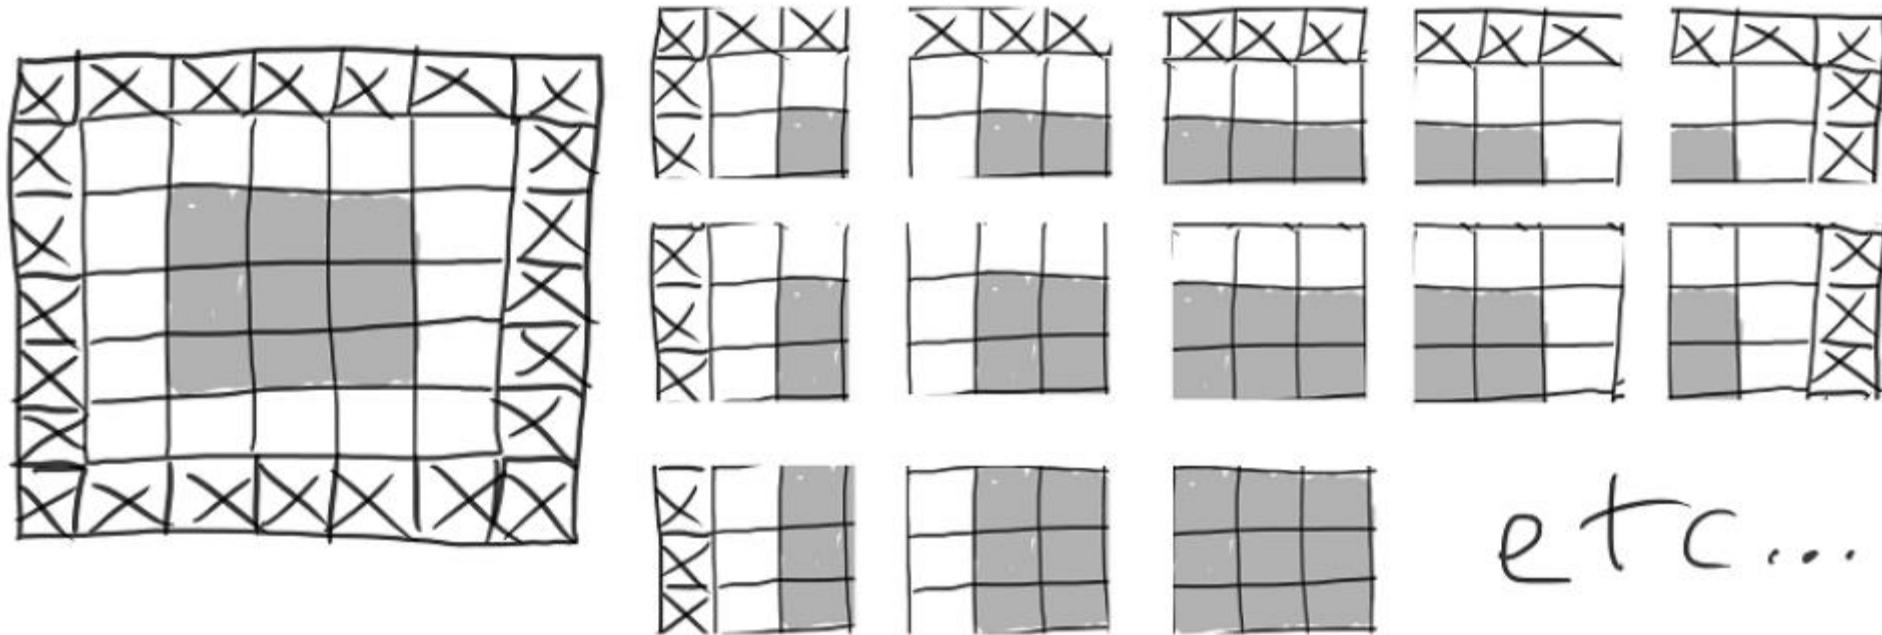
LEARNING REPRESENTATIONS



CONVOLUTION EXAMPLE

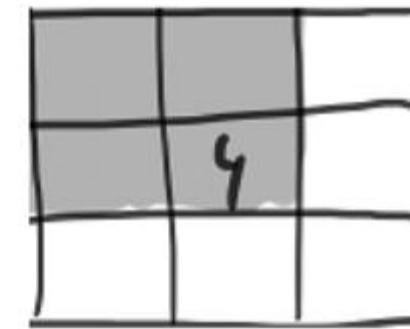
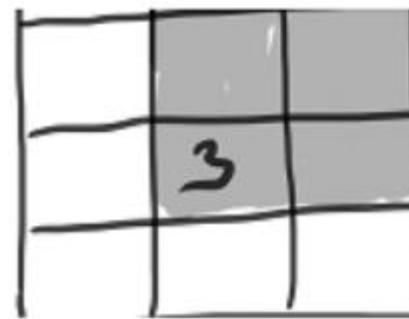
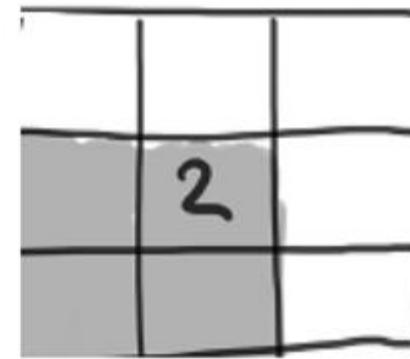
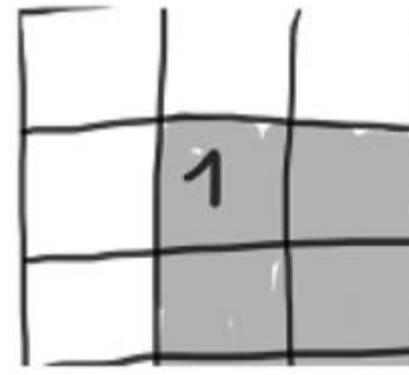
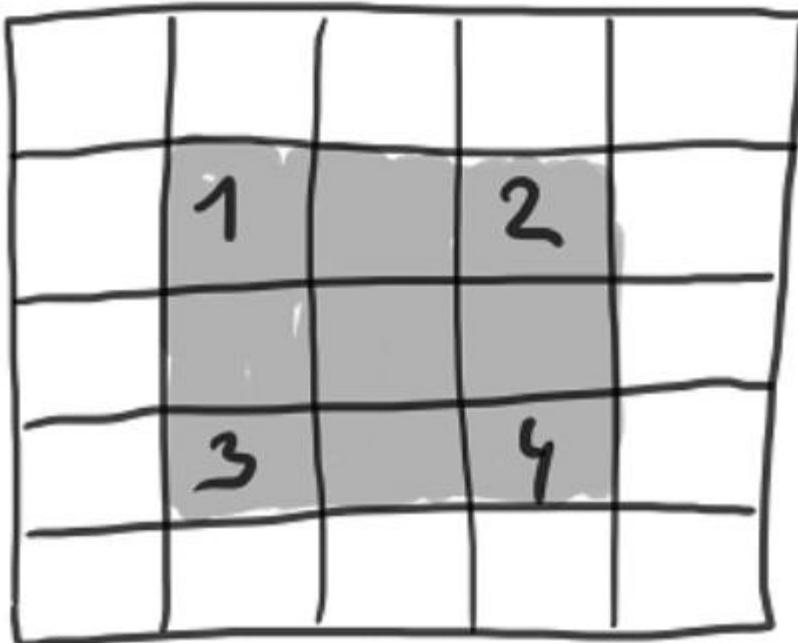
© Francois Chollet

© Francois Chollet



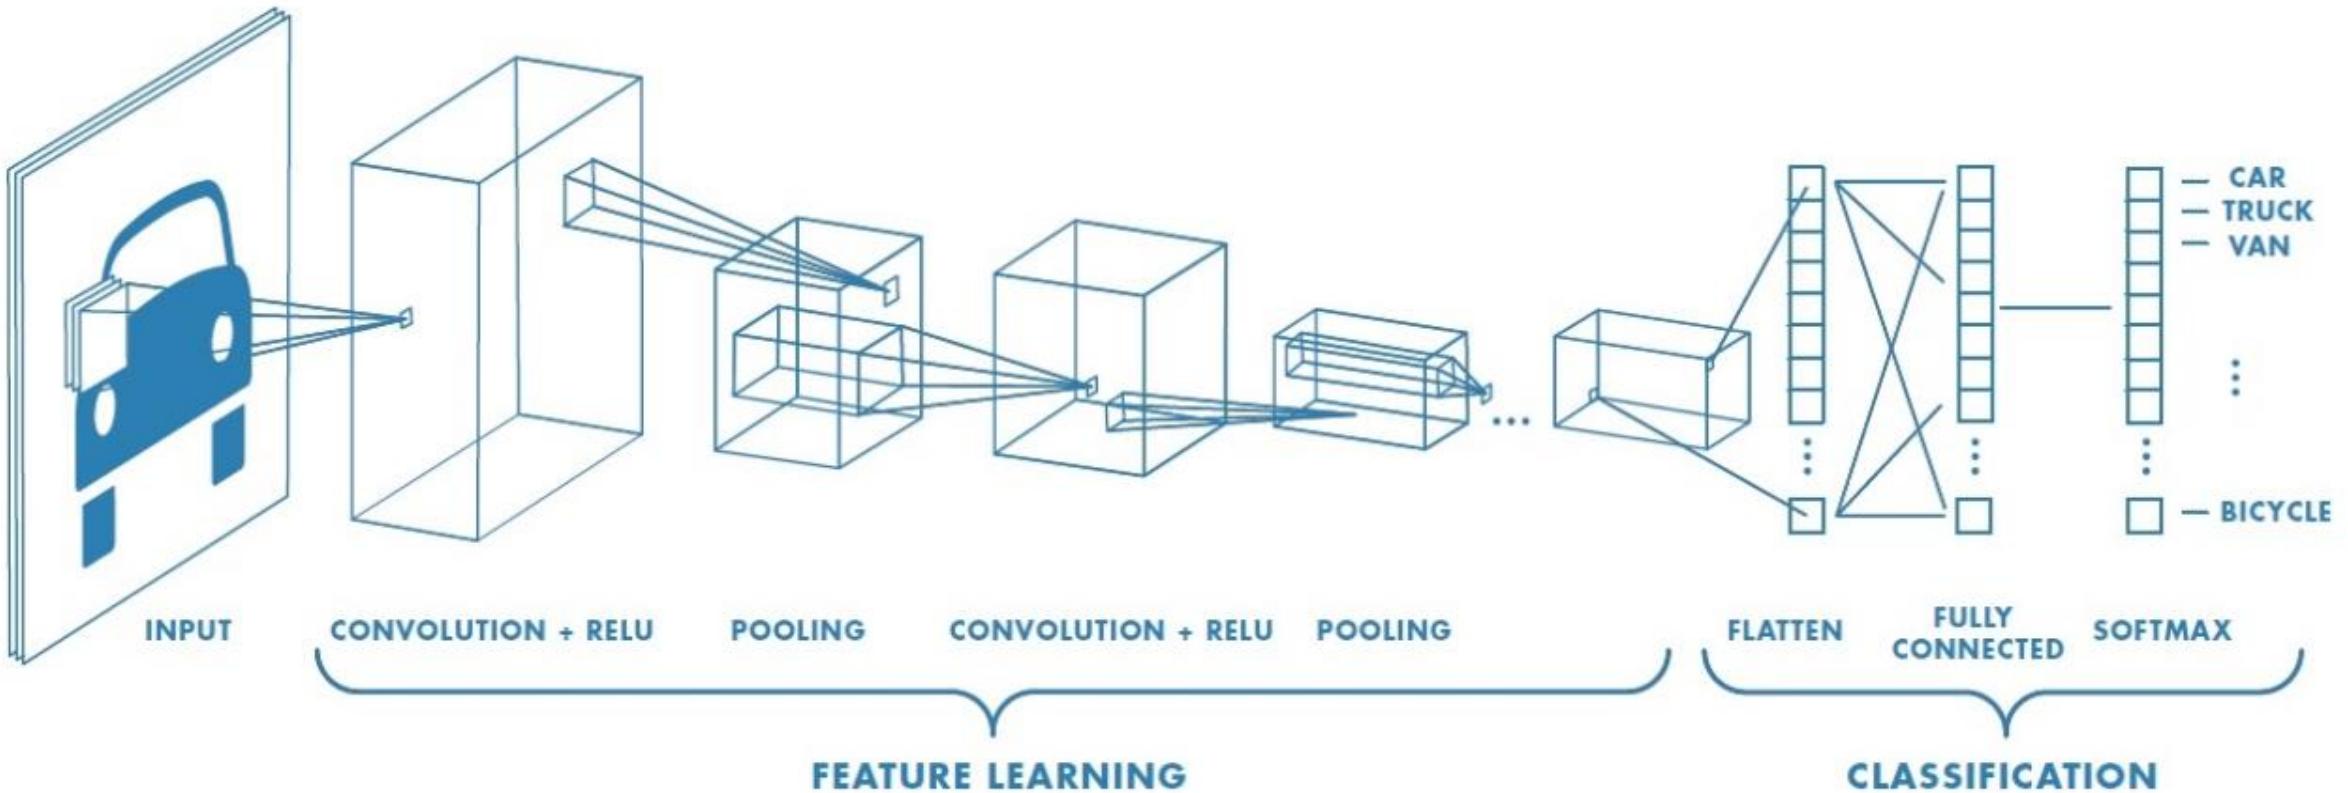
BORDER EFFECTS: 'VALID' VS 'SAME'

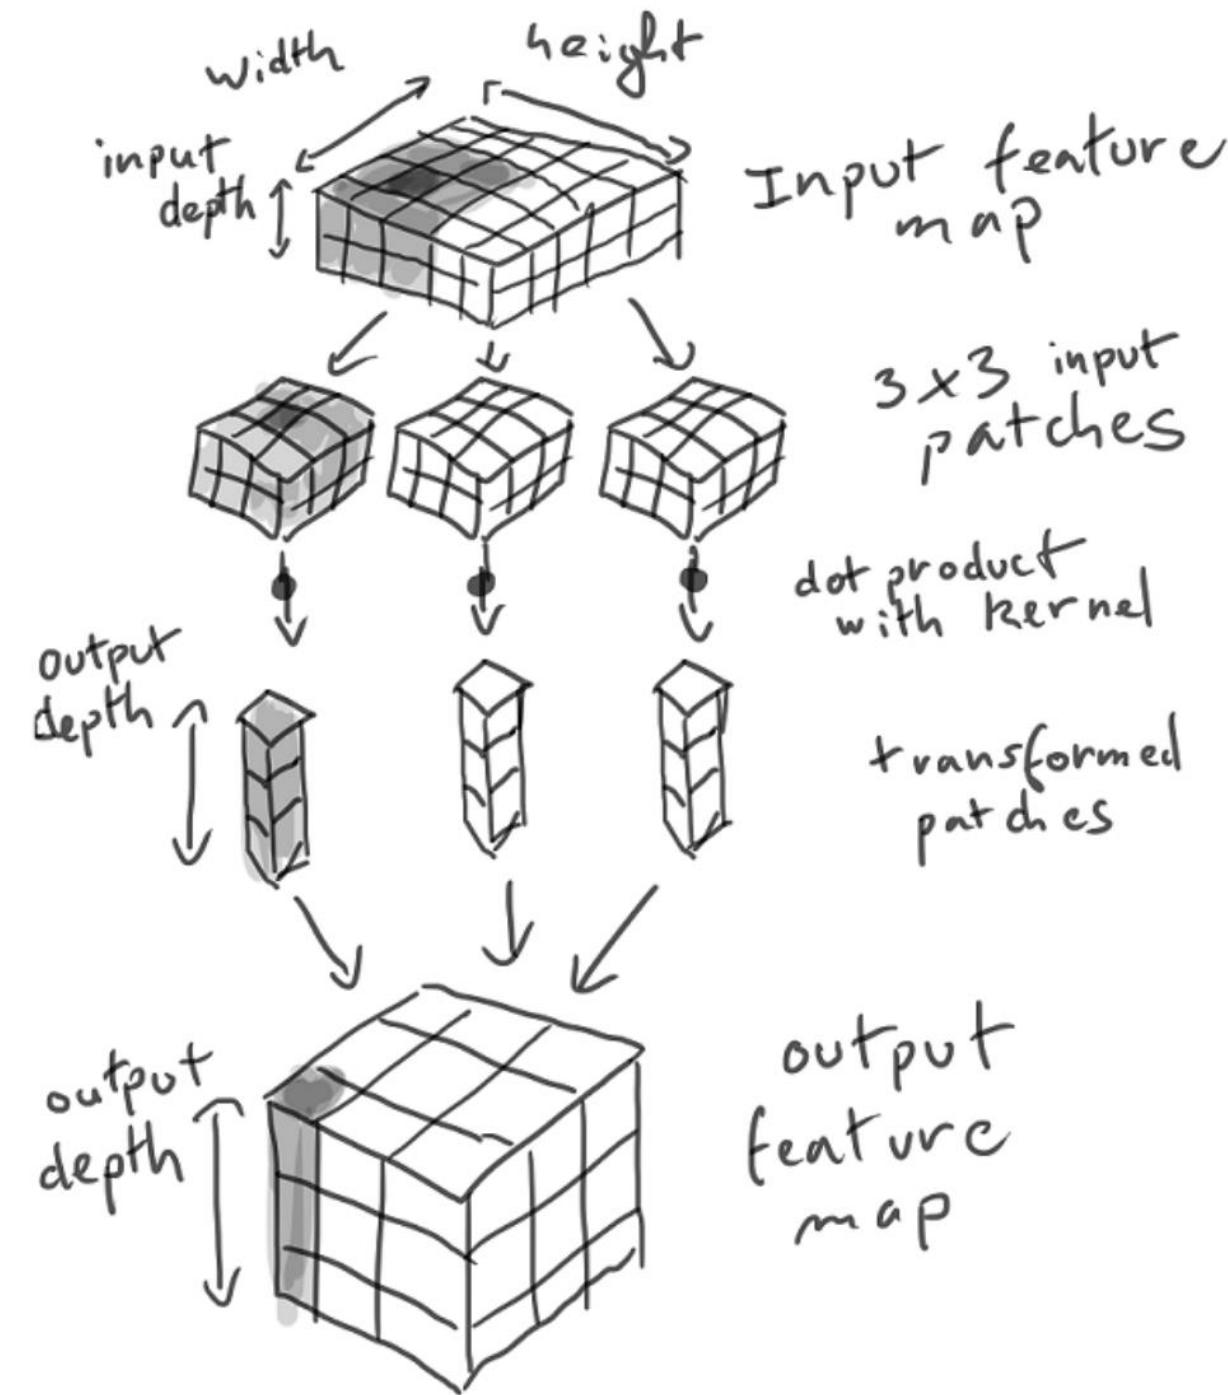
© Francois Chollet



STRIDING

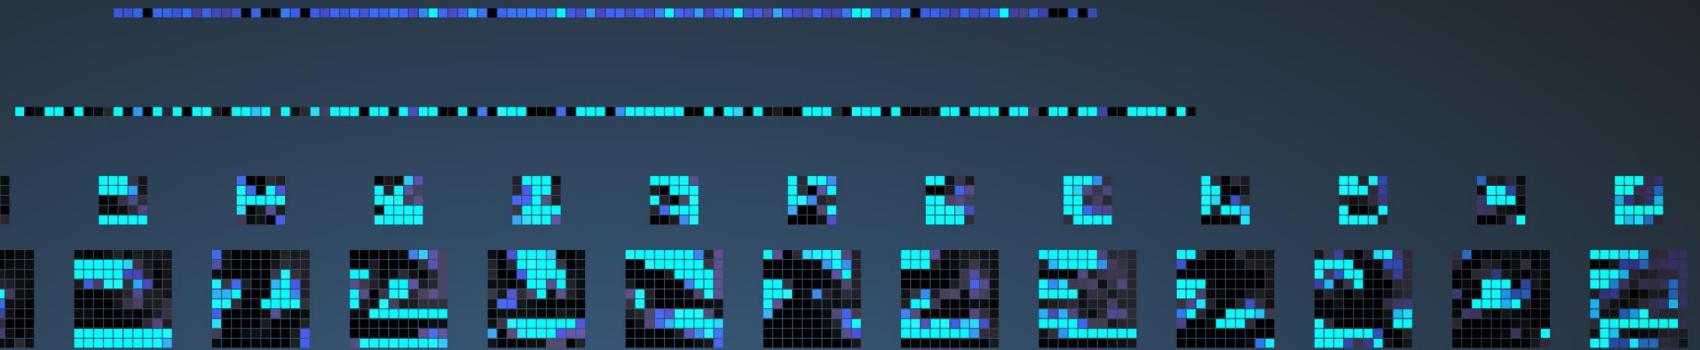
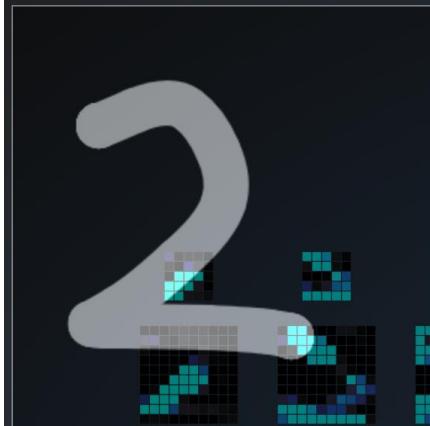
IMAGE CLASSIFICATION EXAMPLE





Draw your number here

0 1 2 3 4 5 6 7 8 9



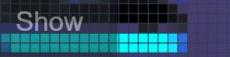
Downsampled drawing:

First guess:

Second guess:

Layer visibility

Input layer



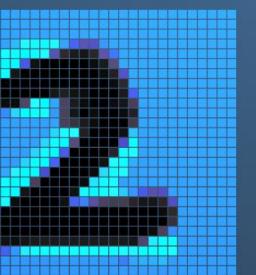
Show

Convolution layer 1



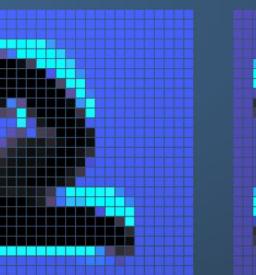
Show

Downsampling layer 1



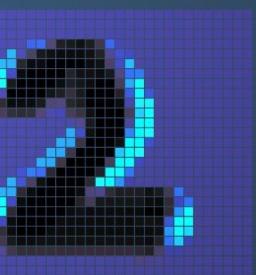
Show

Convolution layer 2



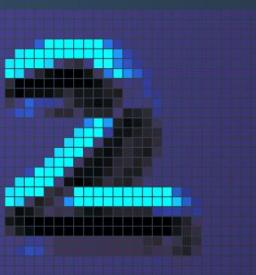
Show

Downsampling layer 2



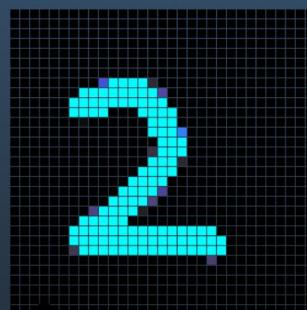
Show

Fully-connected layer 1



Show

Fully-connected layer 2



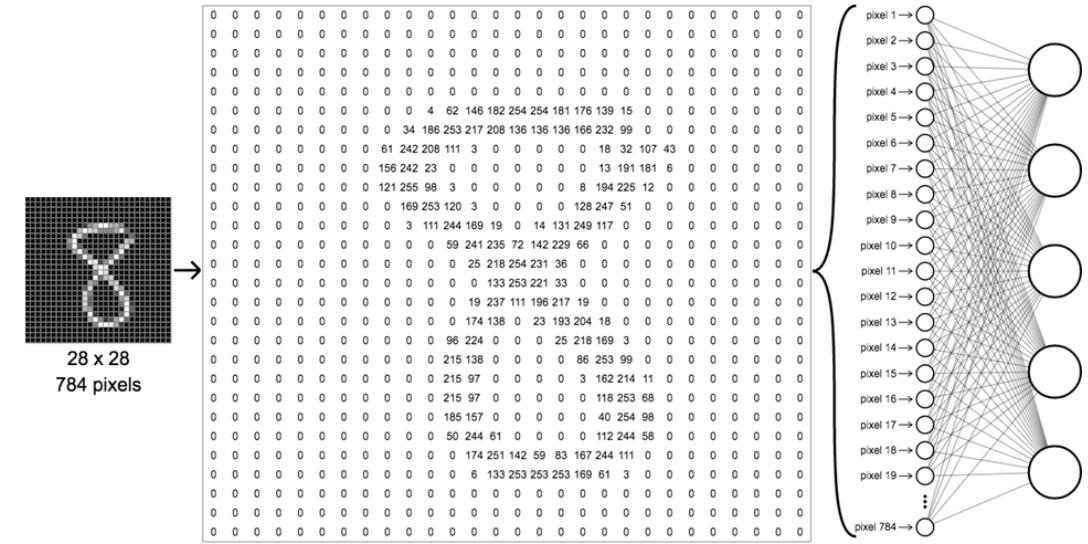
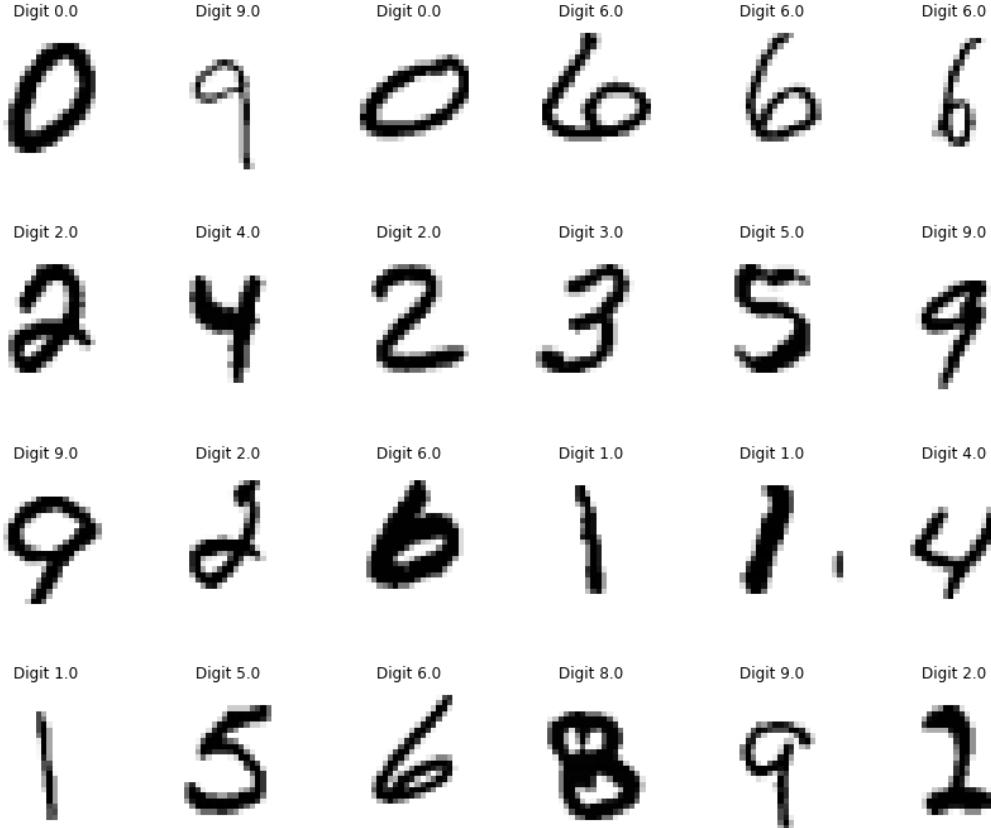
Show

Output layer

Show

<http://scs.ryerson.ca/~aharley/vis/conv>

MNIST DIGIT CLASSIFICATION



Represent input image as a vector $\mathbf{x} \in \mathbb{R}^{784}$
 Learn a classifier $f(\mathbf{x})$ such that,

$$f : \mathbf{x} \rightarrow \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$$

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model.summary()
```

input_shape=(28, 28, 1))

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D) Conv2D(32, (3, 3))	(None, 26, 26, 32)	320 $(32 * 3 * 3) + 32 = 320$
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
=====		
conv2d_2 (Conv2D) Conv2D(64, (3, 3),	(None, 11, 11, 64)	18496 $(3 * 3 * 32 + 1) * 64$
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
=====		
conv2d_3 (Conv2D) Conv2D(64, (3, 3))	(None, 3, 3, 64)	36928 $(3 * 3 * 64 + 1) * 64$
=====		

Total params: 55,744

CNN # params independent of input size

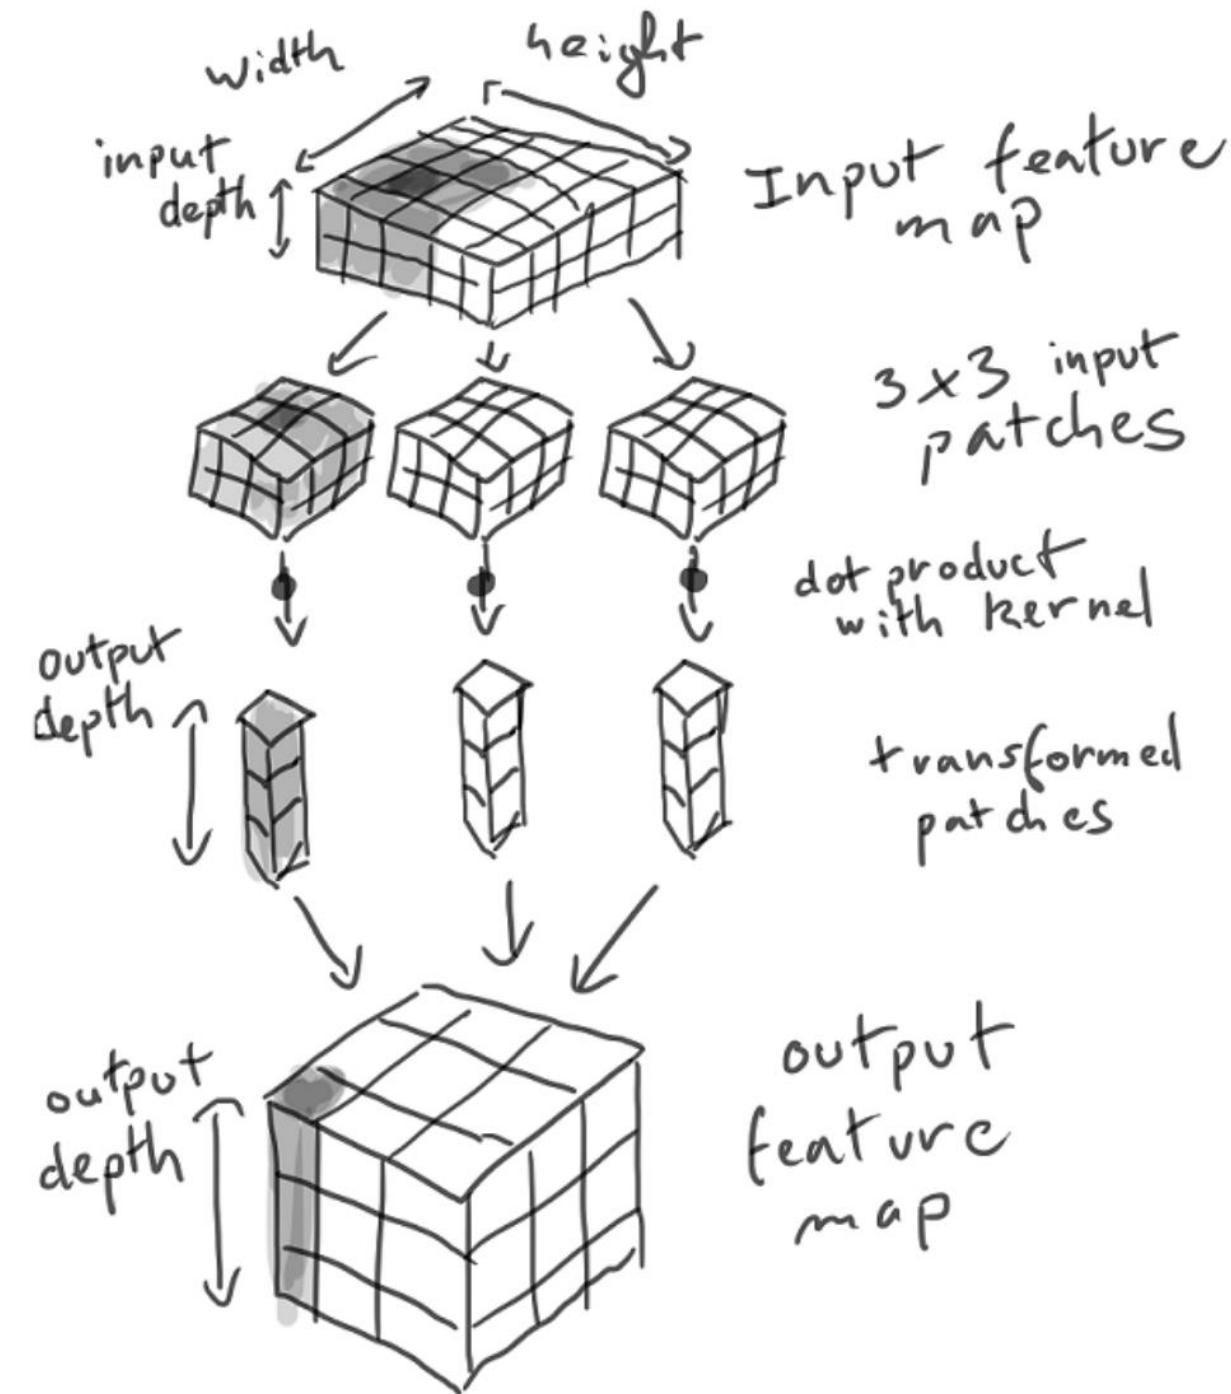
Trainable params: 55,744

total_params =

Non-trainable params: 0

$(\text{filter_height} * \text{filter_width} * \text{input_image_channels} + 1) * \text{number_of_filters}$

Remember this?



EACH LAYER DOES THIS

© Francois Chollet

```
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650
<hr/>		
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

```
from keras.datasets import mnist
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

"one hot encoding"

3 = 0 0 1
2 = 0 1 0

```
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5, batch_size=64)
```

Epoch 1/5

60000/60000 [=====] - 8s - loss: 0.1766 - acc: 0.9440

Epoch 2/5

60000/60000 [=====] - 7s - loss: 0.0462 - acc: 0.9855

Epoch 3/5

60000/60000 [=====] - 7s - loss: 0.0322 - acc: 0.9902

Epoch 4/5

60000/60000 [=====] - 7s - loss: 0.0241 - acc: 0.9926

Let's evaluate the model on the test data:

```
In [8]: test_loss, test_acc = model.evaluate(test_images, test_labels)
```

```
9536/10000 [=====>..] - ETA: 0s
```

```
In [9]: test_acc
```

```
Out[9]: 0.9912999999999996
```

DOGS AND CATS DATASET



- 2000 Dogs
- 2000 Cats
- We use 2K for training

© Francois Chollet

VISION ON SMALL DATA SETS

- “Deep learning only works on small datasets”
 - Only valid if you want to learn representations (i.e. “features”) as part of the training process
- Data Augmentation helps
- Transfer learning helps even more

DOGS CATS MODEL

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

```
: from keras import optimizers  
  
model.compile(loss='binary_crossentropy',  
              optimizer=optimizers.RMSprop(lr=1e-4),  
              metrics=['acc'])
```

```
from keras.preprocessing.image import ImageDataGenerator

# All images will be rescaled by 1./255
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    # This is the target directory
    train_dir,
    # All images will be resized to 150x150
    target_size=(150, 150),
    batch_size=20,
    # Since we use binary_crossentropy loss, we need binary labels
    class_mode='binary')

validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
```

Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50)
```

Epoch 1/30

100/100 [=====] - 9s - loss: 0.6898 - acc: 0.5285 - val_loss: 0.6724 - val_acc: 0.5950

Epoch 2/30

100/100 [=====] - 8s - loss: 0.6543 - acc: 0.6340 - val_loss: 0.6565 - val_acc: 0.5950

Epoch 3/30

100/100 [=====] - 8s - loss: 0.6143 - acc: 0.6690 - val_loss: 0.6116 - val_acc: 0.6650

Epoch 4/30

100/100 [=====] - 8s - loss: 0.5626 - acc: 0.7125 - val_loss: 0.5774 - val_acc: 0.6970

It is good practice to always save your models after training:

```
: model.save('cats_and_dogs_small_1.h5')
```

```
: import matplotlib.pyplot as plt

acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

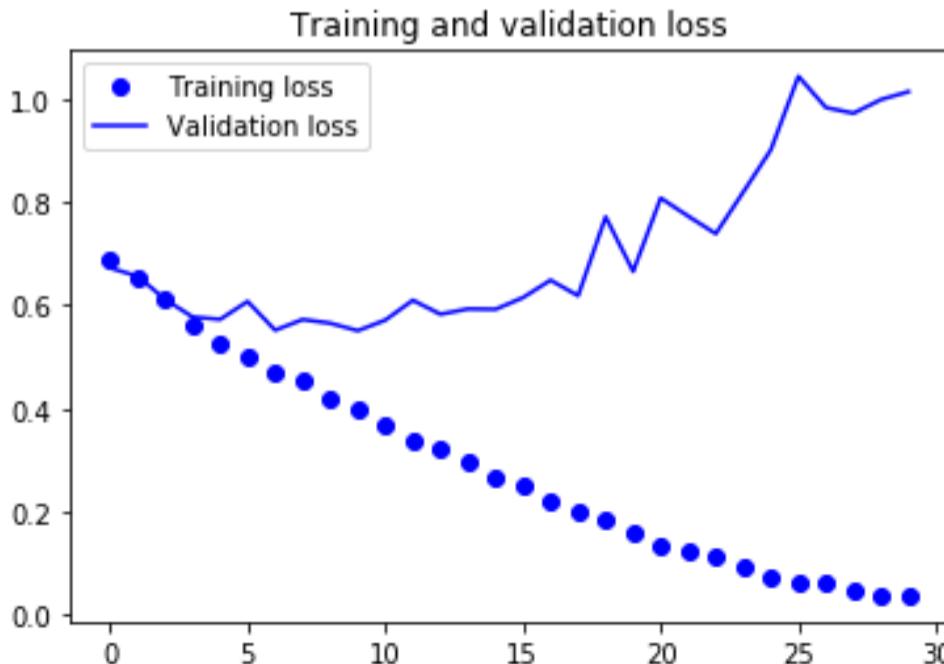
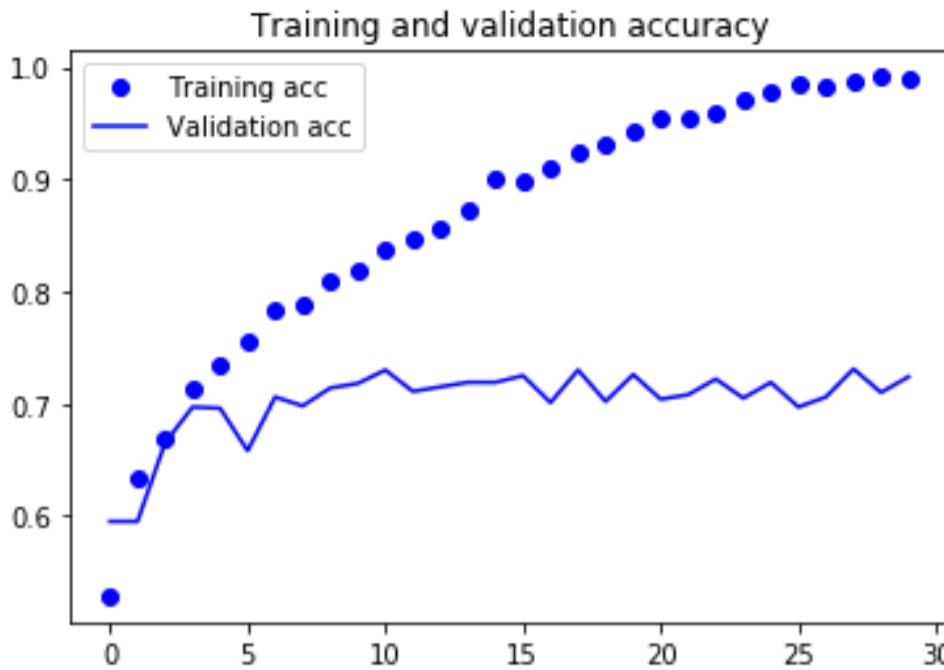
epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

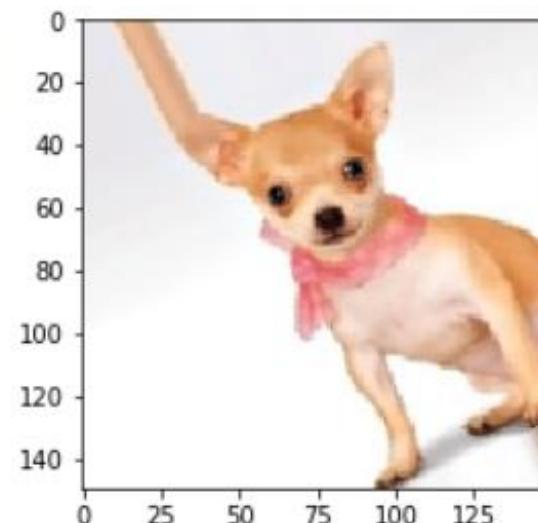
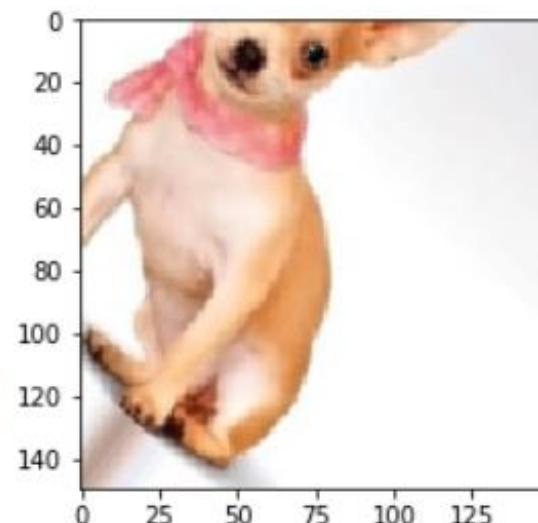
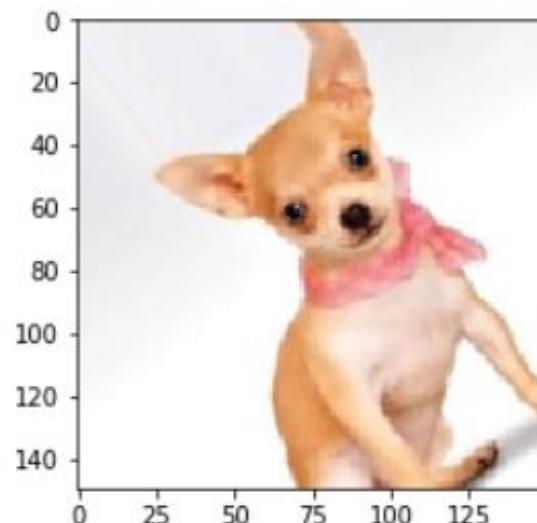
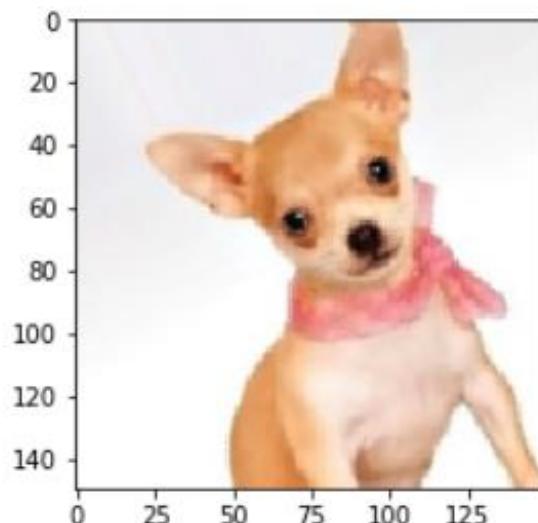
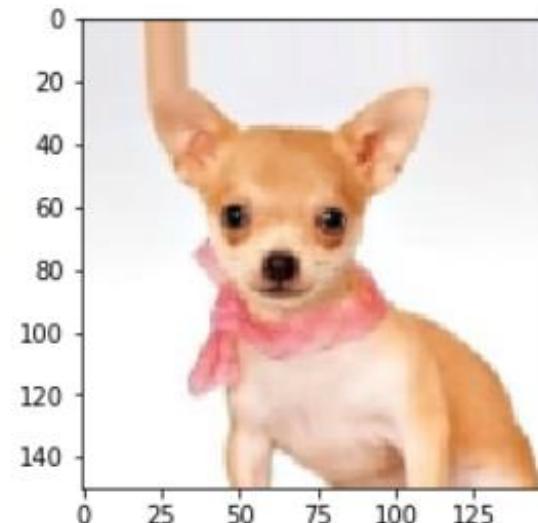
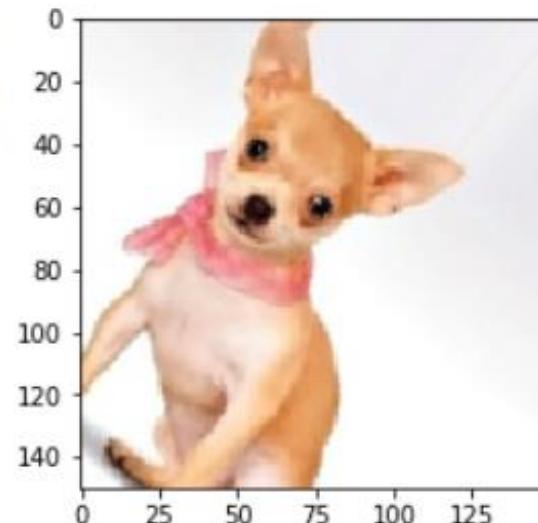
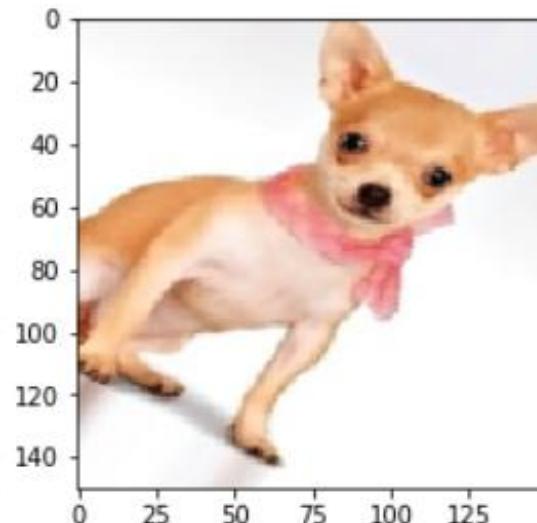
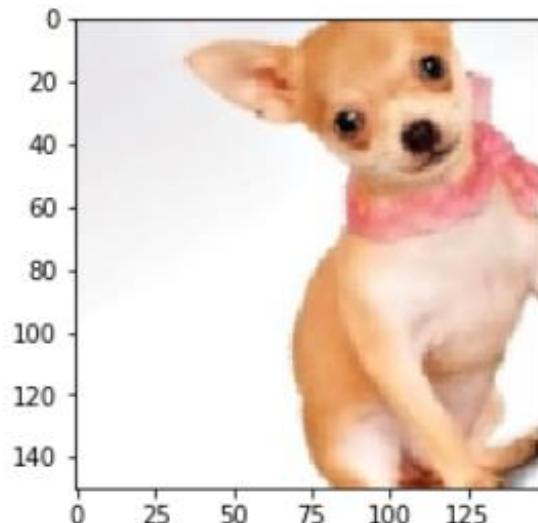
plt.show()
```



- Val acc of 70%
- Overfitting
- Dropout
- L2 regularisation
- Now... Augmentation

Overfitting is caused by having too few samples to learn from, rendering us unable to train a model able to generalize to new data

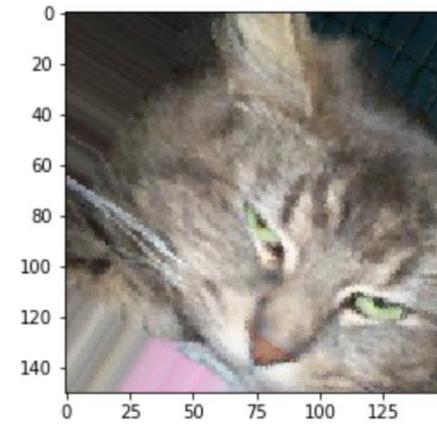
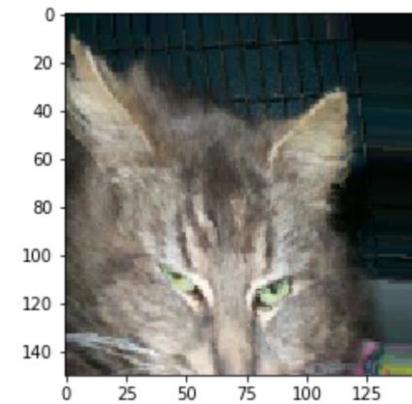
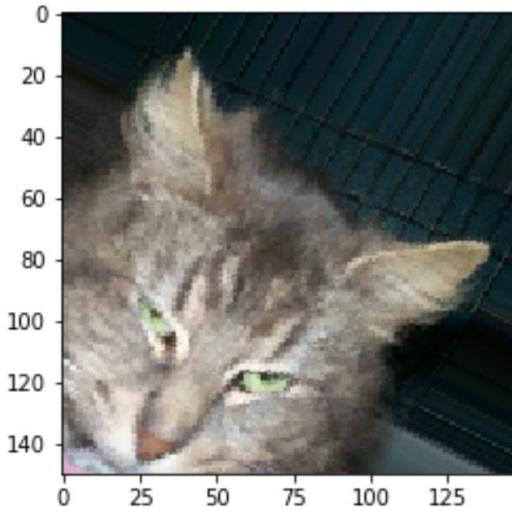
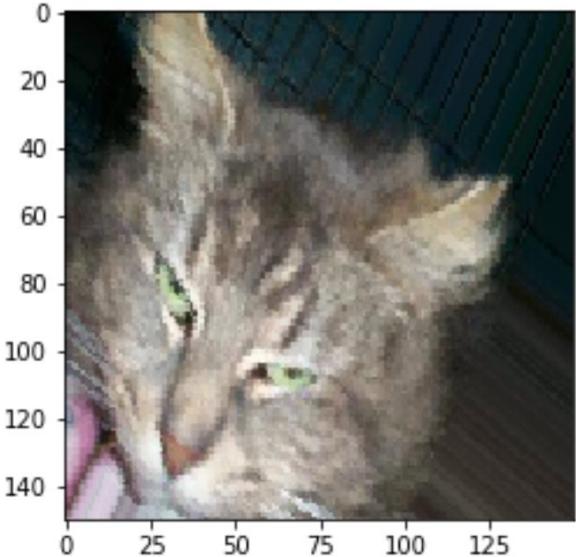
[DATA AUGMENTATION]



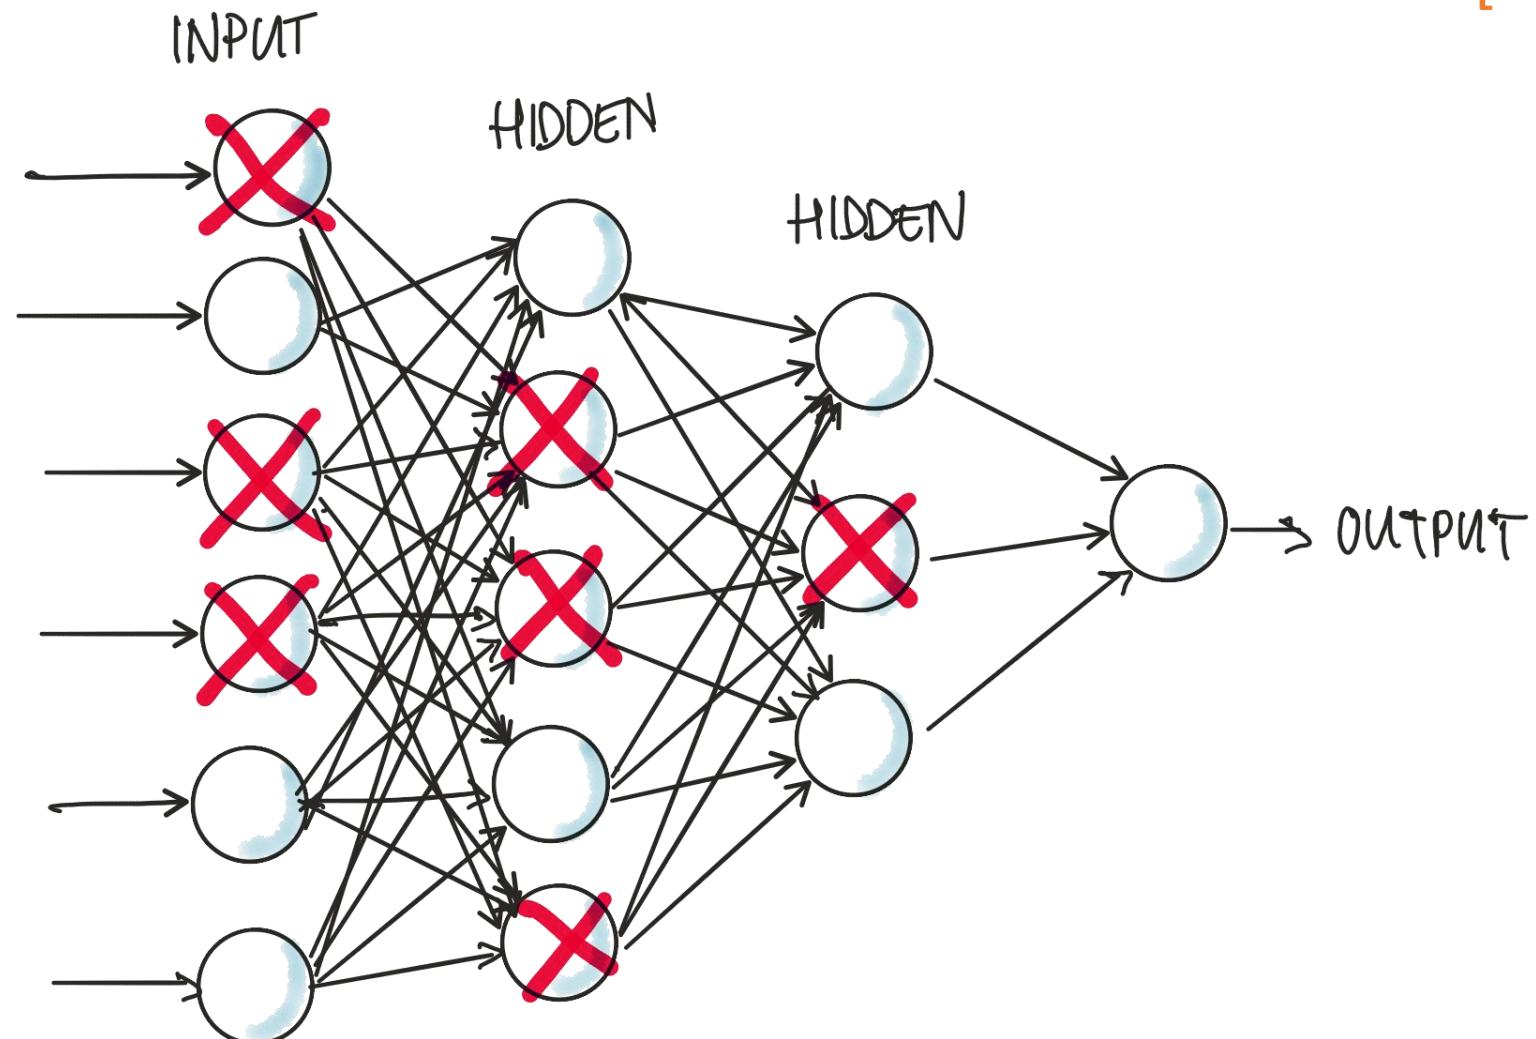
Chihuahua the movie

```
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```

AUGMENTED IMAGES



[DROPOUT]



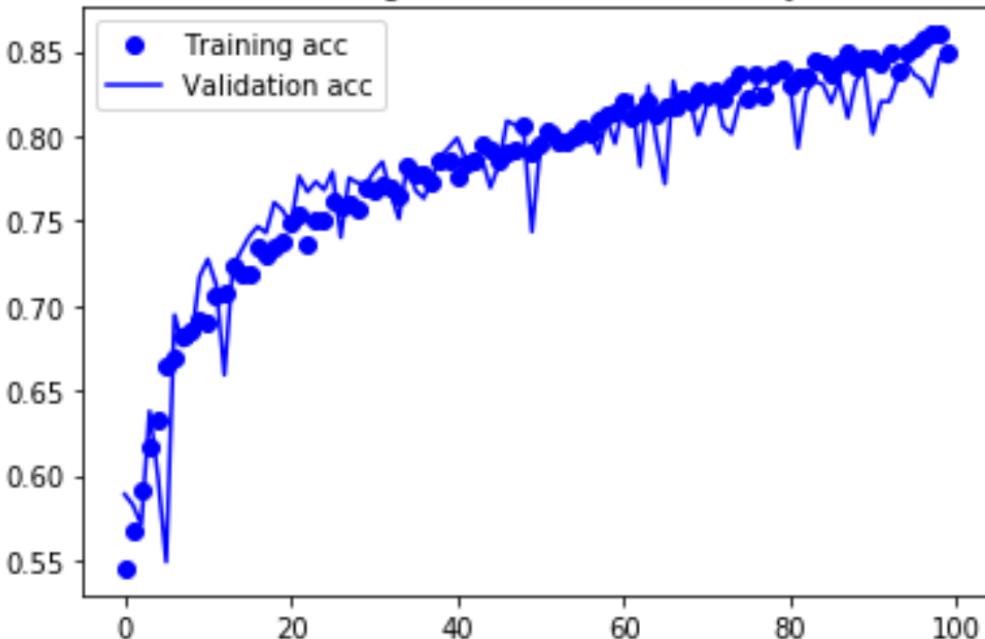
TRAIN AGAIN WITH DROPOUT AND AUGMENTATION

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```

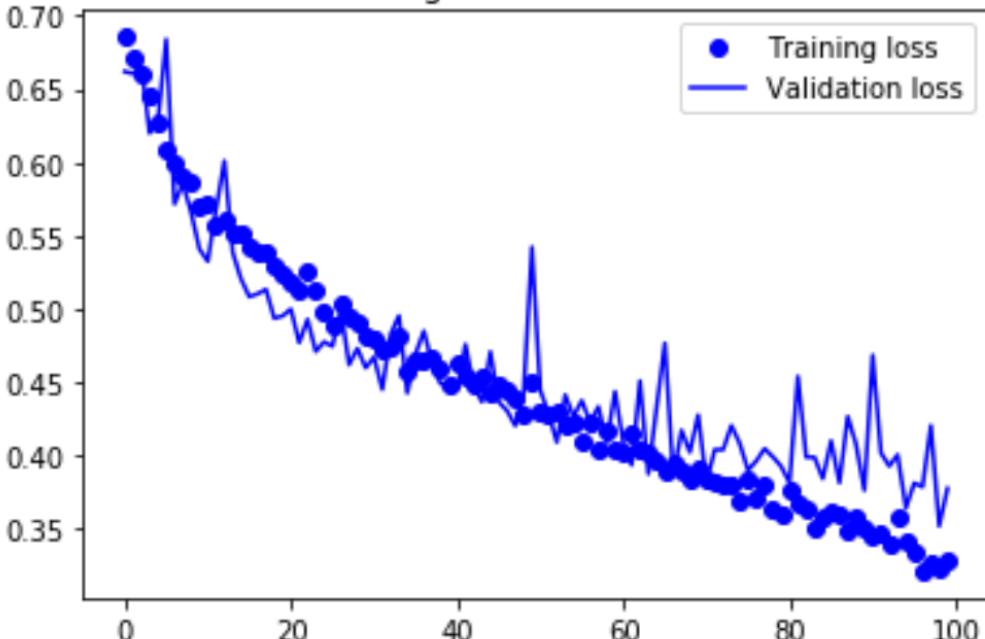
NO OVERFITTING!

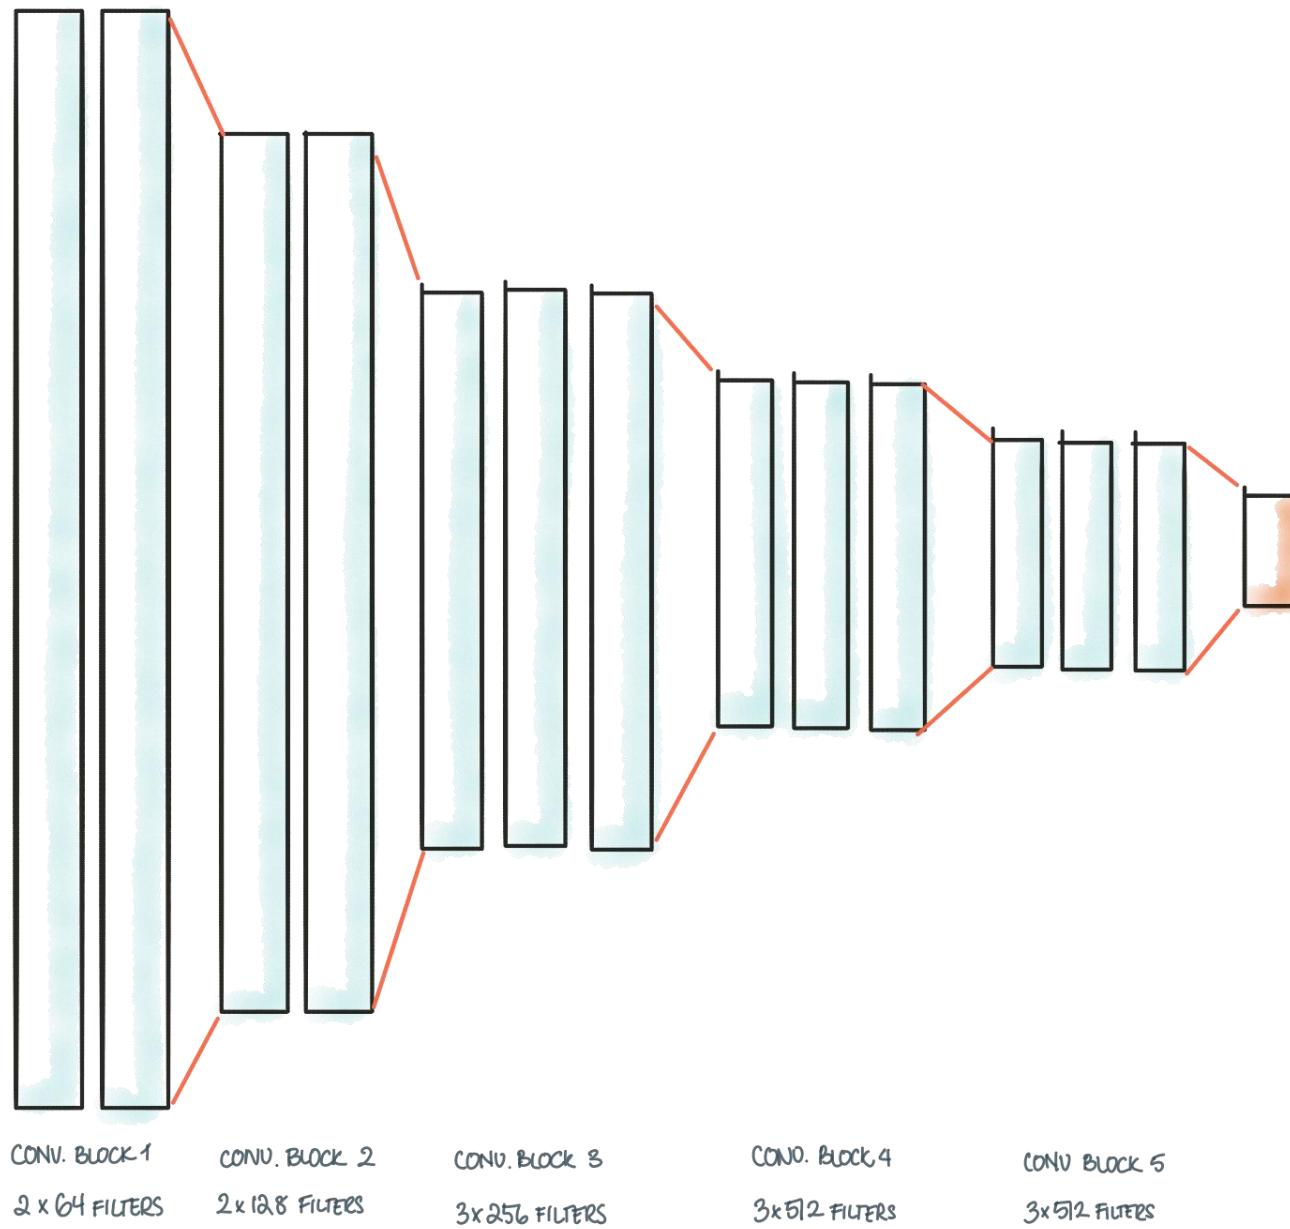
Training and validation accuracy

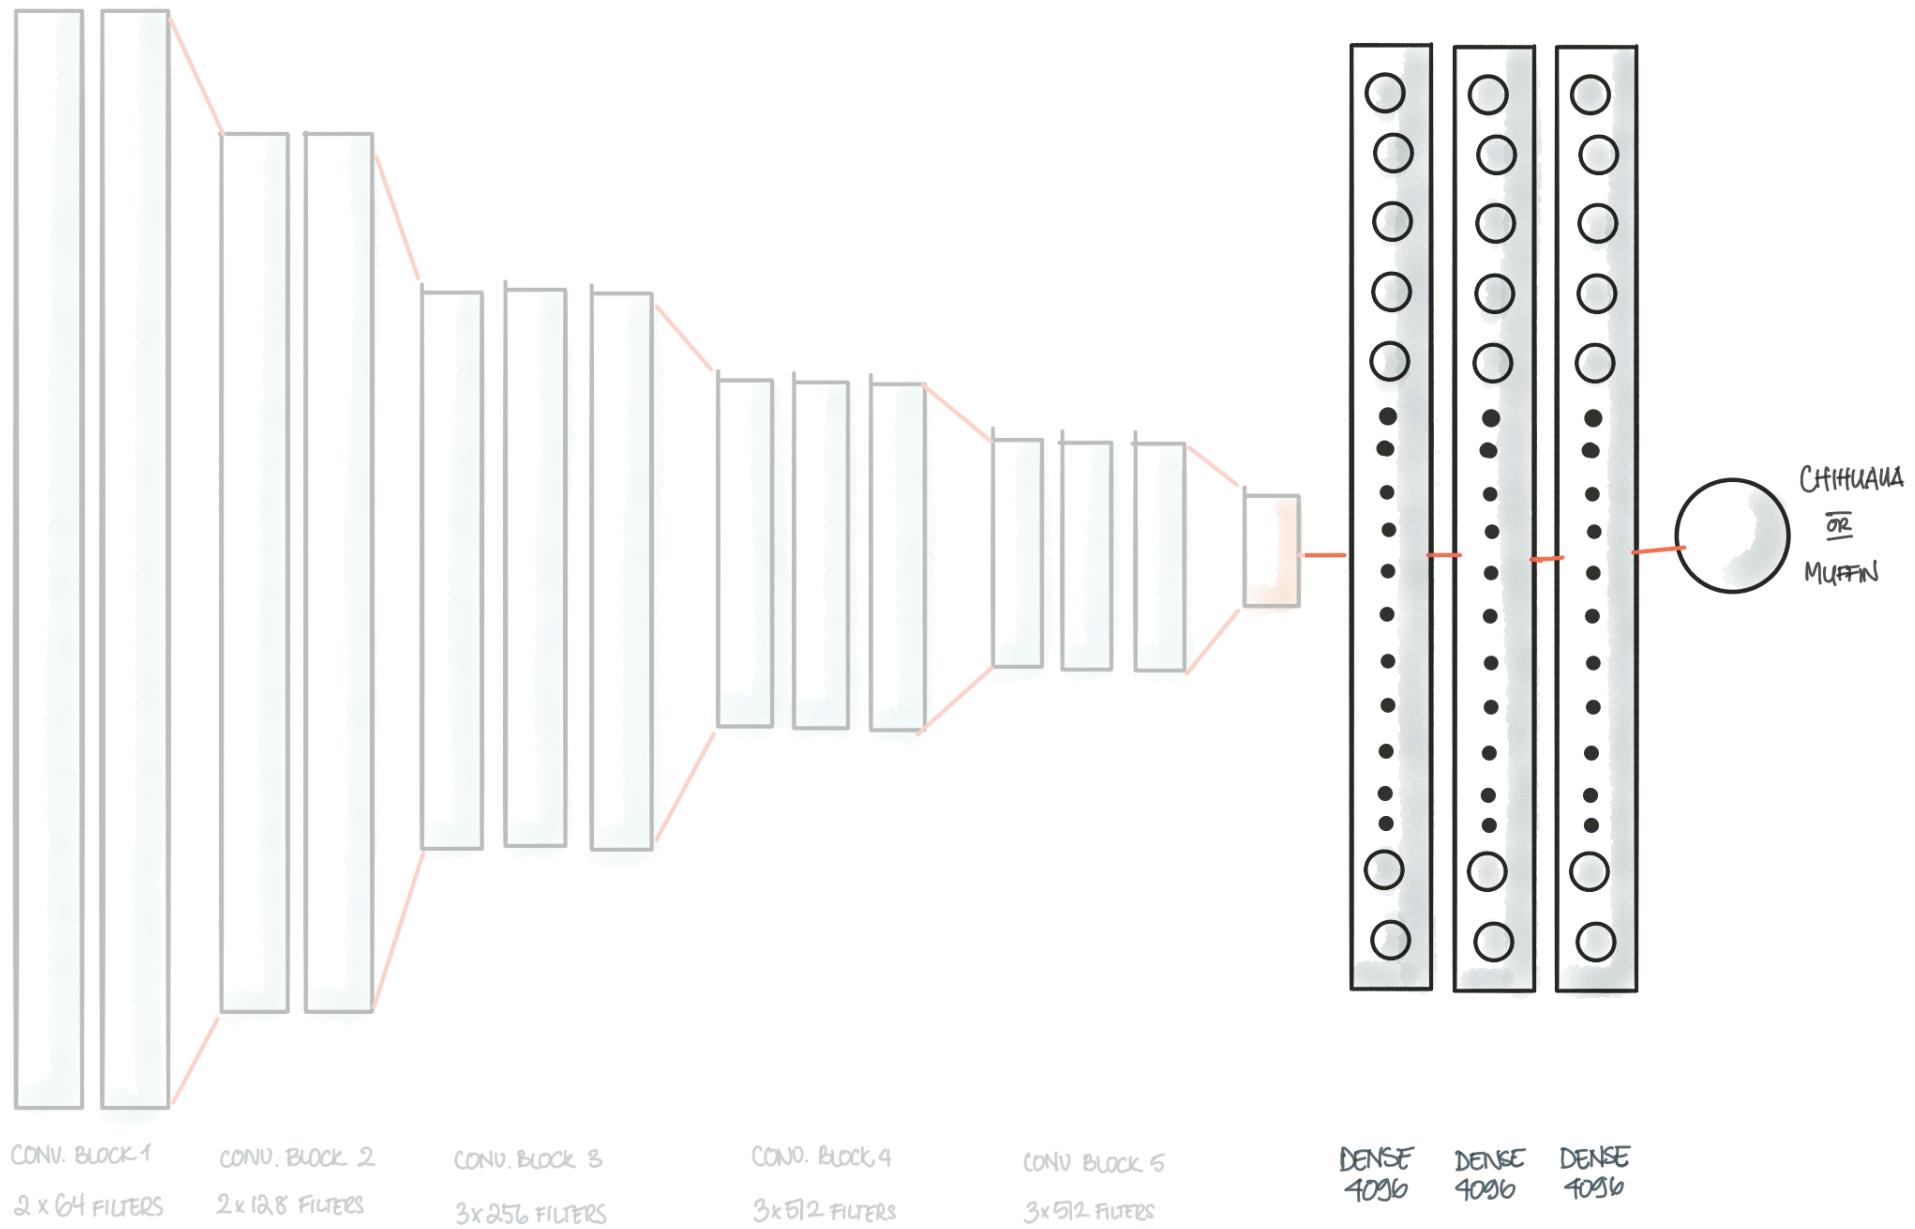


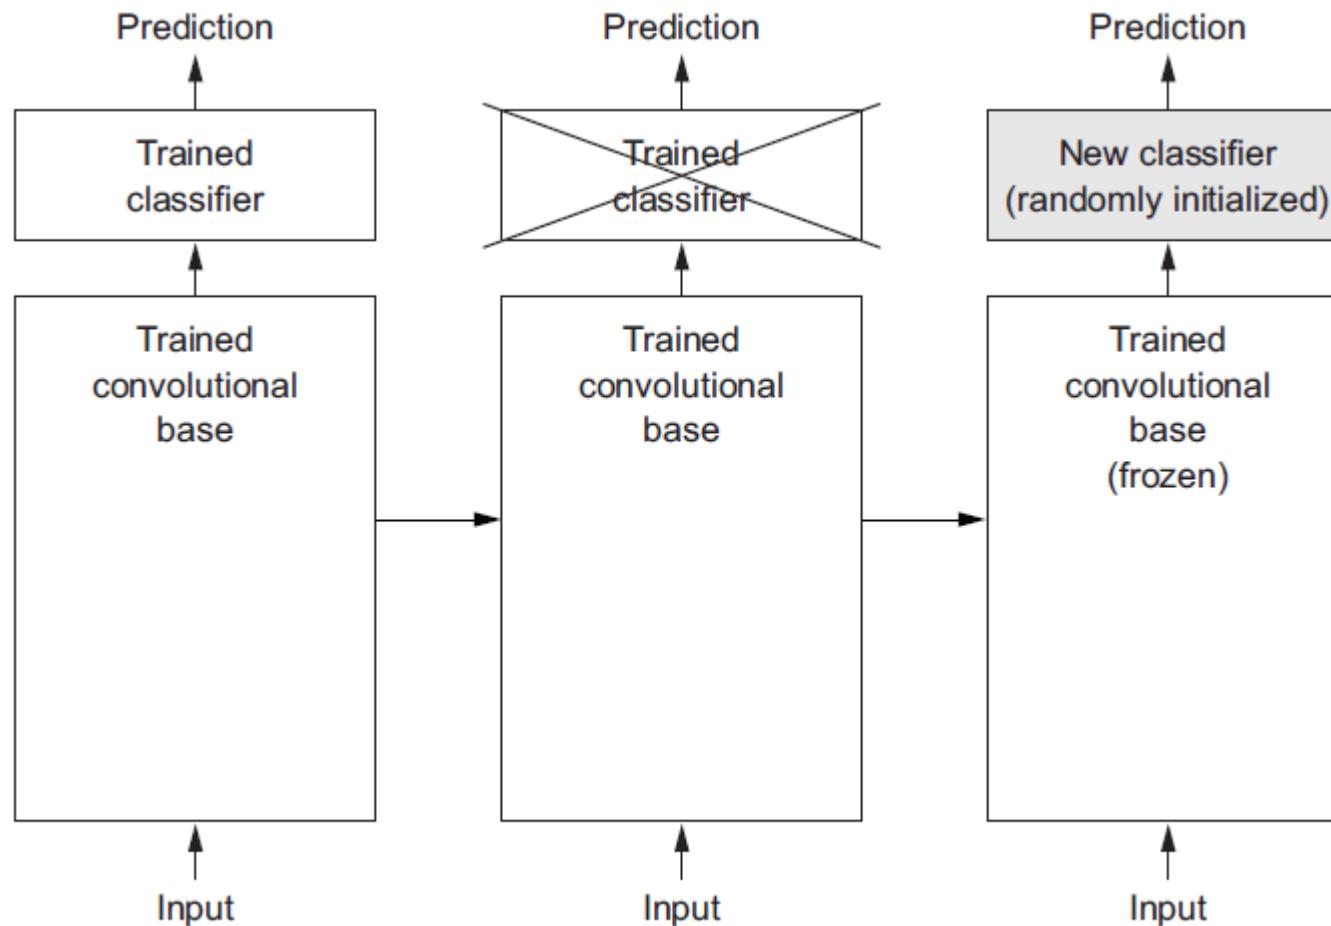
- VAL acc of 82!
- 15% relative

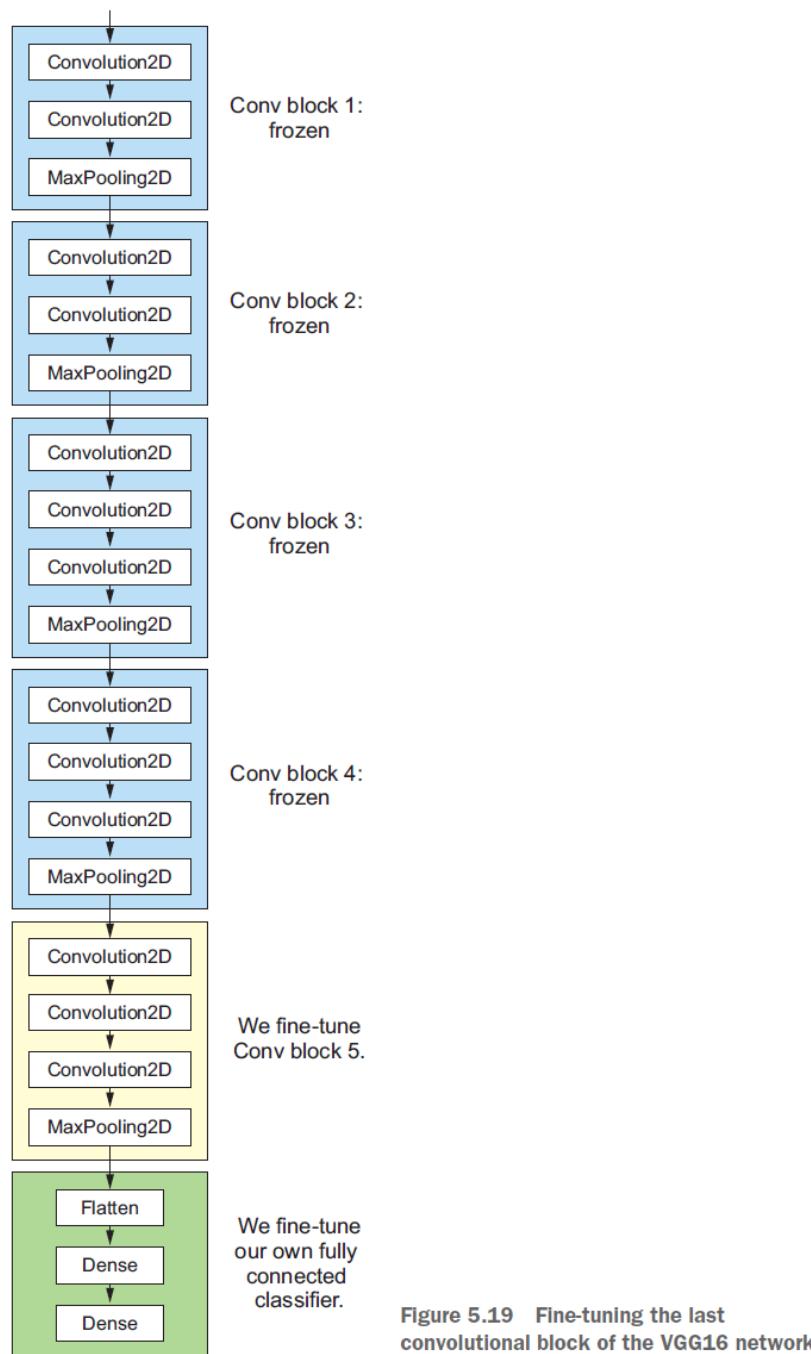
Training and validation loss











- Xception
- InceptionV3
- ResNet50
- VGG16
- VGG19
- MobileNet

Figure 5.19 Fine-tuning the last convolutional block of the VGG16 network

```
from keras.applications import VGG16

conv_base = VGG16(weights='imagenet',
                  include_top=False,
                  input_shape=(150, 150, 3))
```

block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

ADD THE CONV BASE TO MODEL

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [10]: print('This is the number of trainable weights '
      'before freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 30

```
In [11]: conv_base.trainable = False
```

```
In [12]: print('This is the number of trainable weights '
      'after freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 4

We don't train into the conv base until we have trained the classifier,
so we don't destroy the sensitive representations learned in the base

```
model.summary()
```

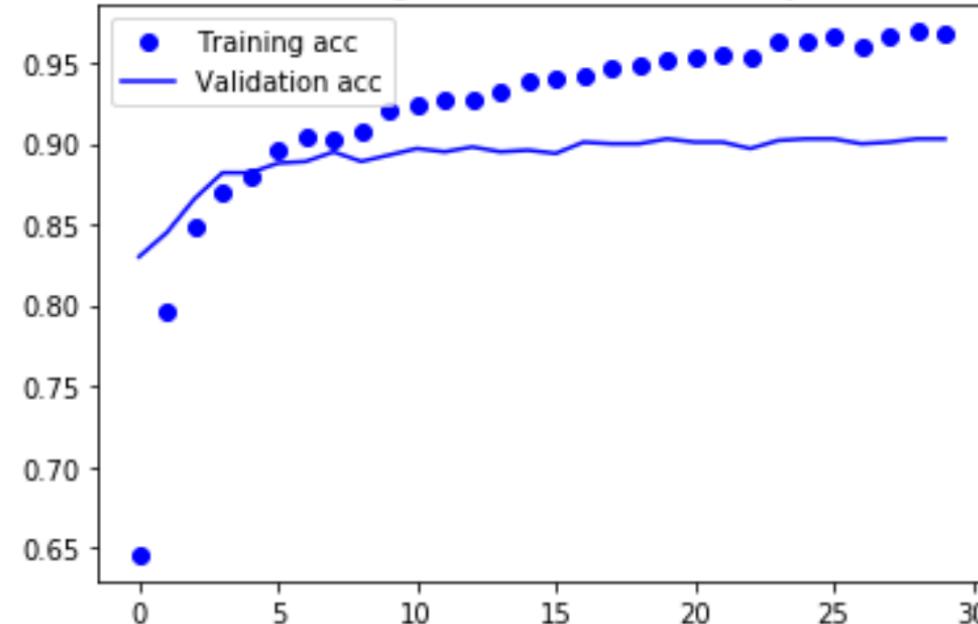
Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_3 (Dense)	(None, 256)	2097408
dense_4 (Dense)	(None, 1)	257

Total params: 16,812,353

Trainable params: 16,812,353

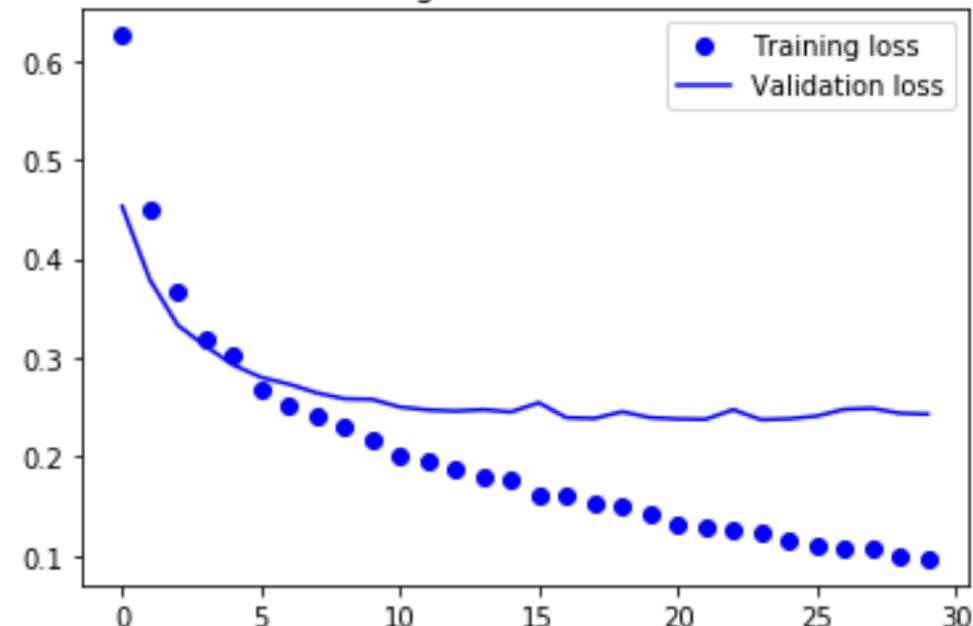
Non-trainable params: 0

Training and validation accuracy

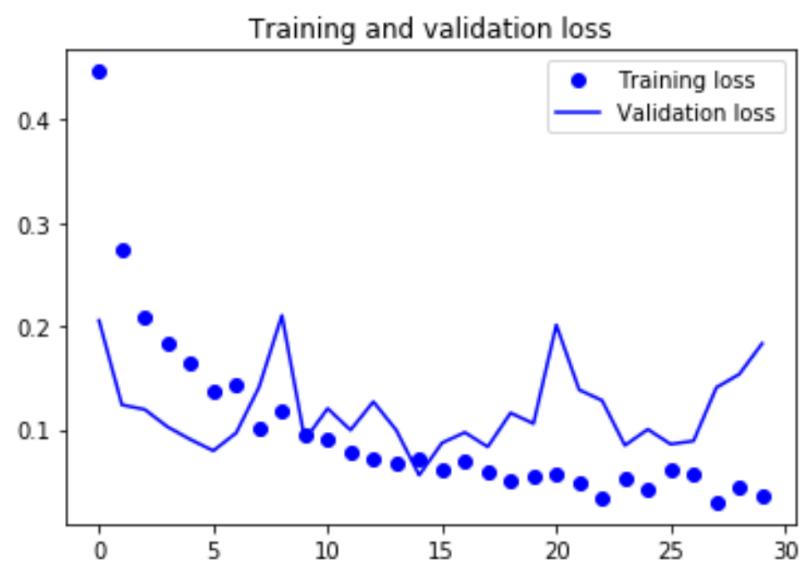
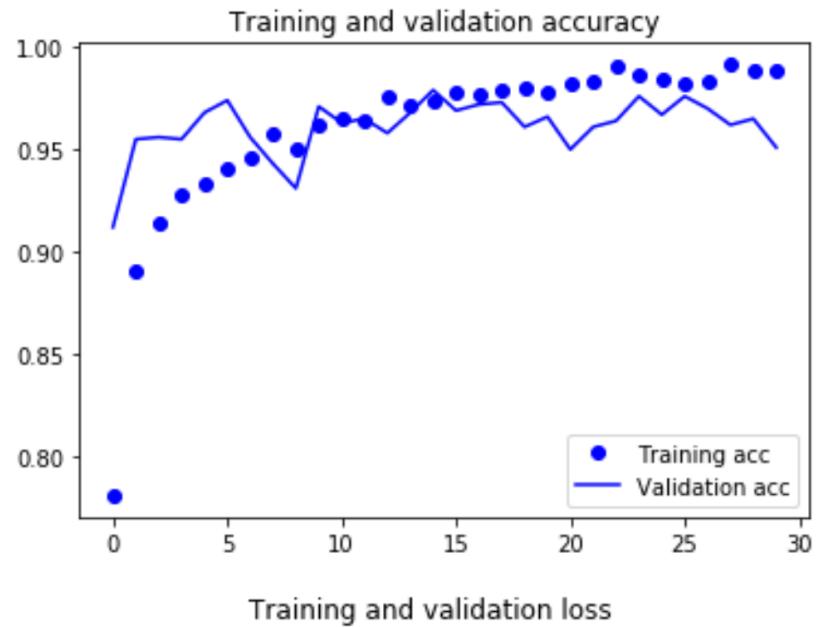


- Val now ~90%!
- Overfitting a lot
- Note this is with NO augmentation!

Training and validation loss



ADD THE AUGMENTATION BACK IN



- Val now ~96%!
- **Not overfitting**

RECAP ON VALIDATION PERFORMANCE

- CATS AND DOGS
 - Basic model 72%
 - Dropout and augmentation 82%
 - Transfer learning 90% (overfitting)
 - Transfer with augmentation: 96% (not overfitting)

TRAINING INTO THE CONV BASE

“FINE TUNING”

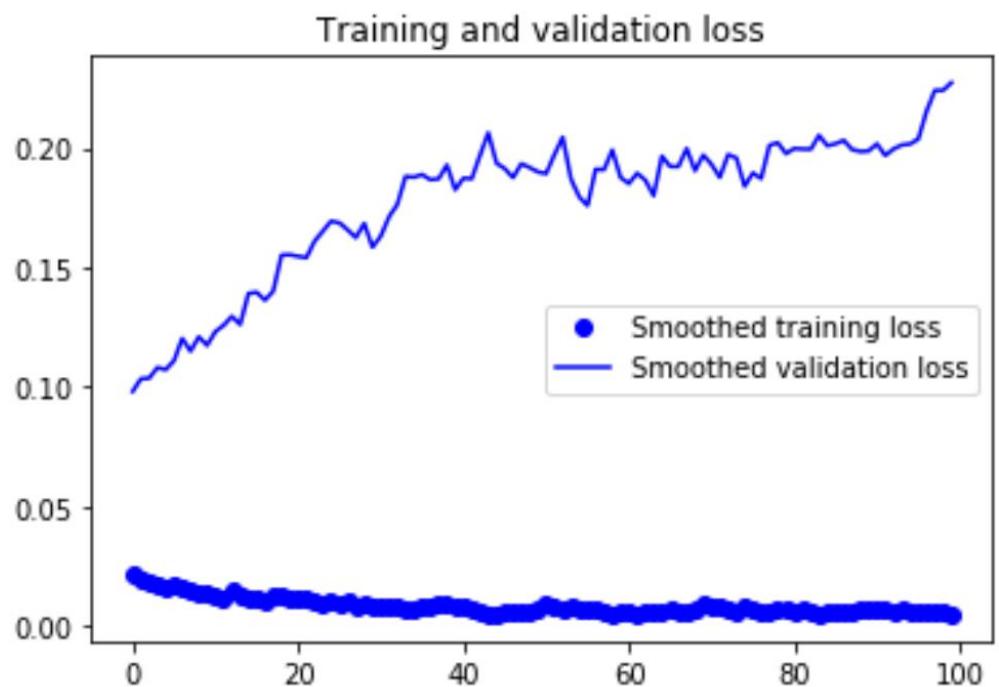
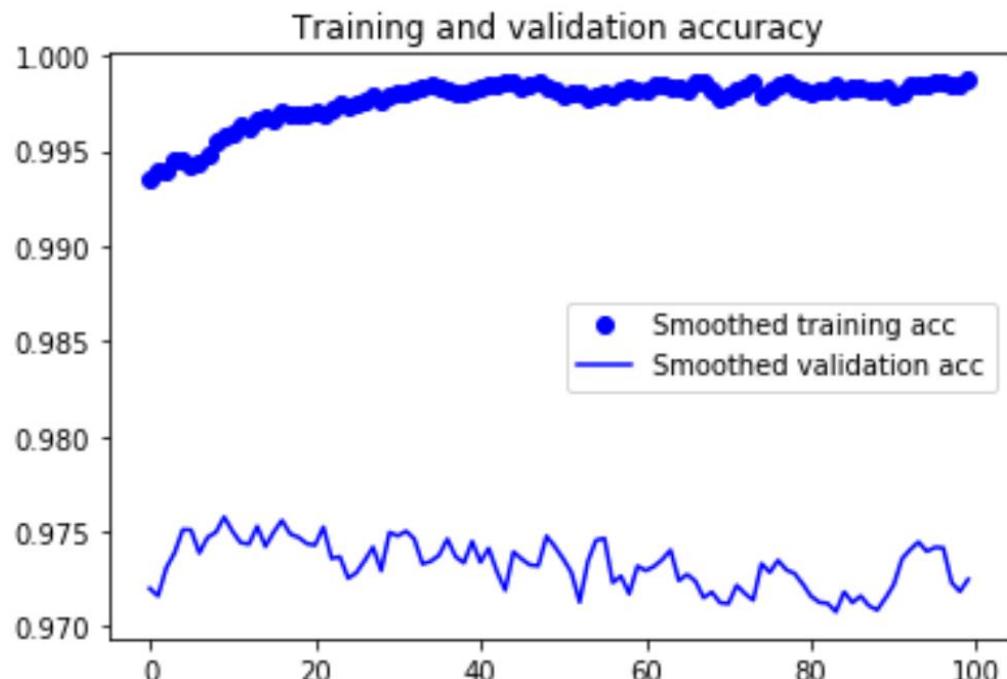
- Add your custom network on top of an already trained base network.
- Freeze the base network.
- Train the part you added.
- Unfreeze some layers in the base network.
- Jointly train both these layers and the part you added.

TRAINING INTO THE CONV BASE (2)

- Why not fine tune all layers?
 - More parameters == more overfitting
- Use a low learning rate == less chance of disaster
- How can loss get worse and accuracy go up?
 - Average of loss vs distribution of loss values might be different i.e. model might still be improving

```
conv_base.trainable = True

set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False
```



- 1% improvement
- 97% would have been a very good score on Kaggle in the day
- Done with only 2000 examples from the 4k!

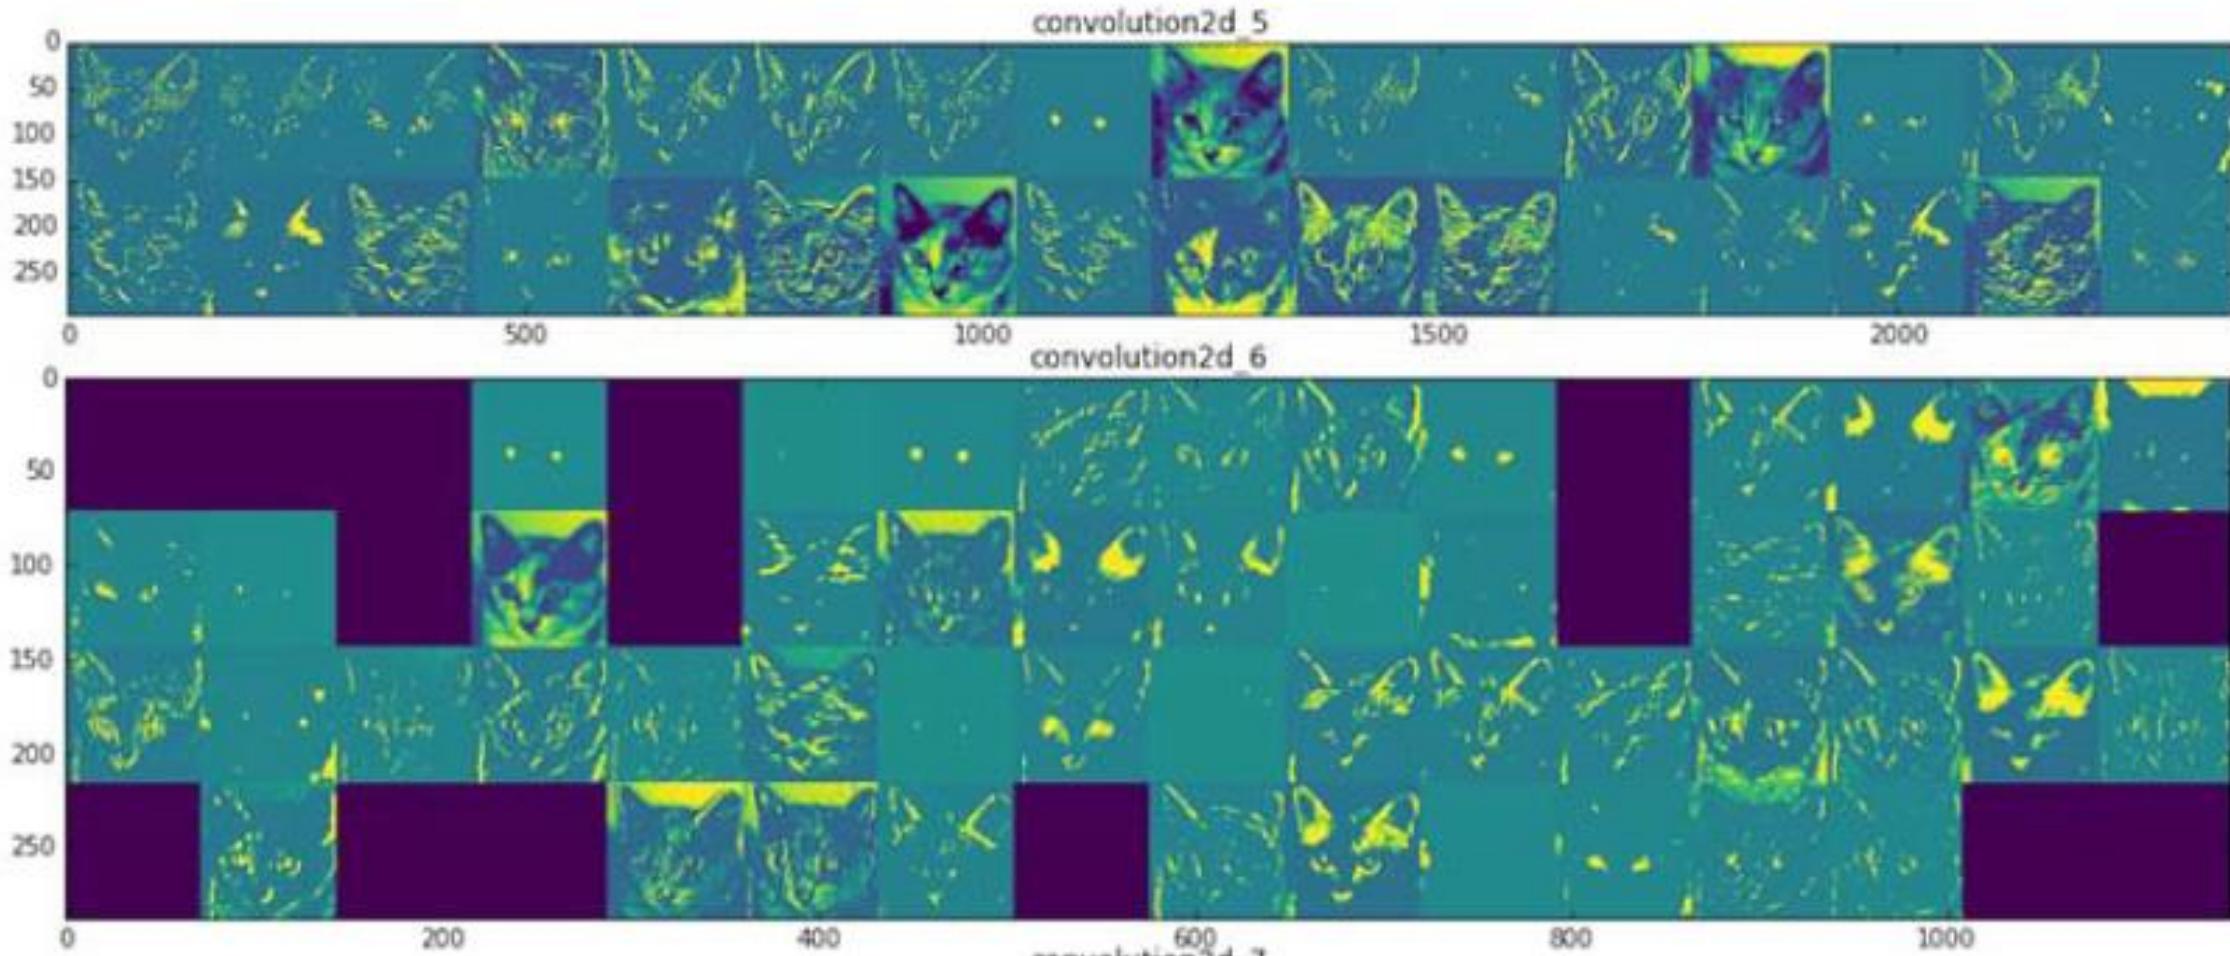
RECAP ON VALIDATION PERFORMANCE

- CATS AND DOGS
 - Basic model 72%
 - Dropout and augmentation 82%
 - Transfer learning 90% (overfitting)
 - Transfer with augmentation: 96% (not overfitting)
 - Training “into” the conv base top layers: 97% aka “fine tuning”

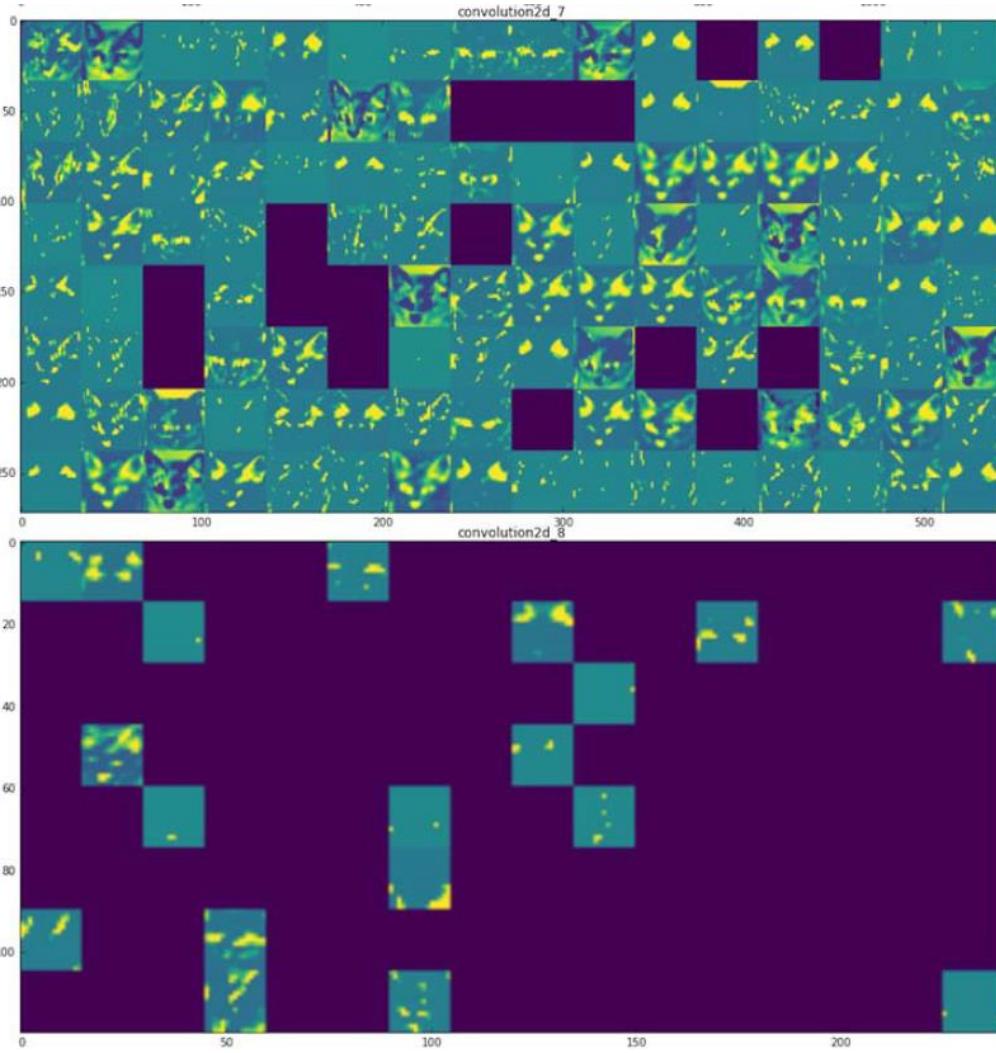
TAKE AWAYS

- Convnets best for vision tasks
- Can train from scratch even with small sets
- Overfitting is a huge issue, regularisation and augmentation helps
- Reuse existing convnets (transfer)
- Can train into conv base “fine tuning” but diminished returns

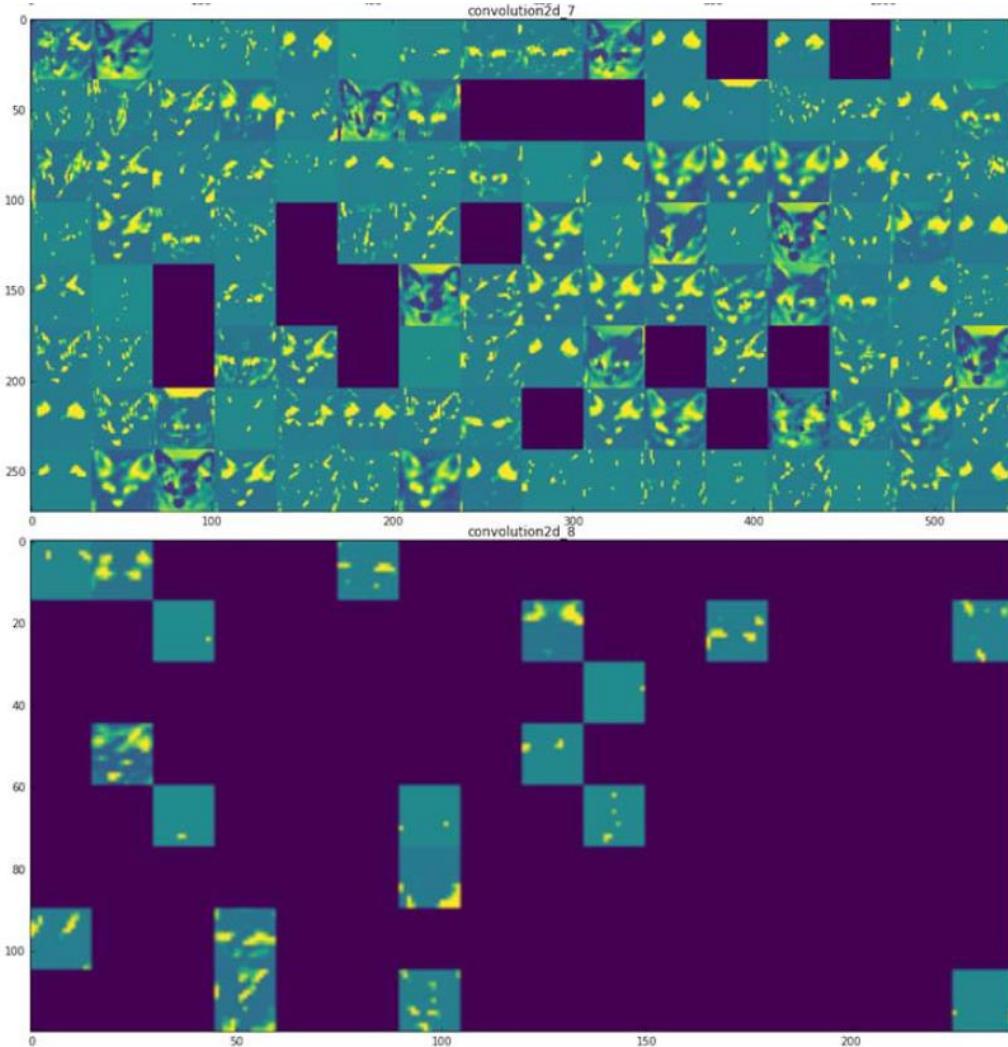
VISUALISING THE FILTERS (1)



VISUALISING THE FILTERS (2)



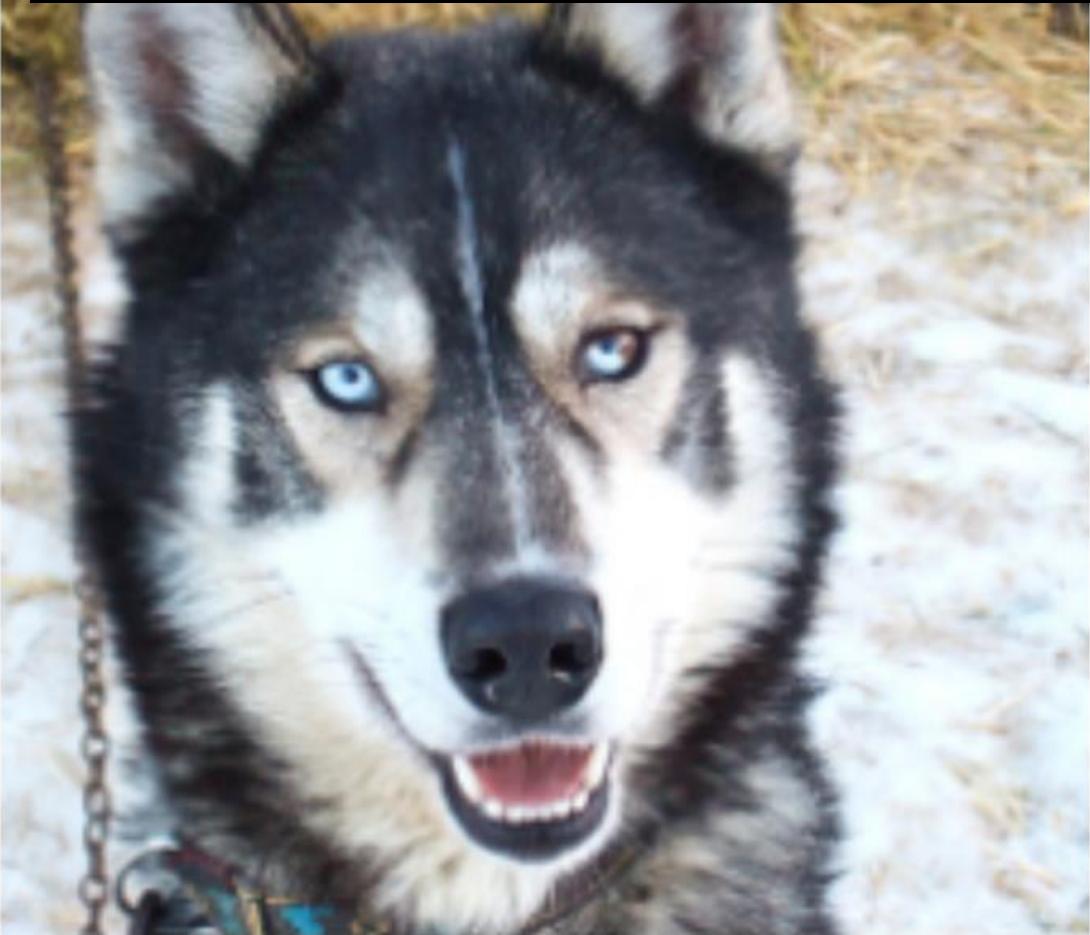
VISUALISING THE FILTERS (3)



- Note these are actually the response maps, not filters
 - Early layers; content
 - Later layers; class
 - Sparsity

INTERPRETABILITY - LIME/SHAP

[LIME]

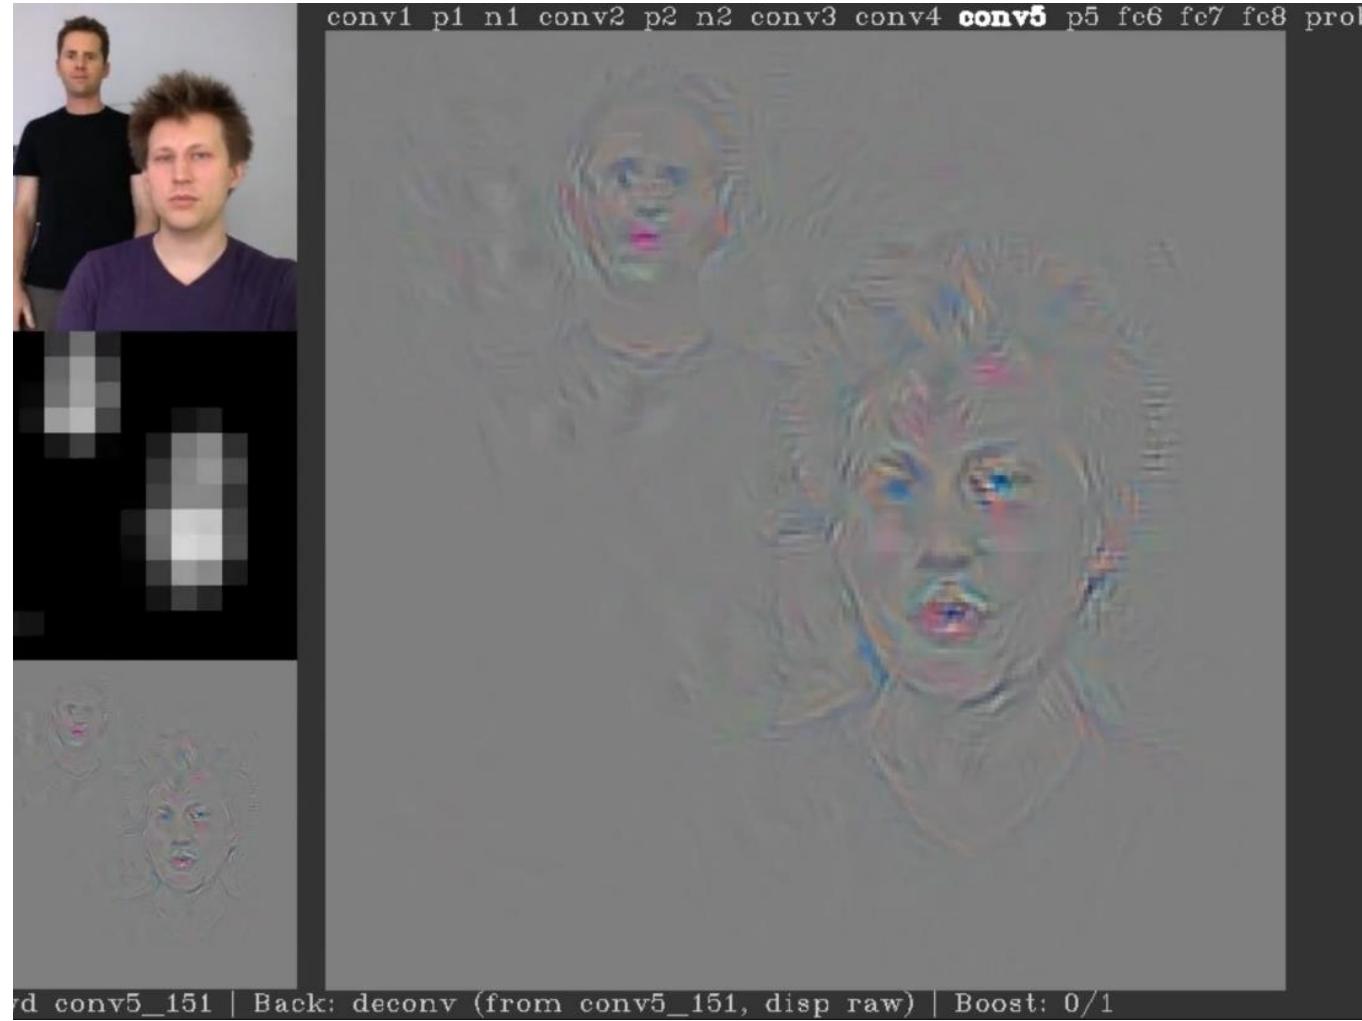


(a) Husky classified as wolf

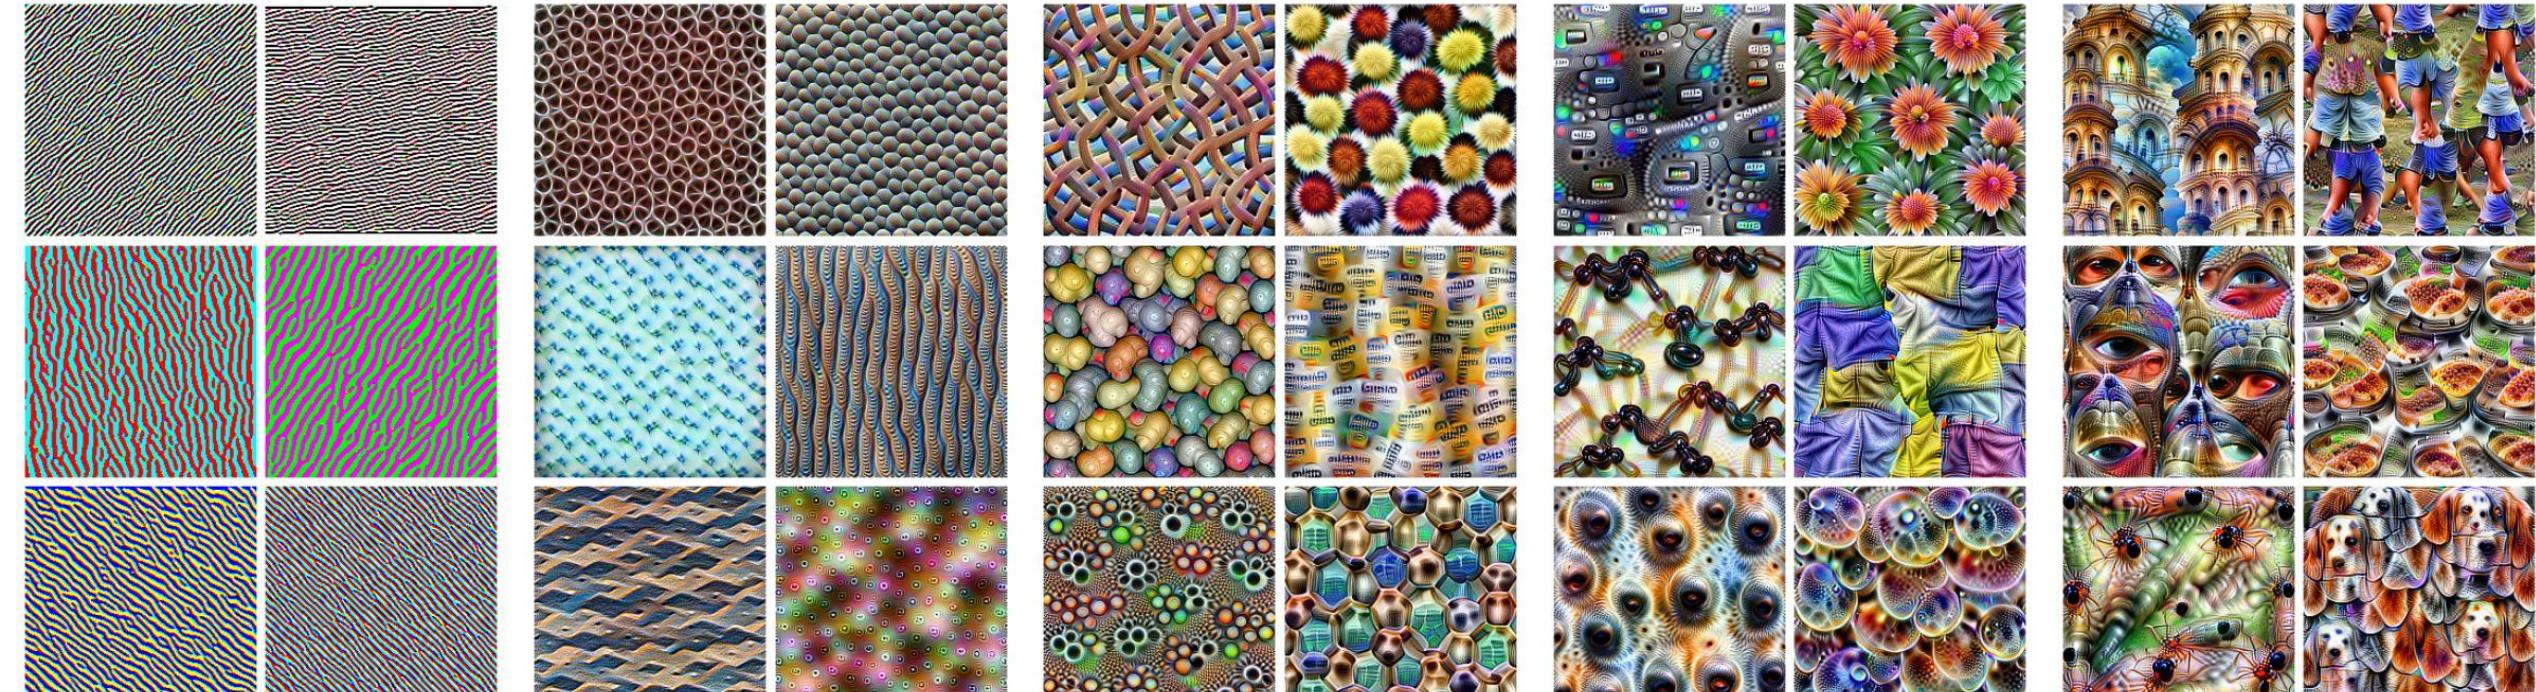


(b) Explanation

DEEP VISUALISATION TOOLBOX "SALIENCY MAPS"



FEATURE VISUALISATION



Edges (layer conv2d0)

Textures (layer mixed3a)

Patterns (layer mixed4a)

Parts (layers mixed4b & mixed4c)

Objects (layers mixed4d & mixed4e)

Tim Scarfe

@ecsquendor

youtube.com/machinelearningatmicrosoft



THANK YOU!

