

Microsoft Future Decoded

Productionising deep learning: The confused.com journey

Who are we?



Dr. Tim Scarfe

Machine learning engineer, data scientist and entrepreneur.

Principal Software Engineer

[@ecsquendor](#)



Dr. Tempest van Schaik

Biomedical engineer, inventor, designer, researcher. Health + AI.

Software Engineer in Data & AI

[@Dr_Tempest](#)

Dr Tim + Dr Tem aka:



Tweat out to [#TimTem](#) !



What you'll hear from us today:

- What we do at Microsoft
- The story of deep learning
- Machine learning is becoming ubiquitous in software
- Many companies are becoming data companies
- Our collaboration with confused.com
- Our deep learning model for car insurance
- The emerging field of DevOps in data science
- Our deep learning DevOps pipeline



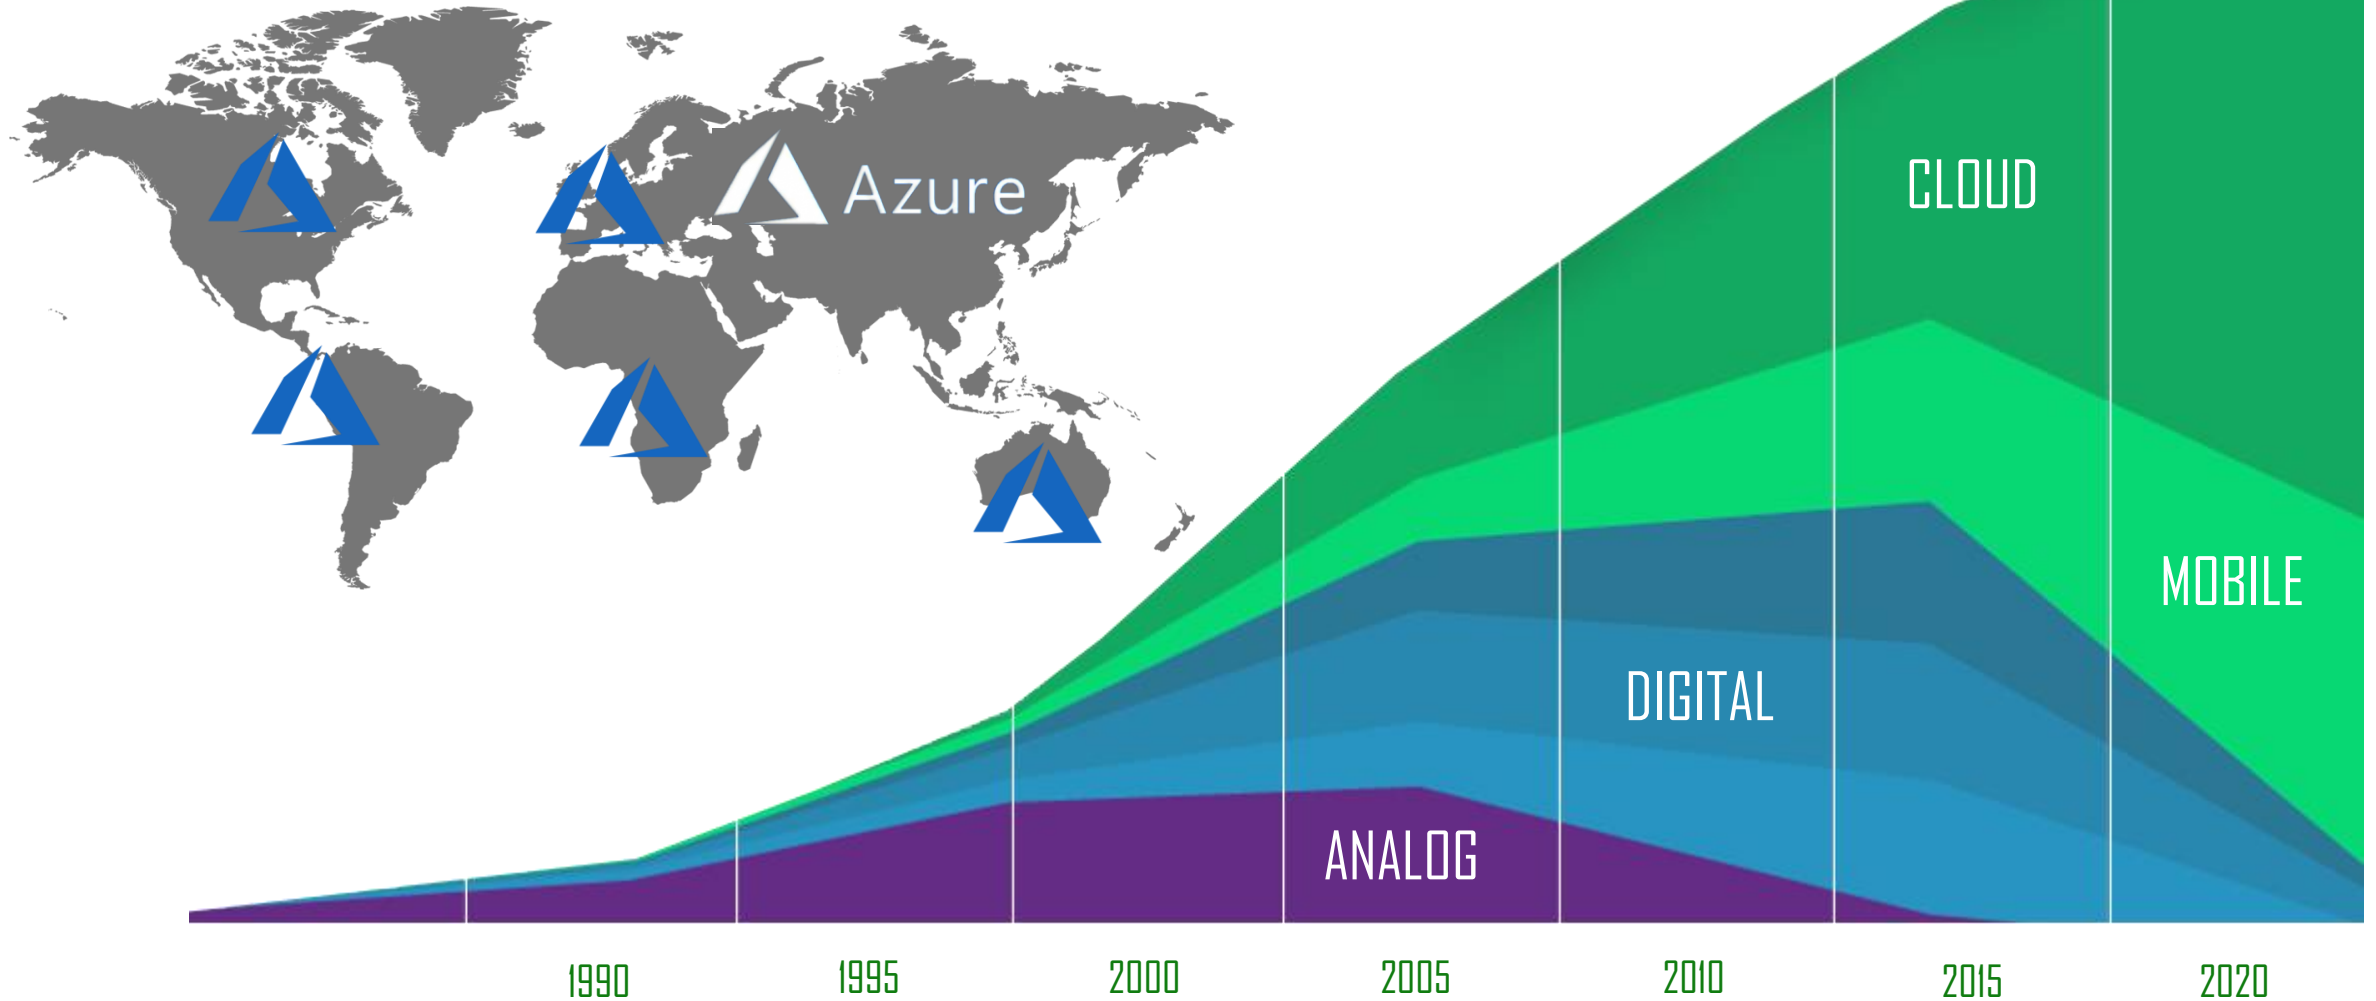


Commercial Software Engineering

Distilled Concepts of Deep Learning




Data + Compute + Algorithms = AI




ImageNet



Alex Krizhevsky



- Alex Krizhevsky wins the ImageNet competition in 2012
- AlexNet
- Heralds a new era of AI
- Same story in language processing, speech recognition etc





Traditional ML

Manual Features



Trainable Classifier



Deep Learning

Representations are hierarchical and trained automatically

Low Level Features



Mid Level Features






High Level Features




Trainable Classifier

Deep Learning = learning a representation




Representations are learned automatically

Universal functional approximators

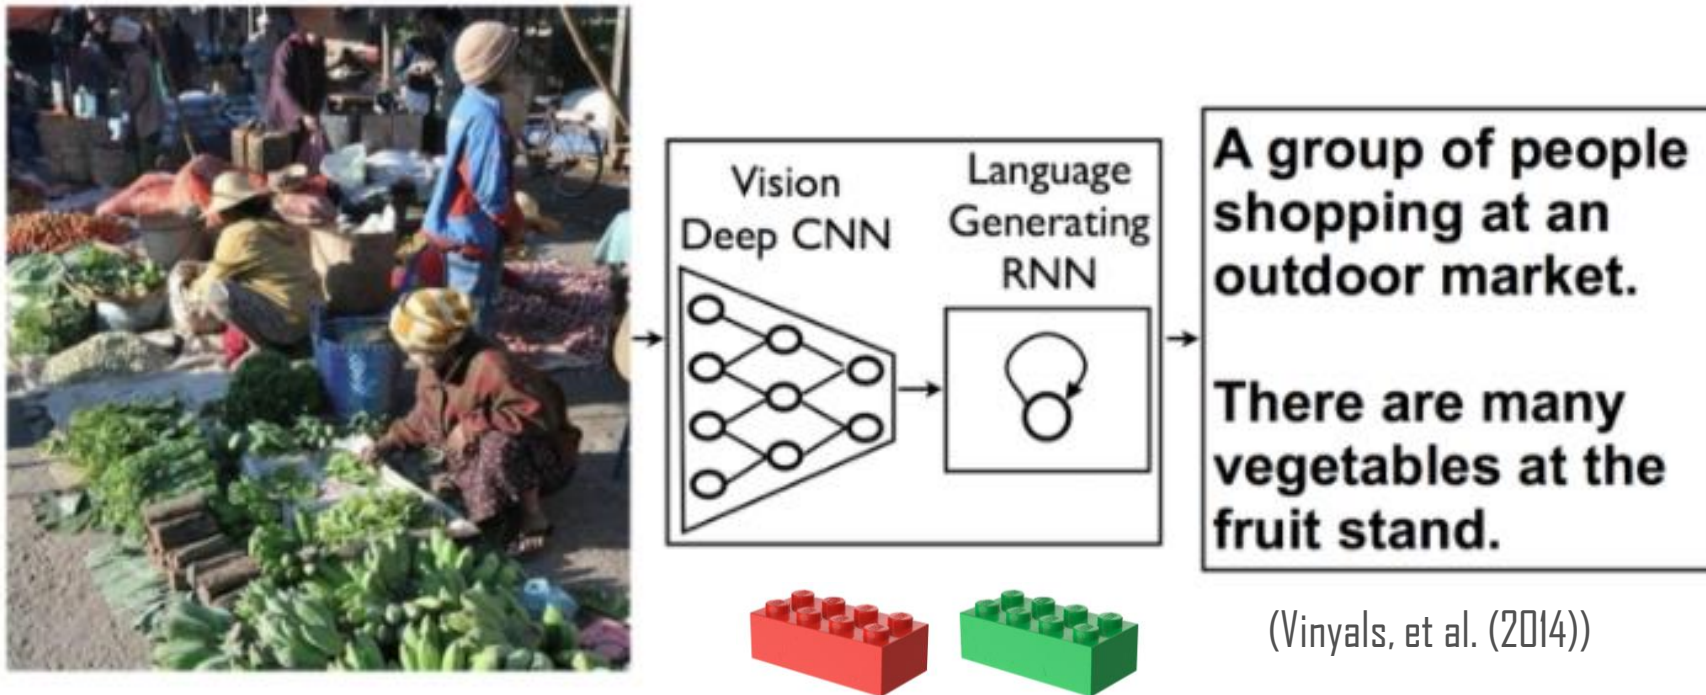


Composability

- Machine learning is becoming a form of software development
- Machine learning models are like software
- There is a dichotomy between the software engineering process and the data science process

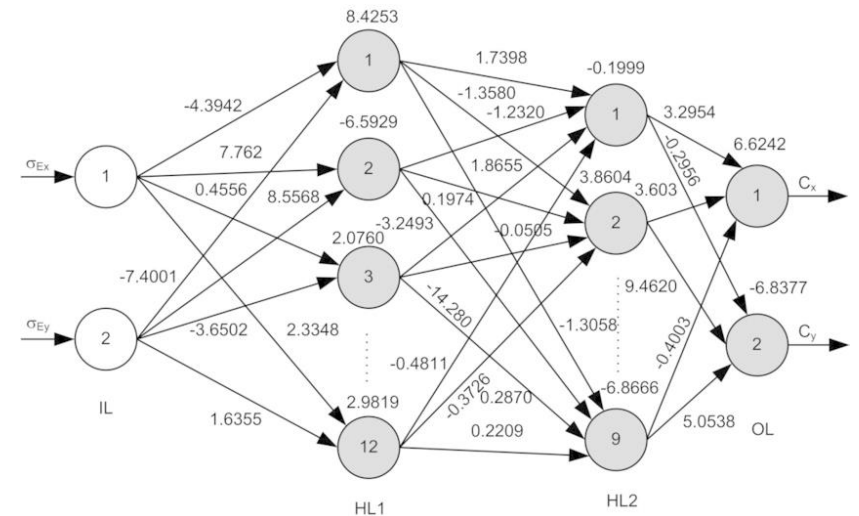


Building predictive architectures like lego blocks



We call this "Software 2.0"


- Software 2.0 is a confection of data generated code and handwritten code
- Some features are **impossible** to write with traditional code i.e. speech synthesis, machine translation, computer vision use cases, game AI
- Computationally homogeneous
- Simple to bake into silicon.
- It is highly portable.
- Composable
- Better than you



Deep learning is not magic; interpolation vs extrapolation

ML algorithms perform very, very badly on data out of range. They don't generalise as well as you think. Trask et al 2018 (DeepMind)

$$f(x)=x$$



When to use deep learning vs classical machine learning

- Large volumes of data
- Unstructured data i.e. images, video, sound
- Modelling temporal or spatial dependencies
- Novel prediction architectures
- Composability, reusability, transferability of models
- Consistency of approach is more conducive to software engineering

When not to use deep learning

- When interpretability is important
- When the data is very small deep learning behaves like classical machine learning i.e. no representation learning
- When you need statistical guarantees

Machine learning models are becoming ubiquitous in software

Software is comprised of some things you can write with code, and some parts which must be generated from data i.e. ML



More companies are becoming data companies

- Companies are becoming data driven
- Diversifying, creating data products & services



Deep learning at confused.com



Our collaboration with confused.com




Confused.com

- Who are confused.com?
- Evolution from BI experts to data science experts
- Old and new skills coming together beautifully
- Millions of records of structured data
- Building new products and services from this data






Confused.com  Beat y the dif
Must be a

1 Car **2** Driver

My registration is:

 Enter registration **FIND CAR**

Make & model

Car make
VOLKSWAGEN ▼

Car model
CC ▼

Year of manufacture
2015 (64, 15, or 65 plate) ▼

Car & driver
details



100+ car
insurers



Car insurance
quotes

Potential services from data

- Only sharing customer data when necessary
- Getting a best-price ballpark in the car showroom with less data
- Annual musical chairs



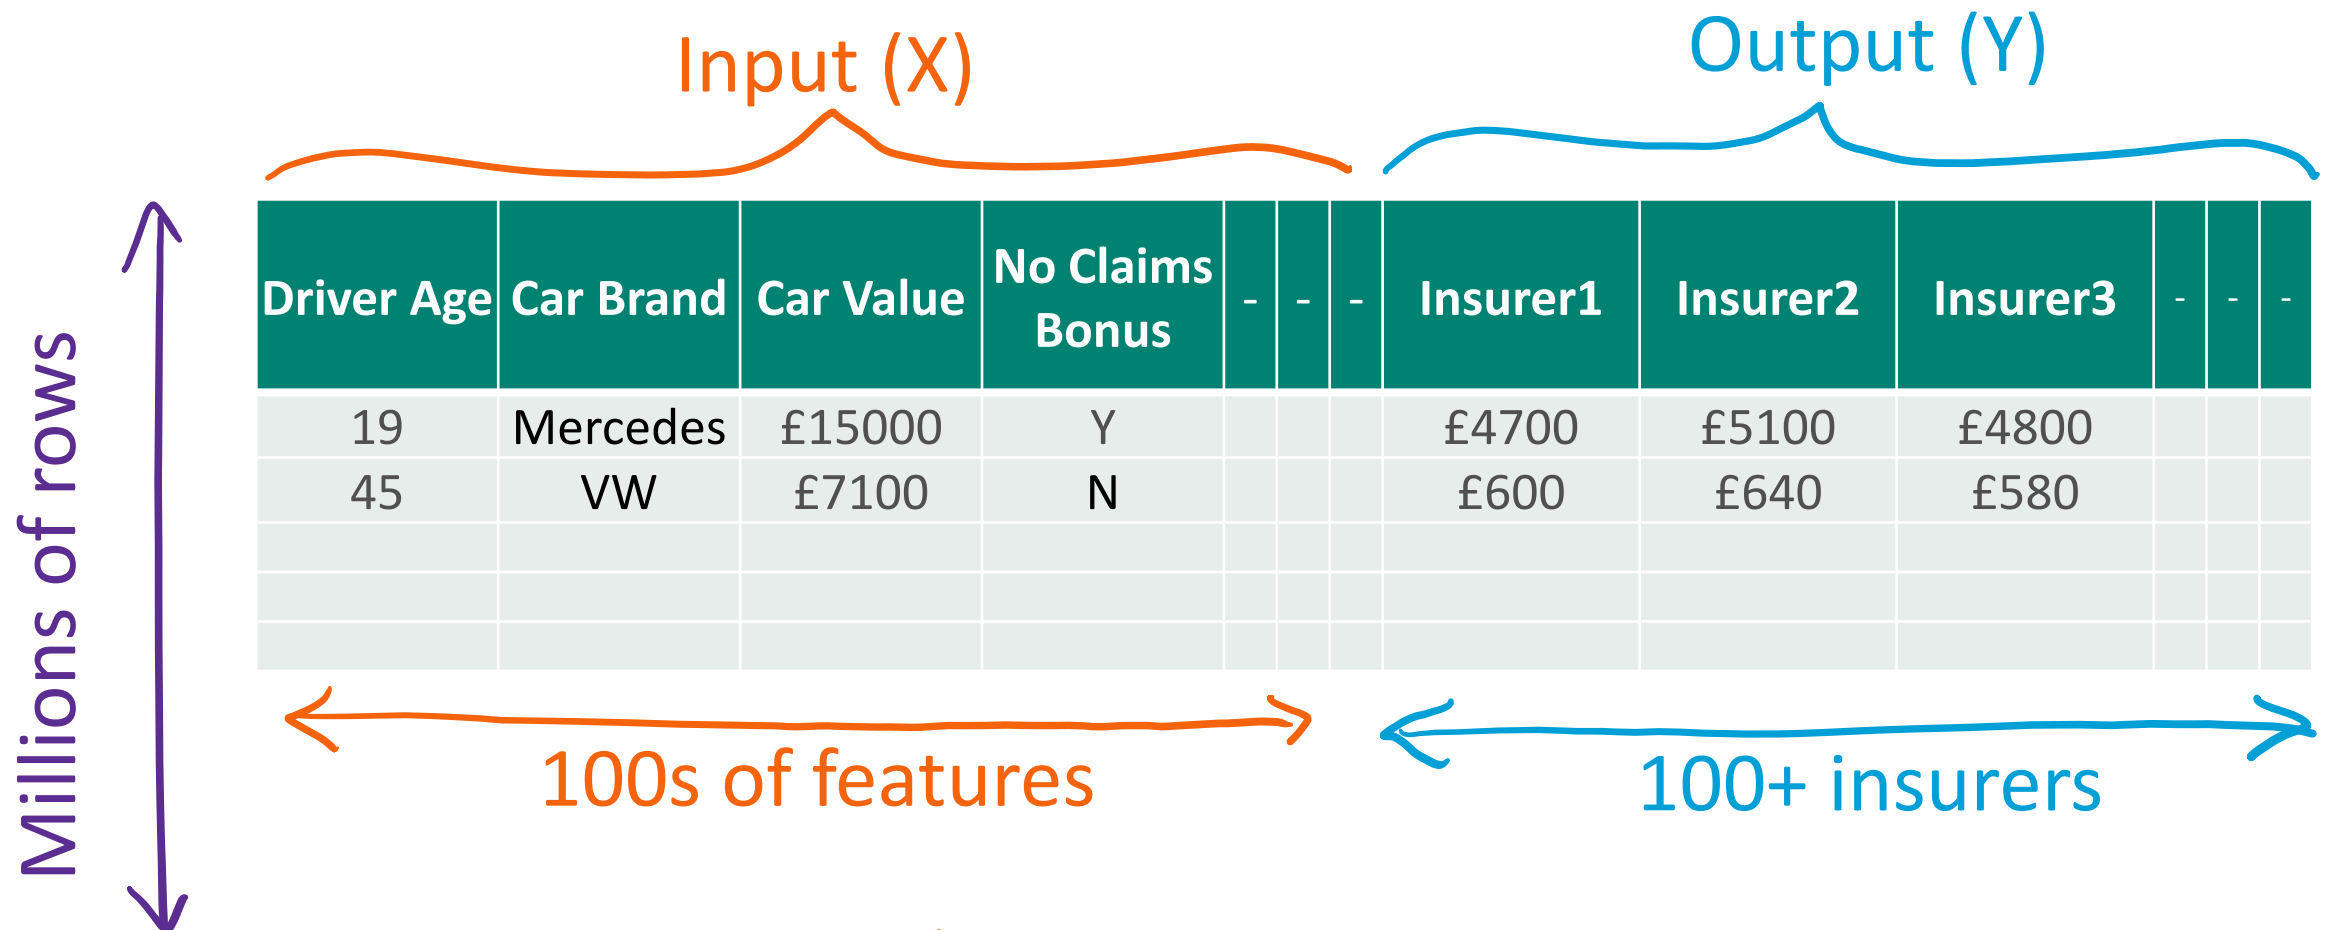
Helping customers get the best insurance premiums with a deep learning model

Our 2-part model:

- Whether or not a customer will be insured (Quote or not?).
- The best insurance quotes that a customer will get (Quote values in £).



Data structure



Prediction model


Model built in:

- Keras, Python, Databricks notebooks
- Keras Functional API to build novel prediction architecture:
DNN with two branches




Driver age	Car brand	Car value	-	-	-	-
23	Audi	£12,000				

Fully
connected
layers

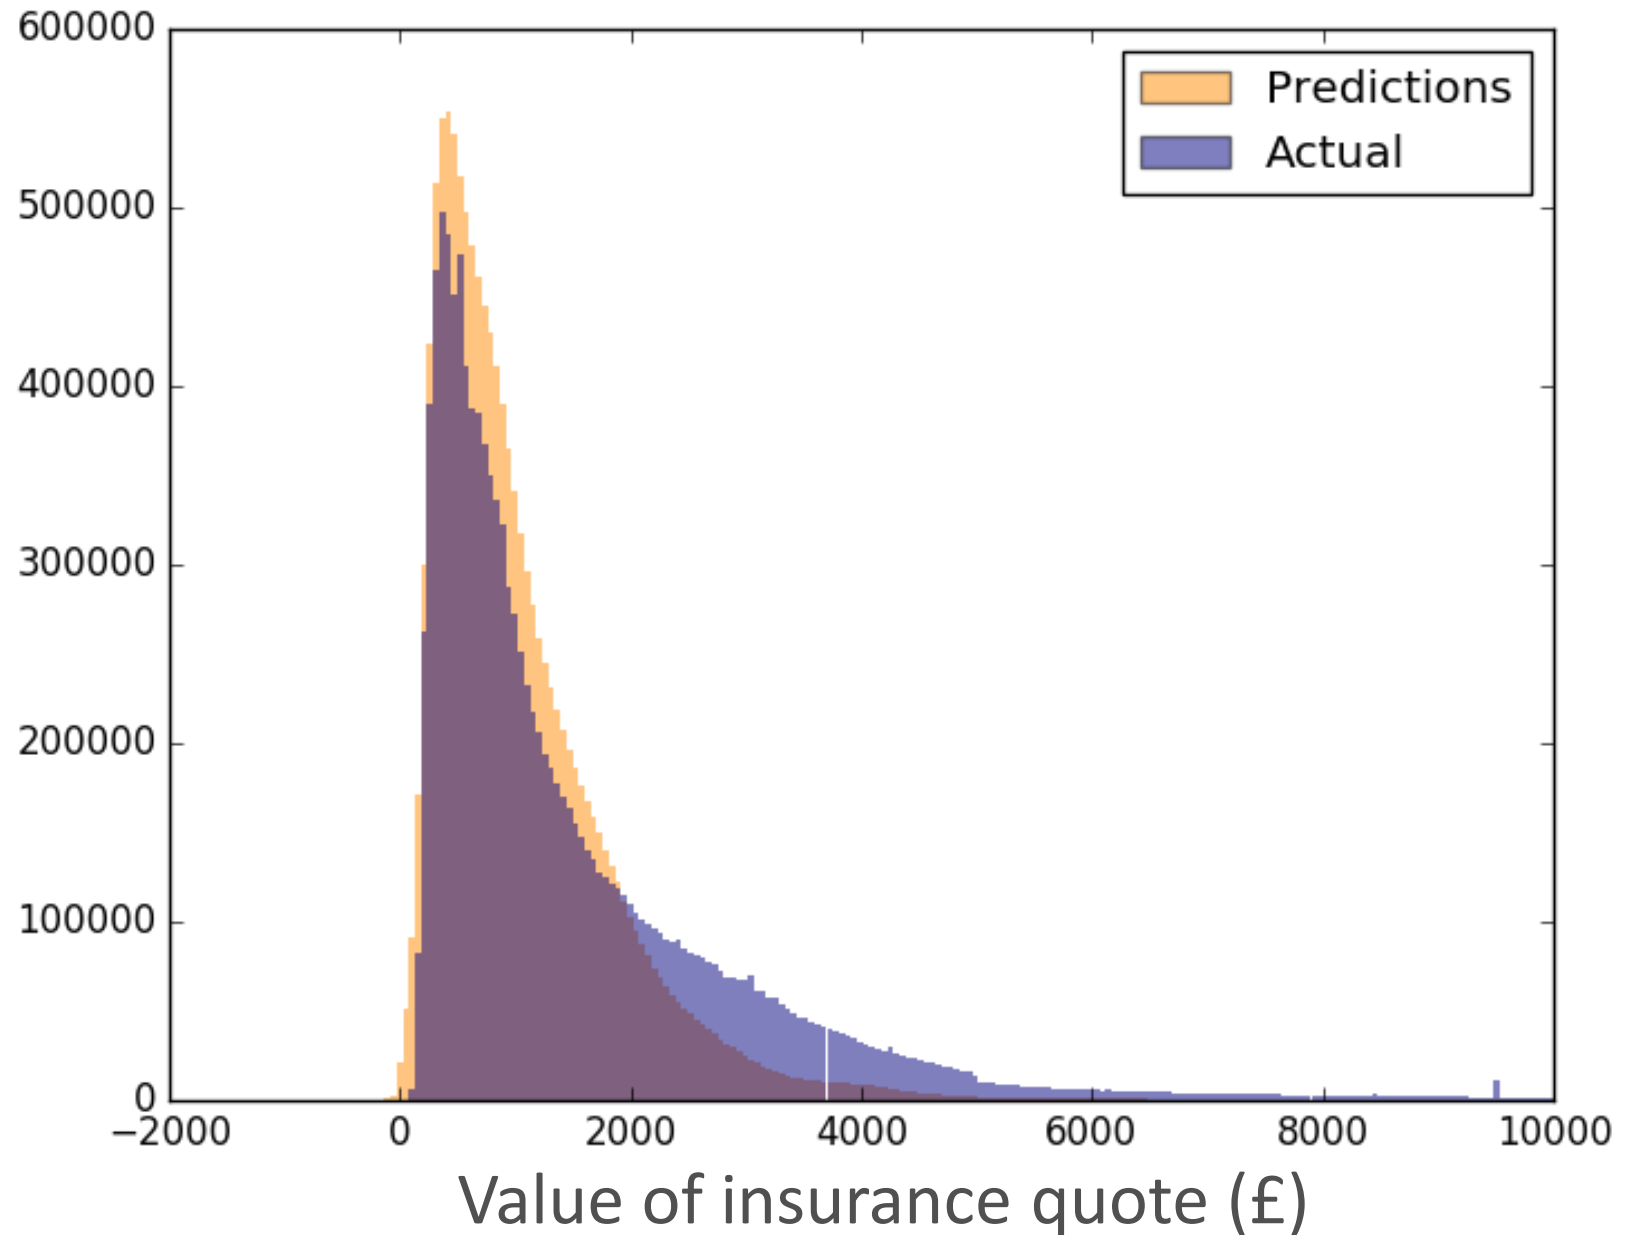



Insurer 1	Insurer 2	Insurer 3	-	-	-
Quote	No Quote	Quote			

Insurer 1	Insurer 2	Insurer 3	-	-	-
£600	-	£1200			



How much will
customers be
insured for?





Nice model, now what?

- Now we leave the world of collaborative notebooks and data science, to the world of deployment.
- How do we get this model into a product?
- How do data scientists tweak the model when it is in production, without breaking everything?



Emergence of DevOps in Data Science




What is DevOps?

People. Process. Products.

DevOps is the union of **people**, **process**, and **products** to enable continuous delivery of value to your end users.

- Donovan Brown







Why DevOps is important

- Your competition is doing it already
- Increase velocity
- Reduce downtime
- Reduce human error
- Any language, any platform



Azure DevOps



Azure Boards

Deliver value to your users faster using proven agile tools to plan, track, and discuss work across your teams.



Azure Test Plans

Test and ship with confidence using manual and exploratory testing tools.



Azure Pipelines

Build, test, and deploy with CI/CD that works with any language, platform, and cloud. Connect to GitHub or any other Git provider and deploy continuously.




Azure Artifacts

Create, host, and share packages with your team, and add artifacts to your CI/CD pipelines with a single click.



Azure Repos

Get unlimited, cloud-hosted private Git repos and collaborate to build better code with pull requests and advanced file management.



<https://azure.com/devops>



Adding data science to DevOps




Culture clash


- Data scientists now ship directly
- DevOps in Data Science is *much* harder



Data Scientist



Software Engineer



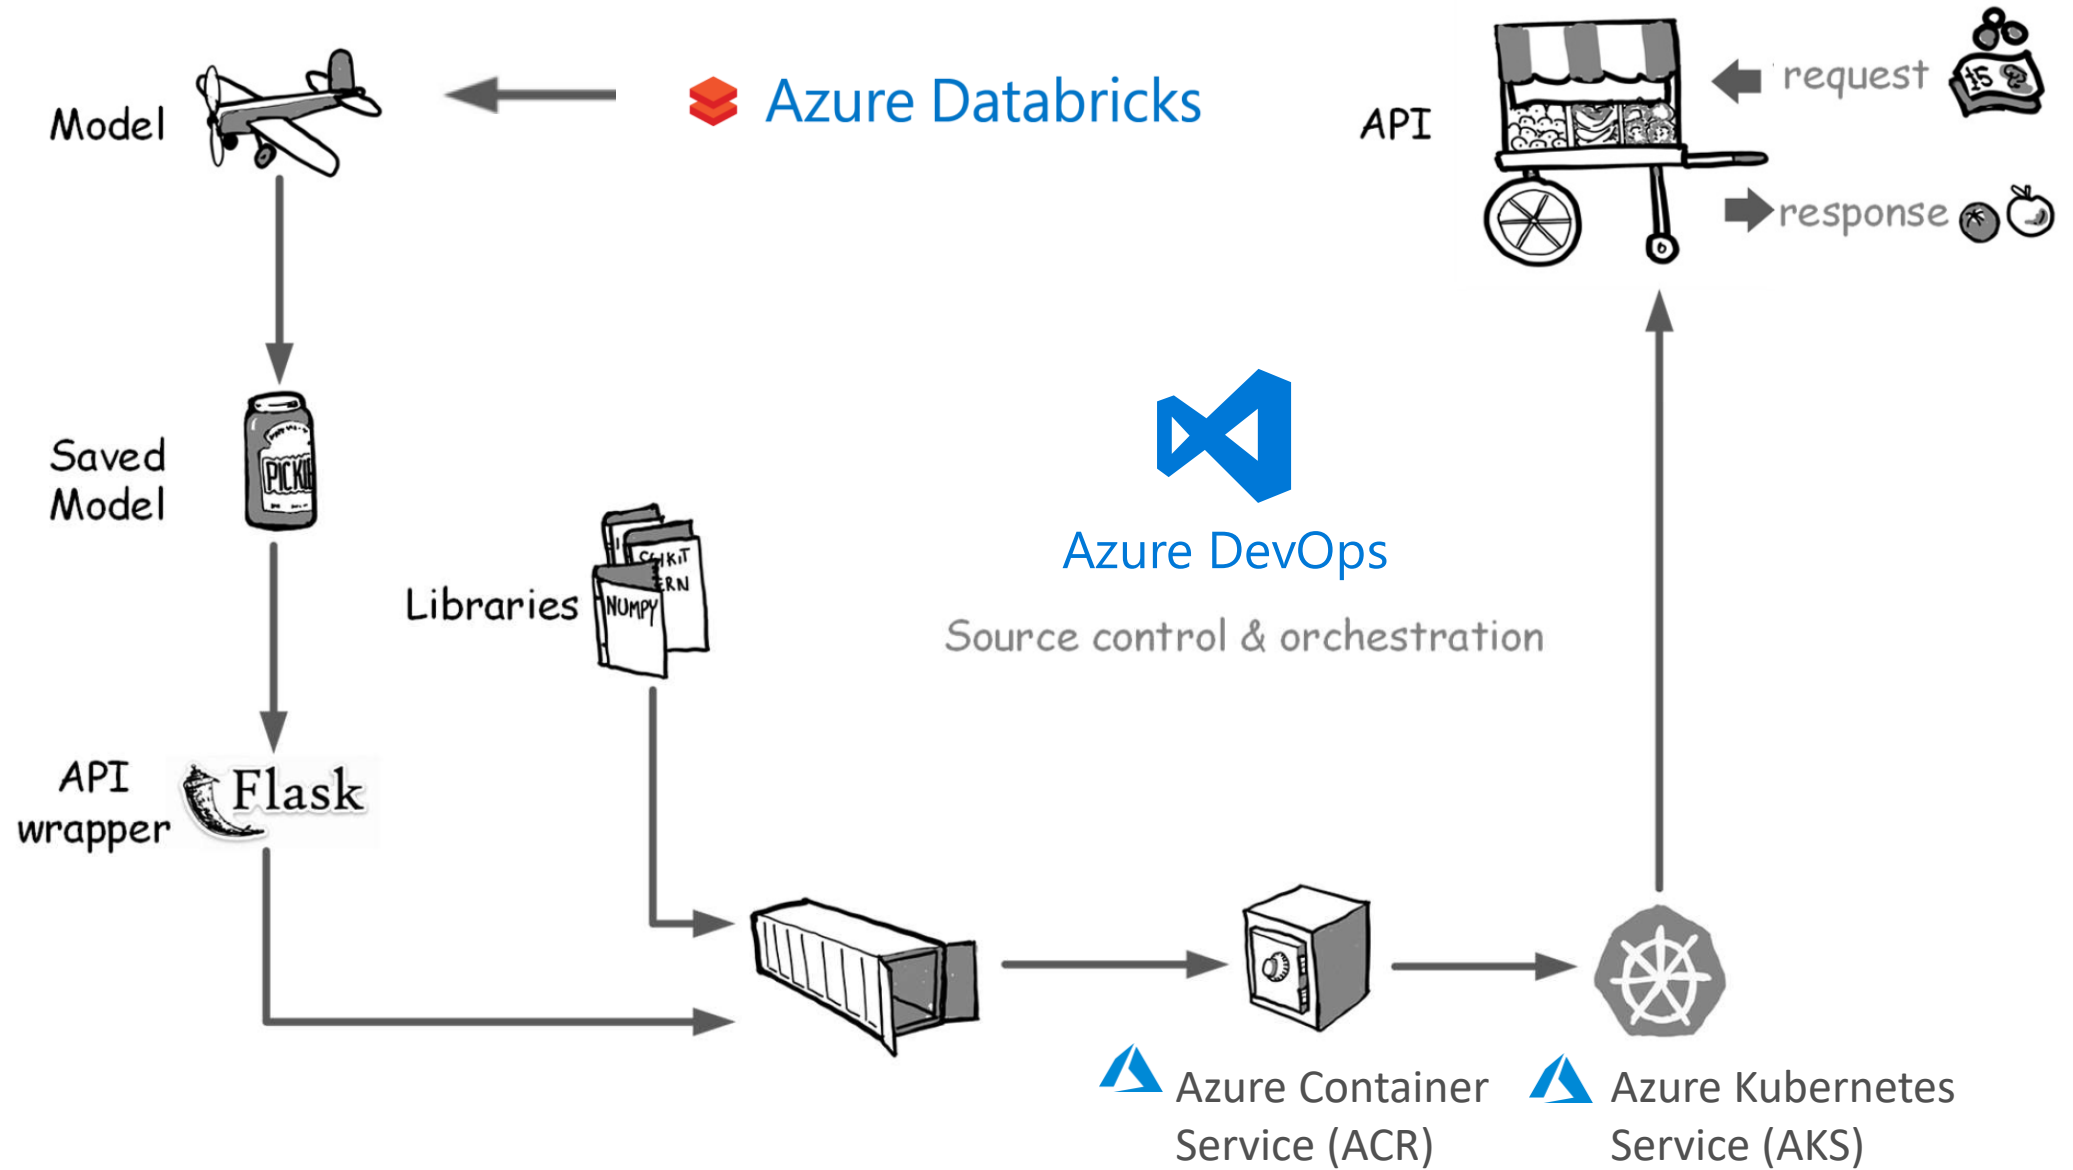
Operations



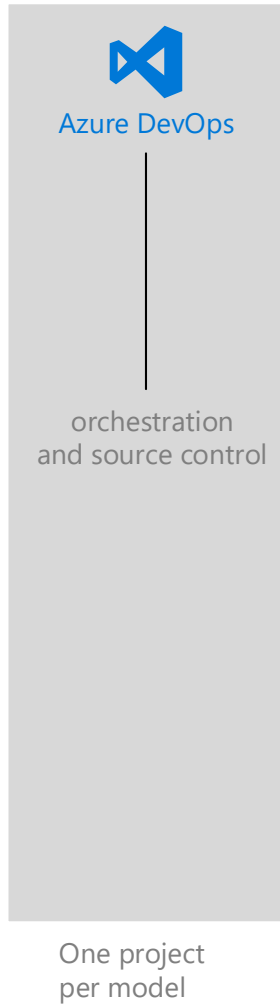
How did confused.com build a pipeline that enabled:

- Software quality/engineering
- Locked down dependencies
- Scalability
- GDPR awareness
- Monitoring
- CICD for models

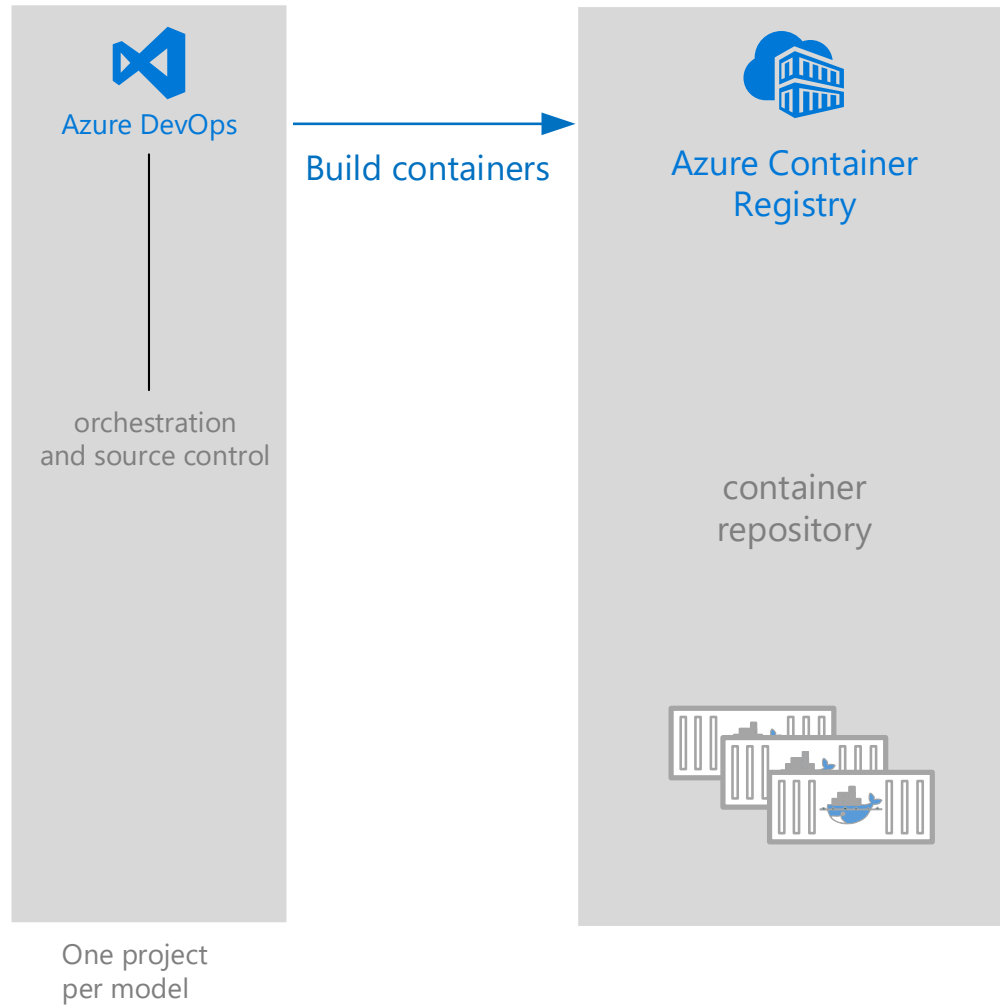




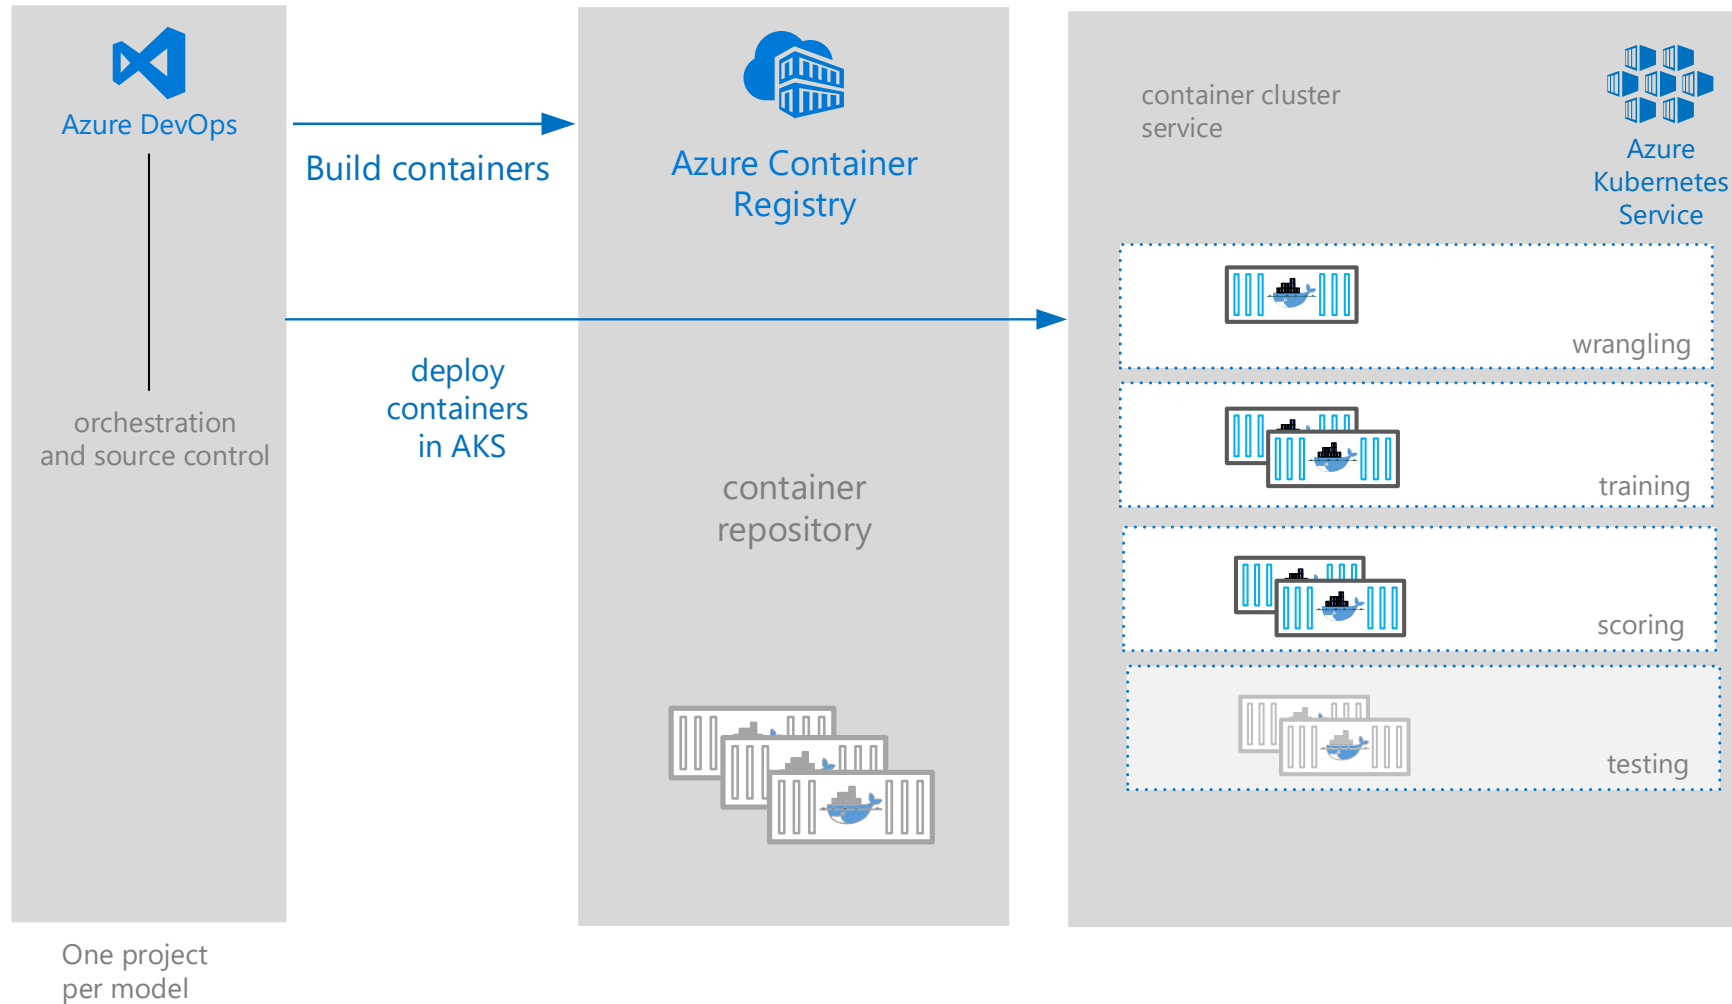
DevOps Architecture



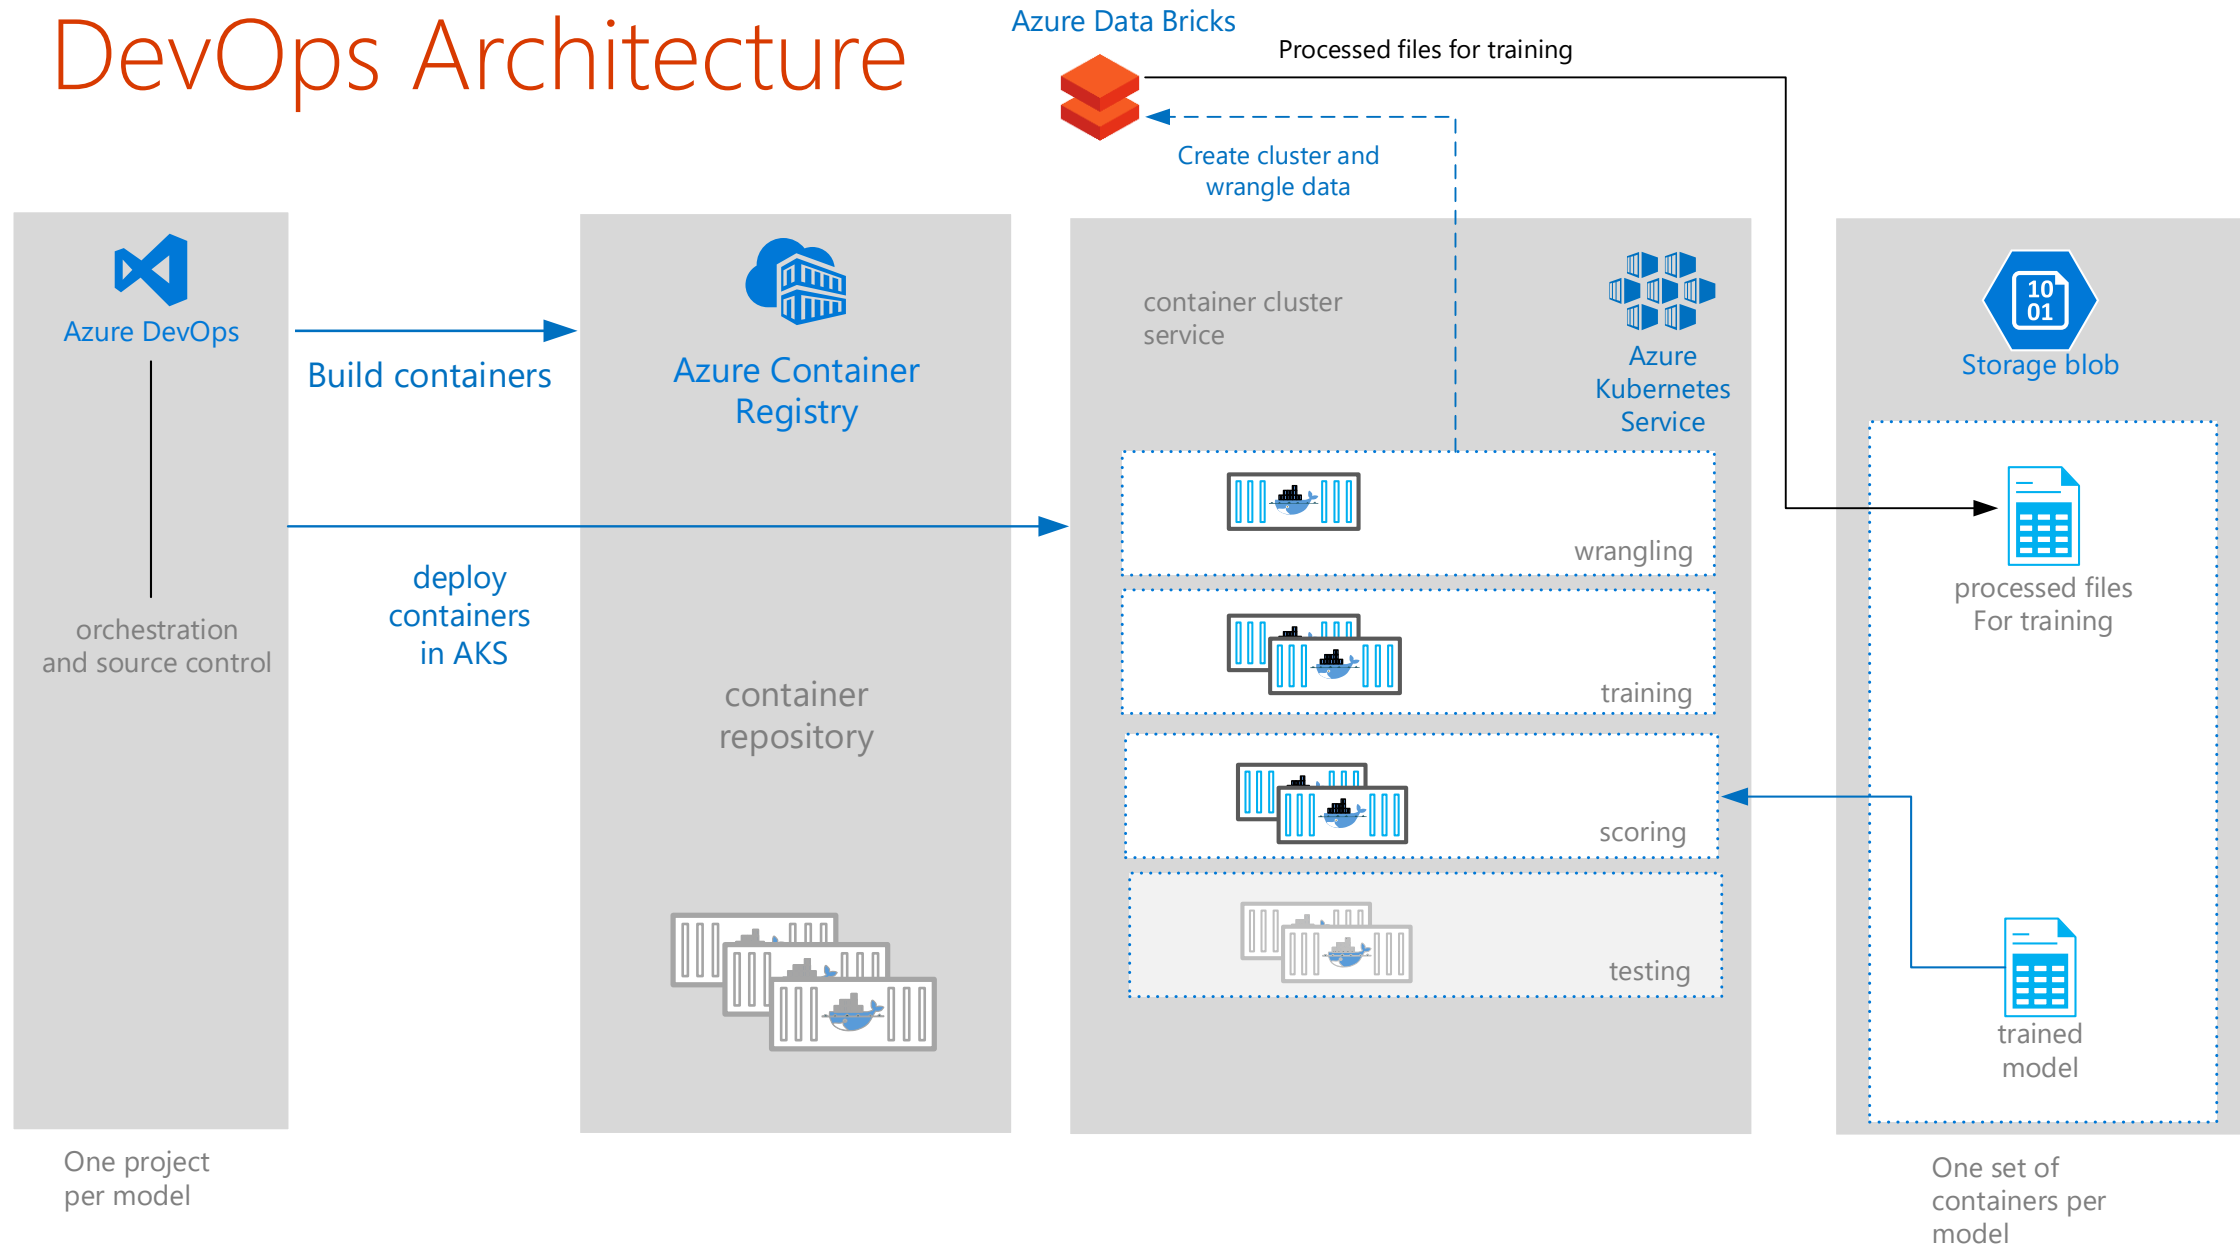
DevOps Architecture



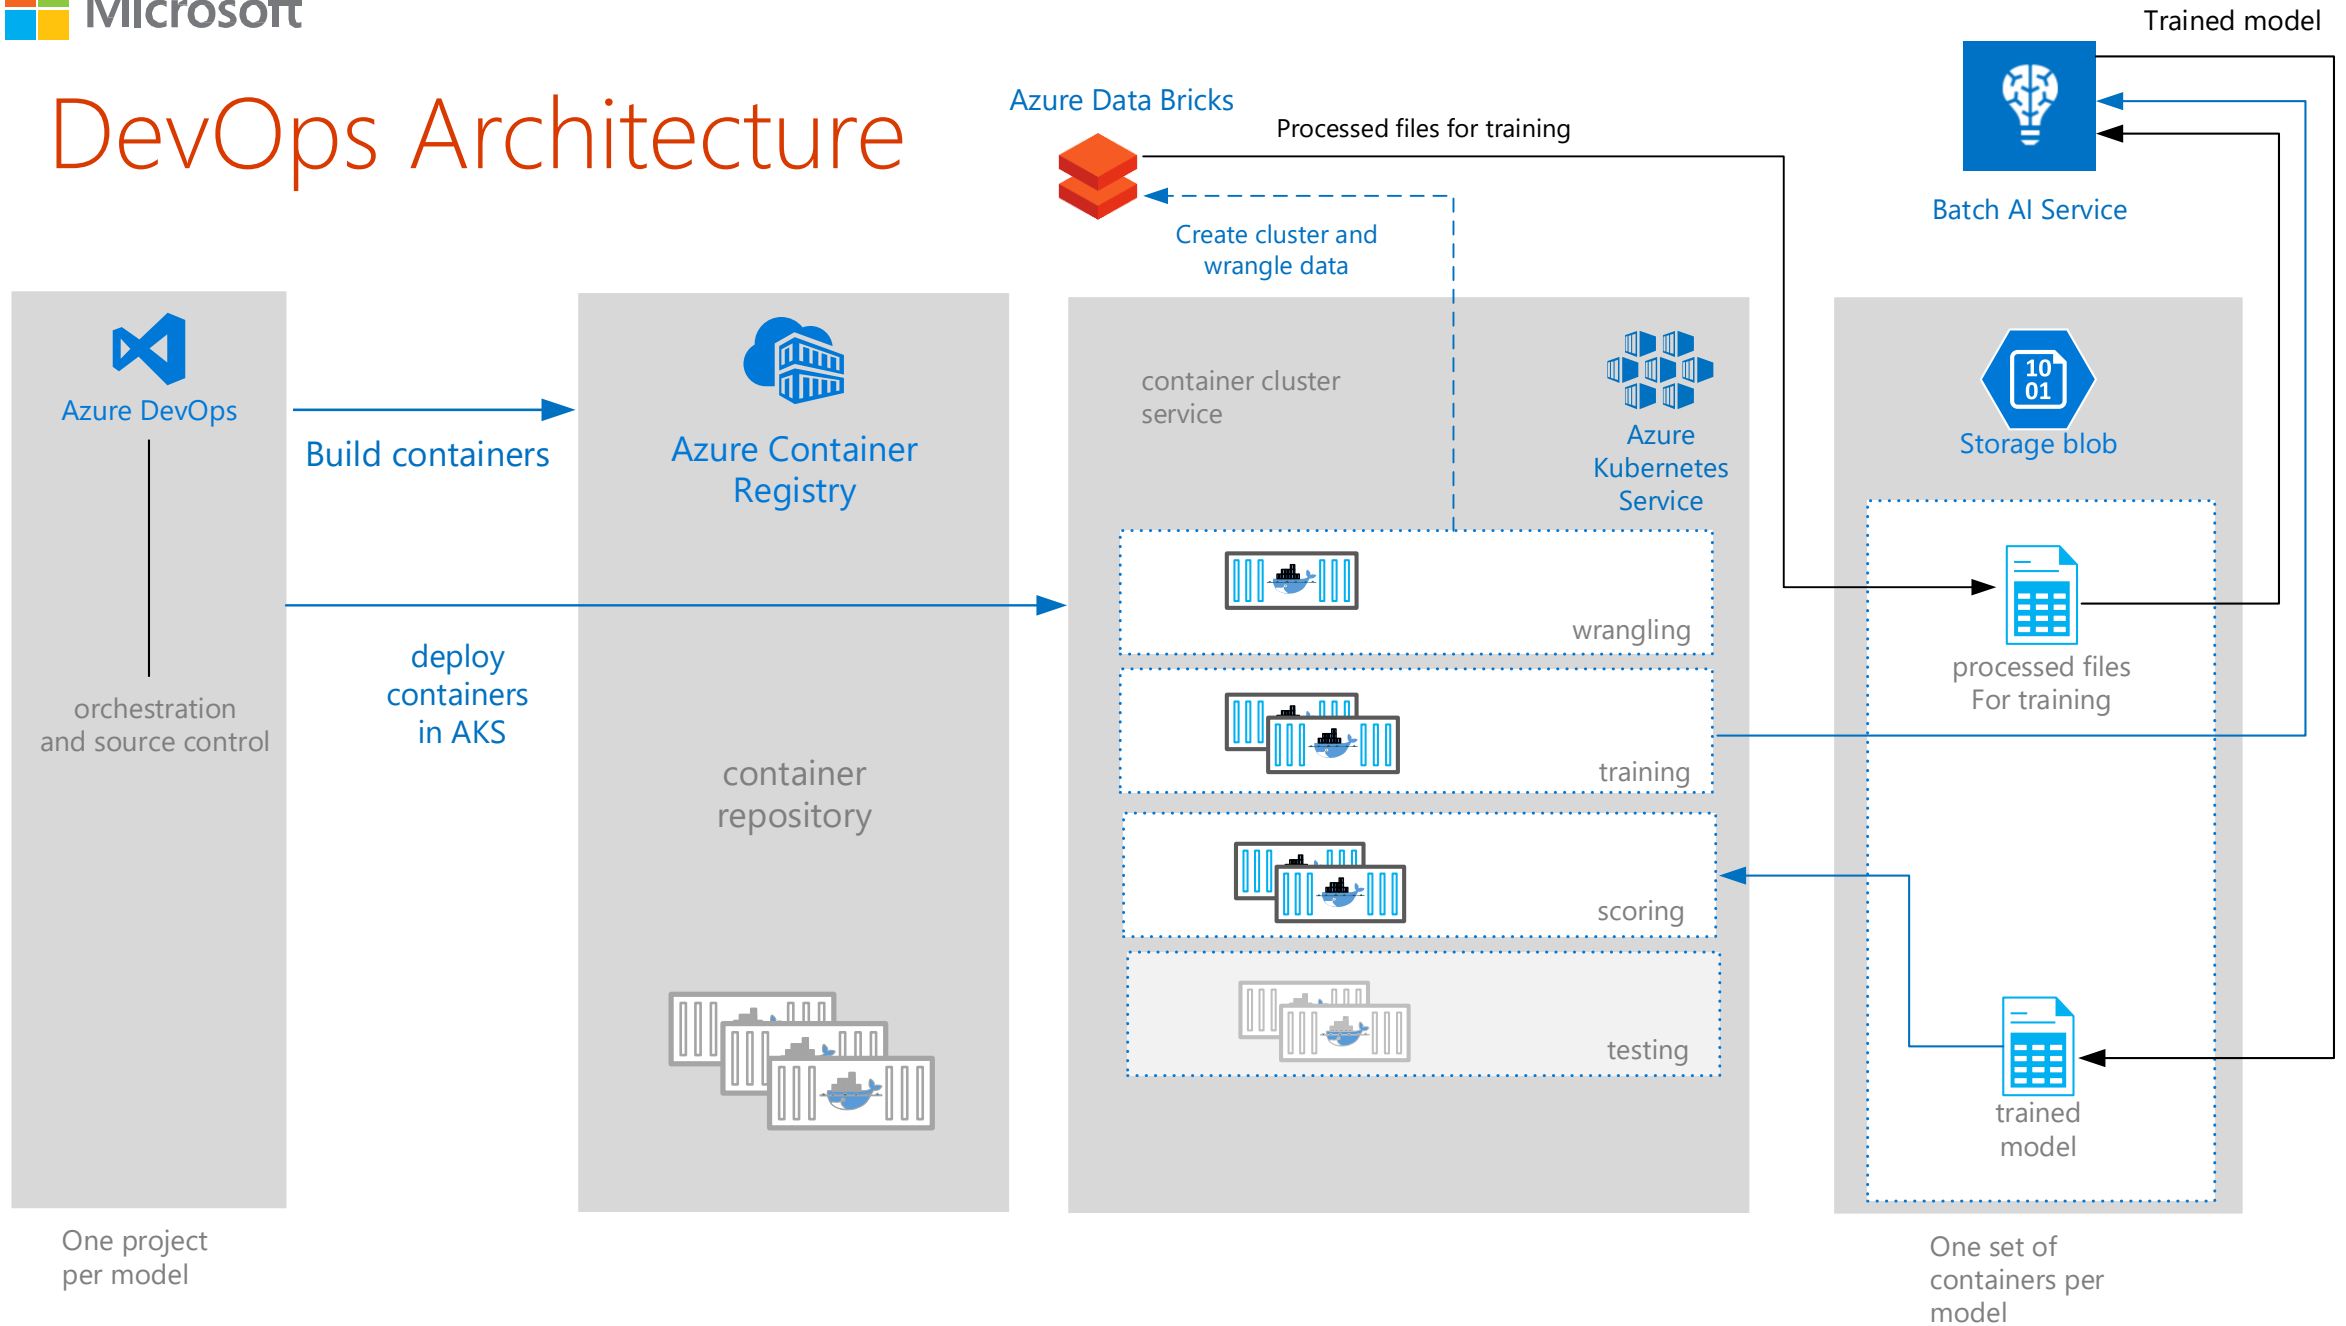
DevOps Architecture



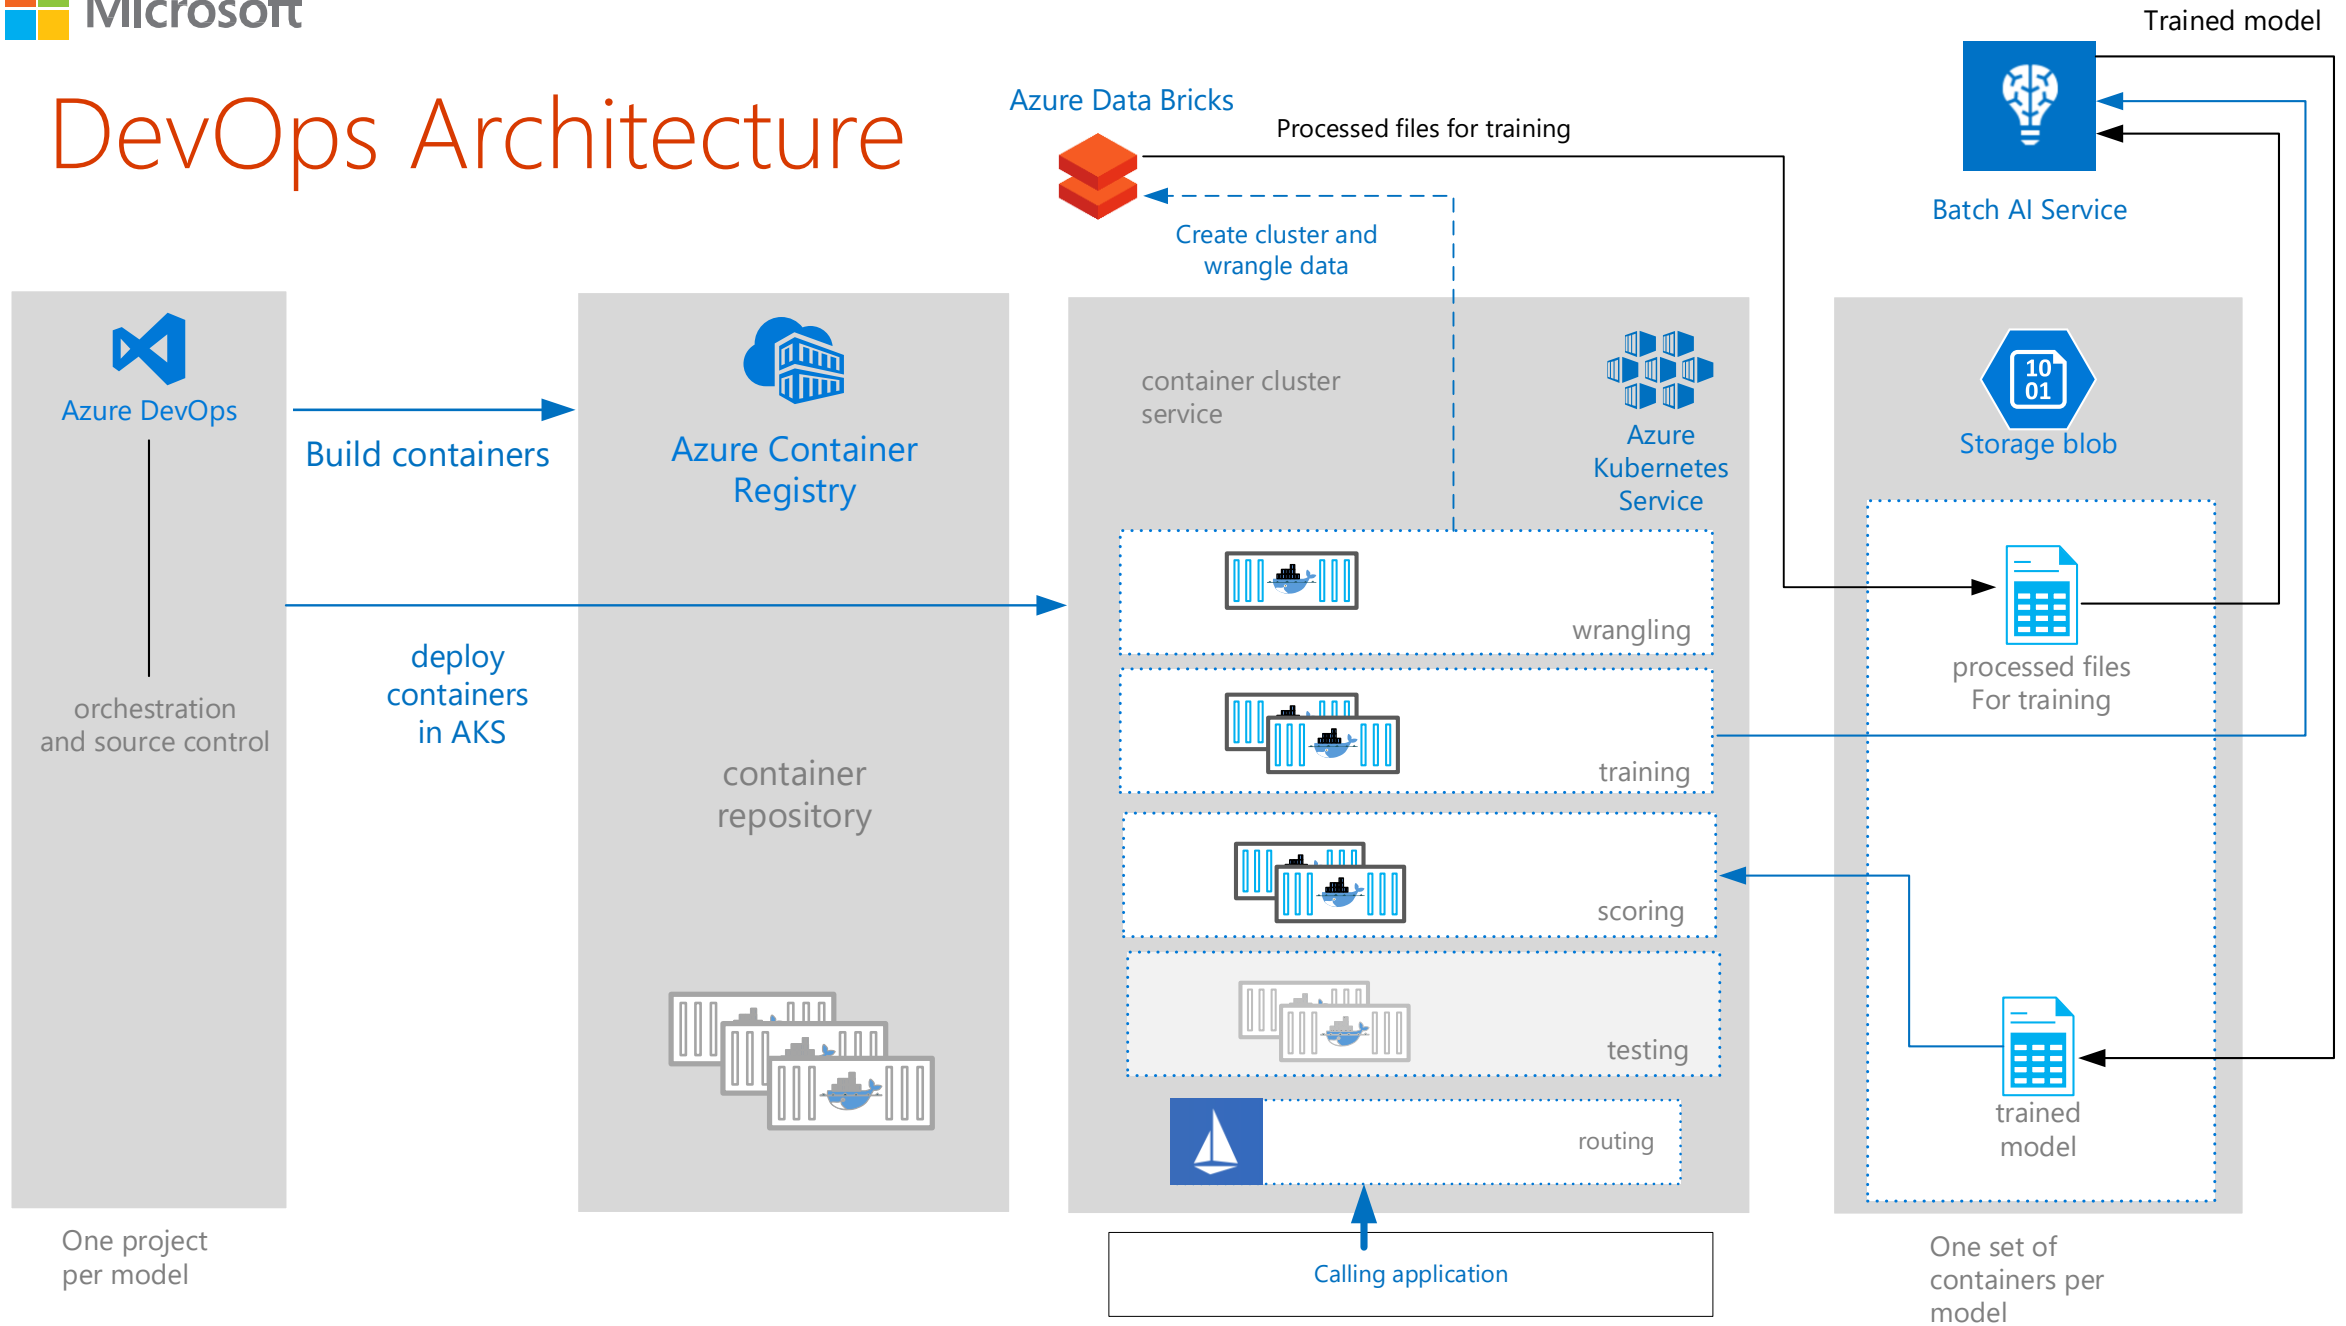
DevOps Architecture



DevOps Architecture



DevOps Architecture





Confused.com's journey

- Data
- DevOps
- Team



Wrapping up



Conclusion & recap

- What we do at Microsoft
- The story of deep learning
- Machine learning is becoming ubiquitous in software
- Many companies are becoming data companies
- The confused.com journey
- Our deep learning model for car insurance
- The emerging field of DevOps in data science
- Our deep learning DevOps pipeline



What's your data story?

- Data?
- DevOps?
- Team?
- Ready!





Things to do next

Check out the Azure docs

DevOps: docs.microsoft.com/en-us/azure/devops

Databricks: docs.azuredatabricks.net/

Containers: azure.microsoft.com/en-us/services/container-registry


Kubernetes: azure.microsoft.com/en-us/services/kubernetes-service

Session Feedback

Please rate this session in the Future Decoded app!



Best intro book on deep learning



Francois Chollet @fchollet

<https://www.manning.com/books/deep-learning-with-python>

<https://github.com/fchollet/deep-learning-with-python-notebooks/>



Thank You!

@Ecsquendor @Dr_Tempest #timtem



at CSE team:

Dr. Tempest van Schaik, Andrew Fryer, Luke Devonshire, Anita Ramanan, Alysha

ed.com team:

ding, Tim Lang, Ben Weatherburn, Vaughan Pugsley, Rex Johnson, Simon
h, Danielle Mansfield, Alice Burnett, David Trickey

Productionising deep learning: The confused.com journey



Dr. Tim Scarfe
@ecsquendor



Dr. Tempest van Schaik
@Dr_Tempest

Dr Tim + Dr Tem aka #TimTem

